



Introduction

Operating Systems

Yu-Ting Wu

1

Operating Systems 2022

Outline

- What is an operating system
- Computer system organization
- Computer system architecture
- Operating system operations
- Resource management
- Security and protection
- Kernel data structures
- Computer system environments

2

2

Operating Systems 2022

What is an Operating System ?

3

3

Operating Systems 2022

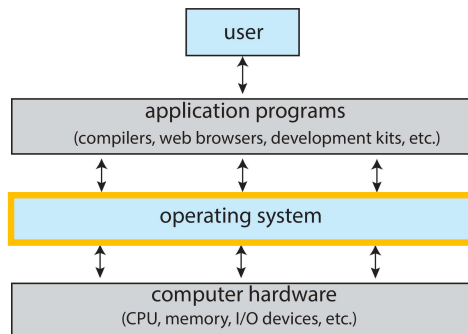
What is an Operating System?

- Operating system (OS) is a **software program** that acts as an intermediary between a user and the computer hardware
 - Execute user programs
 - Make the computer system convenient to use
 - Such that users can focus on their problems
 - Use the computer hardware in an efficient manner

4

4

Computer System Components



An operating system can be considered as a government or environment provider

5

5

User View (Features)

- Varies by the types of the computer



6

6

System View (Tasks)

- A resource allocator**
 - CPU time
 - Memory space
 - File storage
 - I/O devices
- A control program**
 - Control execution of user programs
 - Prevent errors and misuse

7

7

Definition of an OS

- No universally accepted definition**
 - Because of the myriads designs and uses of OSes
 - US Dept. of Justice against Microsoft (1998)
 - The stuff shipped by vendors as an OS
 - The one program **running all the times** on the computer

8

8



Operating Systems 2022

The Goals of an OS

Convenient
for the users

\longleftrightarrow
Conflict !

Efficiency
for the computer system

Windows Vista Windows XP

9

9

Operating Systems 2022

OS and Hardware

- Can change due to new computer architectures and hardware devices

(multi-core CPU, GPU)

OS

resource
management
 \longrightarrow
 \longleftarrow
new features
support
(multi-tasking)

Hardware

Learn OSs by tracing their evolution enables us to predict what they will become !

10

10

Operating Systems 2022

The Development of an OS

- In the past time, operating systems are usually implemented with low-level languages
- Later, **high-level languages** are used for developing an operating system because of the increasing complex functionalities and porting issues
 - But system calls are still implemented by assembly language

\Rightarrow As a result, more and more people can involve in the development of operating systems

- Another trend is **modulation**


11

11

Operating Systems 2022

Booting: How an OS is Launched

- Bootstrap program**
 - Simple program to initialize the system and load the kernel
 - Typically stored in ROM or EPROM
 - Initialize the entire system, including CPU registers, device controllers, memory, ... etc.
- Steps**
 - Booting
 - Load kernel
 - Start system daemon (e.g., login)
 - Kernel (HW/SW) interrupt driven



12

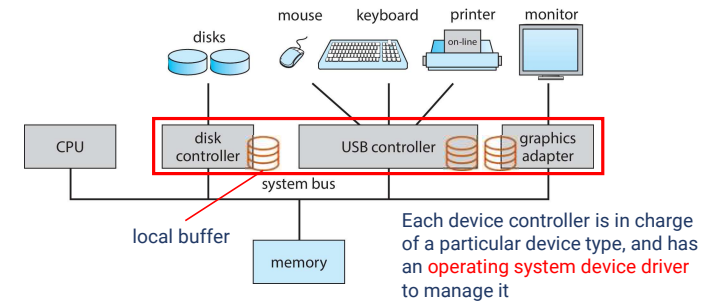
12

Computer System Organization

13

Computer System Organization


- CPU (or CPUs) and device controllers connect through common **bus**, which provides access to memory



14

Computer System Operations


- CPU moves data from/to main memory to/from local buffer for executing programs



15

Computer System Operations (cont.)

- IO: from the device to local buffer of controller
 - Use **interrupt** to inform CPU that it has finished its task



16

Computer System Operations (cont.)

- A simplest design for accessing devices
 - Use instructions to test when a device is ready
 - Busy/wait

```
#define OUT_CHAR 0x1000 // device data register
#define OUT_STATUS 0x1001 // device status register
```

```
current_char = mystring;
while (*current_char != '\0') {
    poke(OUT_CHAR, *current_char);
    while (peek(OUT_STATUS) != 0); // busy waiting
    current_char++;
}
```

17

17

Interrupt

- Busy/wait is inefficient
 - CPU cannot do other task while testing device
- Interrupt provides a way to change the **flow of control** in the CPU

Hardware interrupt (signal)

- Service requests from one of the devices
 - Ex: keyboard, mouse click, etc.



Software interrupt (trap)

- Invalid memory access
- Software error
 - Ex: division by zero
- System calls
 - Request for system services

```
#include <stdio.h>

int main(int argc, const char * argv[]) {
    FILE* fp = fopen("test.txt", "r");
    if (fp) {
        printf("Not NULL");
    } else {
        printf("NULL");
    }
}
```


→ system call (open)

18

18

Interrupt Timeline

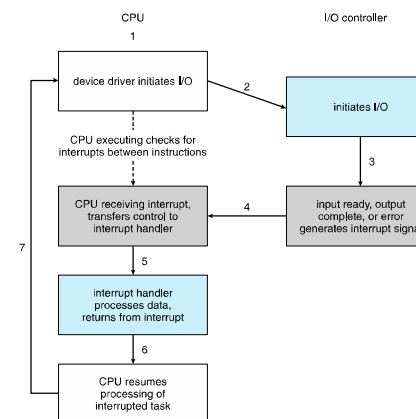
- Transfer control to the interrupt service routine
- Must save the address of the interrupted instruction



19

19

Interrupt-driven I/O Cycle

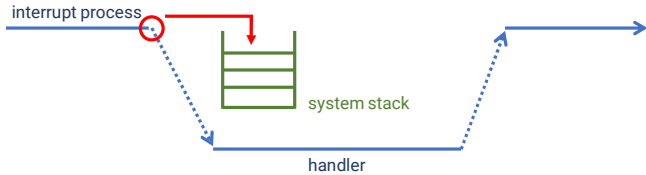


20

20

Operating Systems 2022

Interrupt Handling



The diagram illustrates the interrupt handling process. It shows a horizontal line representing the 'interrupt process'. A red circle marks a point on this line. A red arrow points from this circle to a green box labeled 'system stack'. A blue arrow points from the 'system stack' to a horizontal line labeled 'handler'. A dashed blue arrow points from the 'handler' back to the red circle on the 'interrupt process' line. A solid blue arrow continues from the 'interrupt process' line after the red circle.

- Saving of the address of the interrupted instruction
 - Fixed location
 - Fixed location per interrupt type
 - Stacks
- Interrupt disabling /enabling issues

21

21

Operating Systems 2022

Interrupt Handling

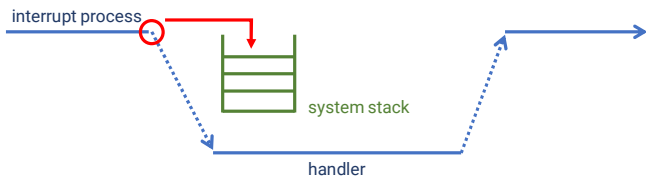
- Determine which interrupt service routine should be called
 - **Generic handler**
 - A specific program for handling all types of interrupts
 - Call the corresponding interrupt service routine after checking
 - Inefficient but can handle infinite types of interrupts
 - **Interrupt vector**
 - Use a vector to store all interrupt service routines
 - Hardware jumps to the corresponding interrupt service routine based on an **interrupt number** (ID)
 - Efficient but can only handle a fixed number of interrupt types
- Current OSs usually hybrid of the two strategies

22

22

Operating Systems 2022

Interrupt Handling Procedure (Summary)



The diagram illustrates the interrupt handling procedure. It shows a horizontal line representing the 'interrupt process'. A red circle marks a point on this line. A red arrow points from this circle to a green box labeled 'system stack'. A blue arrow points from the 'system stack' to a horizontal line labeled 'handler'. A dashed blue arrow points from the 'handler' back to the red circle on the 'interrupt process' line. A solid blue arrow continues from the 'interrupt process' line after the red circle.

- Steps:
 - Save interrupt information
 - OS determine the interrupt type
 - Call the corresponding handlers
 - return to the interrupted job by restoring the information of original process

23

23

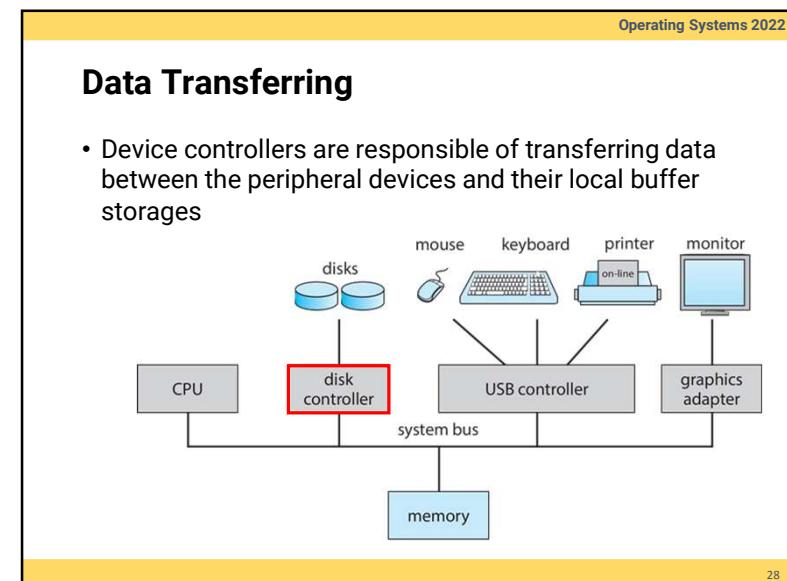
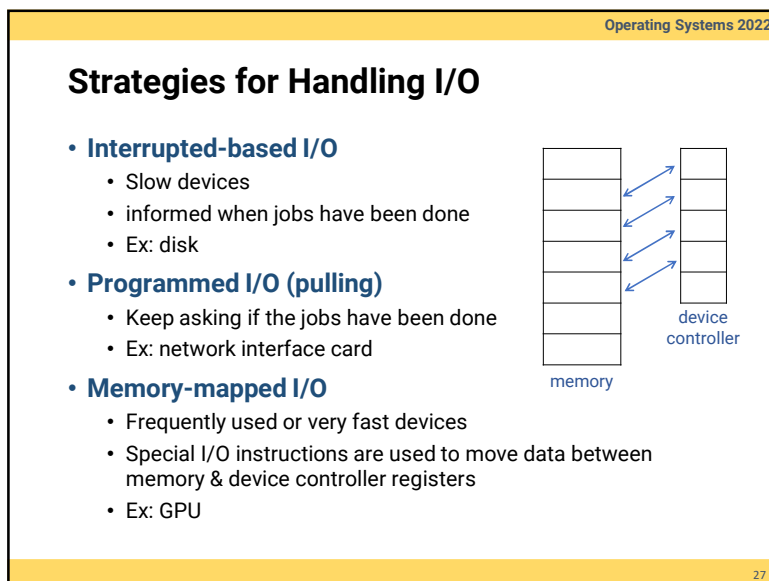
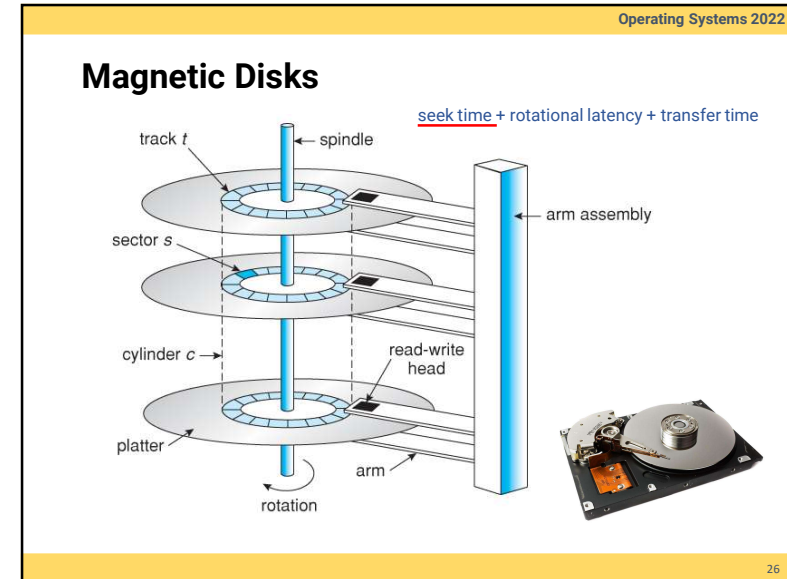
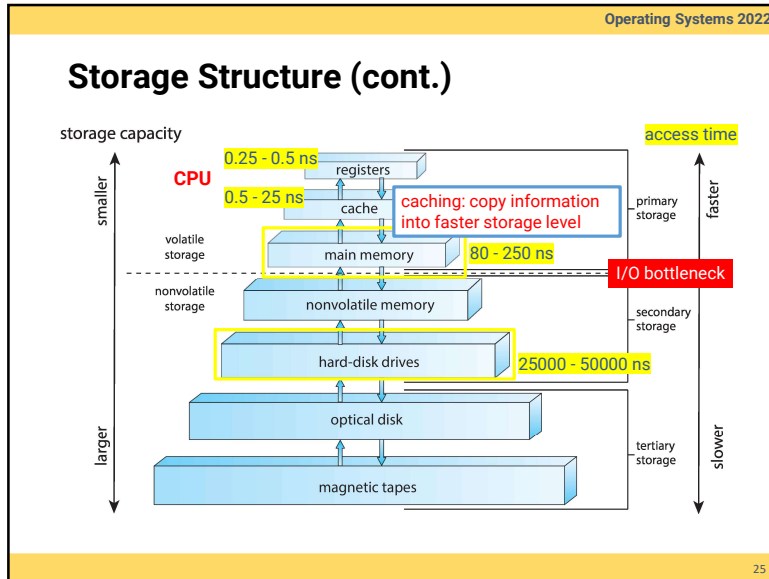
Operating Systems 2022

Storage Structure

- **Main memory**
 - The only large storage media that the CPU can access directly
 - Typically volatile
- **Secondary storage** (ex: HDD, USB sticks, CD, DVD, ...)
 - Extension of main memory that provides large storage capacity
 - Typically nonvolatile
- Organized in hierarchy based on
 - **Speed**
 - **Cost**
 - **Volatility**

24

24



I/O Operation Procedure

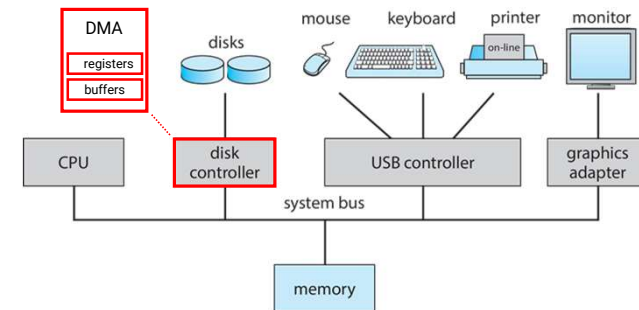
- CPU setups specific controller registers within the controller
- Read / Write
 - Read: devices → controller buffers → memory
 - Write: memory → controller buffers → devices
- Notify the completion of the operation by triggering an interrupt

29

29

DMA: Direct Memory Access

- Transfer blocks of data without bothering CPU



30

30


DMA: Direct Memory Access (cont.)

- **Goal**
 - Device controller can transfer blocks of data from buffer storage to main memory without CPU intervention
 - Only one interrupt is generated per block (rather than per byte), thus avoiding CPU handling excessive interrupts
- **Procedure with DMA**
 - Execute the device driver to setup the registers of the DMA controller
 - DMA moves blocks of data between the memory and its own buffers
 - Transfer from its buffers to its devices
 - Interrupt the CPU when the job is done

31

31

Storage Structure Summary



32

32

Computer System Architecture

33

33

Single-Processor Systems

- One main CPU per system
 - Control other low-end processors, e.g., disk controller microprocessors
 - Ex: earlier desktop or mobile devices




34

34

Multi-Processor Systems


• Tightly coupled

- More than one processor in close communication sharing bus, memory, and peripheral devices



• Loosely coupled

- Otherwise (such as distributed systems)
- Each machine has its own memory



35

35

Multi-Processor Systems (cont.)

• Symmetric model

- Each processor in the system runs an identical copy of the OS

• Asymmetric model

- master-slave
- Commonly seen in extremely large systems

• Task allocation strategies

- Dynamic allocation
- Pre-allocation

36

36


Multi-Processor Systems (cont.)

• Advantages of multi-processor systems

- **Speedup**: better throughput
- **Lower cost**: building one small fast chip is very expensive
- **More reliable**: Graceful degradation and fail soft

• The recent trend: from a fast single processor to lots of processors

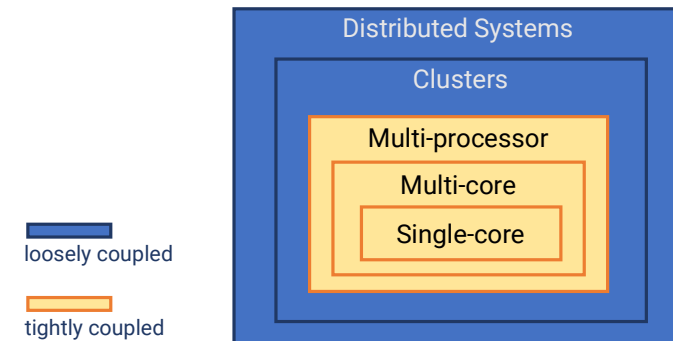
- **Multiple cores** over a single chip
- Hyperthreading (logical core)



37

37

Computer System Architecture Summary



38

38

Operating System Structure


39

39

Simple Batch System

• Workflow

- Users submit data (program, data, control card)
- Operator sort jobs with similar requirement
- OS simply transfer control from one job to the next
 - Resident monitor: automatically transfer control from one job the next



40

40


Simple Batch System (cont.)

• Problem of batch systems

- One job at a time → multi-programming
- No interaction between users and jobs → time sharing
- CPU is often idle

• Spooling (Simultaneous Peripheral Operation On-Line)

- Replace sequential-access devices with random-access devices (disks)



- Ex: printer
 - Data for printing is first written into a directory as files, then printed by the printer

41

41

Multi-Programming

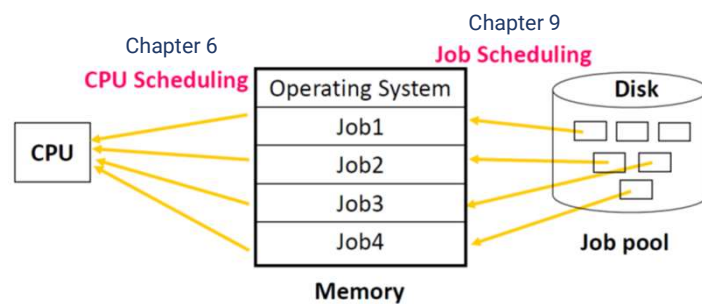
- Single user cannot always keep CPU and I/O devices busy
 - Even with spooling, disk I/O is still too slow compared to CPU and memory
- Put multiple programs in memory
- OS organizes jobs so that the CPU always has one to execute
 - When job has to wait (e.g., for I/O), OS switches to another job
 - Increase CPU utilization
 - Issue: job and CPU scheduling

42

42

Multi-Programming (cont.)

- Job scheduling and CPU scheduling




43

43

Time-Sharing (Multi-Tasking)

- CPU switches jobs frequently so that users can interact with each job while it is running
 - A logical extension of multi-programming
 - **Interactivity!**
 - Response time should be less than 1 sec.
- Brings lots of new issues
 - Online file system (chapter 11 and chapter 12)
 - Virtual memory (chapter 10)
 - Allow execution of processes not completely in memory
 - Job synchronization (chapter 7 and chapter 8)
 - Protection and security



44

44

Batch, Multi-Programming, Time-Sharing

	Batch	Multi-Programming	Time-sharing (Multi-Tasking)
System Model	single user single job	multiple programs	multiple Users multiple programs
Purpose	simple	resource utilization	interactive response time
OS Features	N.A	CPU scheduling I/O system	file system virtual memory synchronization

45

45

Resource Management

46

46

Process Management

- Process (running program) is an **active** entity using physical and logical resources
 - Memory, I/O buffers, data, ...etc.
 - Data structure representing current activities:
 - Program counter
 - Stack
 - Data section
 - CPU registers
- Activities
 - Process creation, suspension, resumption, and deletion
 - Process synchronization
 - Process communication
 - Deadlock handling

context

47

47

Memory Management

- To execute a program, all (or part) of the instructions and its needed data must be in memory
- Memory management
 - Determines what is in memory and when
 - Optimizes CPU utilization and response time
 - Sometimes need hardware support
- Activities
 - Keep track of which parts of memory are currently being used and by whom
 - Memory scheduling (move process and data in/out memory)
 - Memory allocation and deallocation

48

48

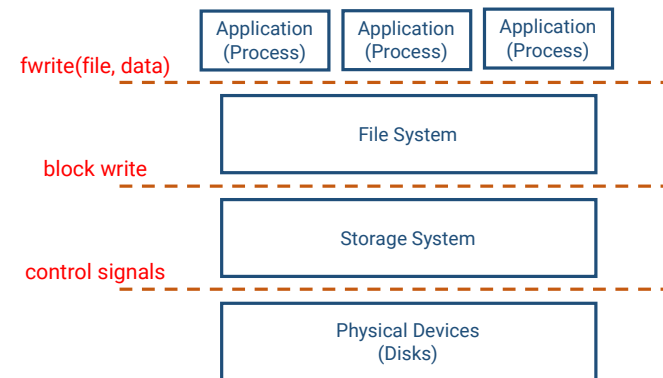
File System Management

- OS provides a uniform logical view of information storage
 - **File: a logical storage unit**
 - Need to be cross-platform
 - Treated as a sequence of bits, bytes, lines, records
 - The format of a file is determined by the application
- File system services provided by OS
 - File creation and deletion
 - Directory creation and deletion
 - Directory hierarchy maintenance
 - File backup

49

49

File System Management (cont.)



50

50

Mass and Tertiary Storage Management

- Usage
 - Store data that does not fit in main memory or data that must be kept for a long time (e.g., disks)
 - Backups of disk data, seldom-used data (e.g., tape drives and tapes, CD/DVD drives and platters)
- Activities
 - Mounting and unmounting
 - Free-space management
 - Storage allocation
 - Disk scheduling (e.g., FCFS)

51

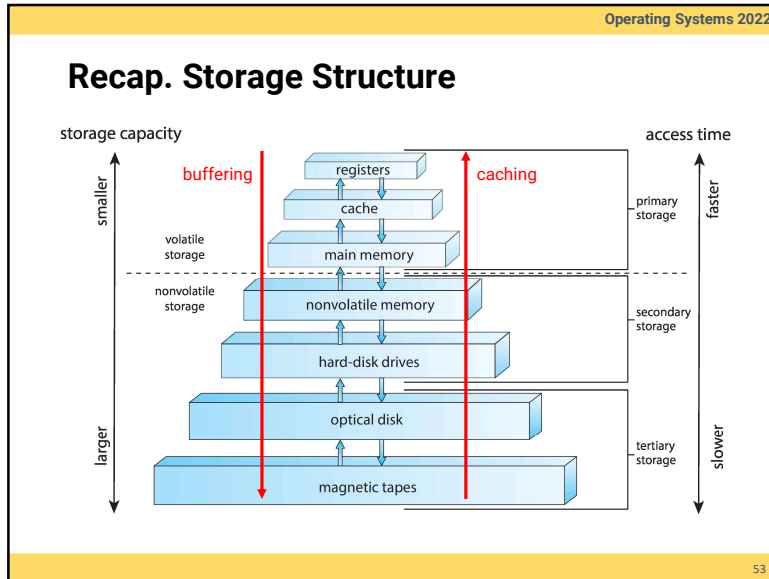
51

I/O System Management

- Hide the peculiarities of specific hardware devices from users
 - Provide a uniform interface for new devices (driver) so they can "plug and play"
- Components
 - **Buffering, caching, and spooling** system
 - A general device-driver interface

52

52

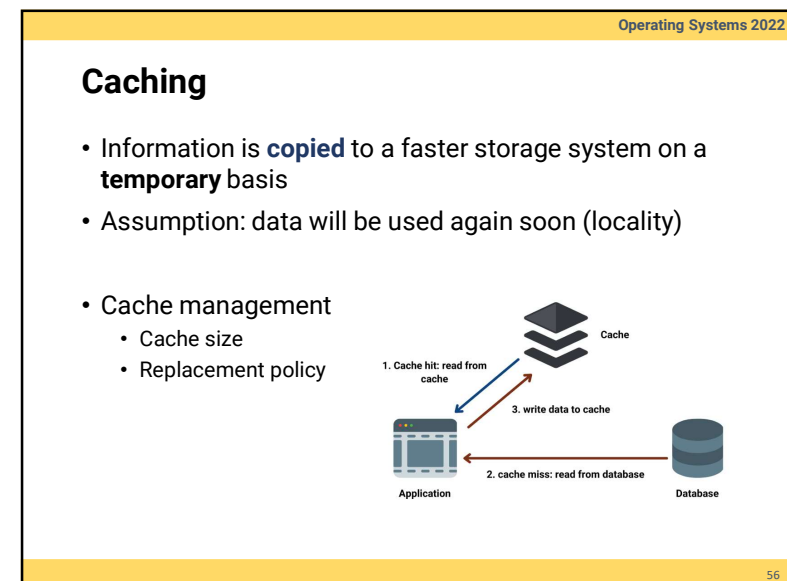
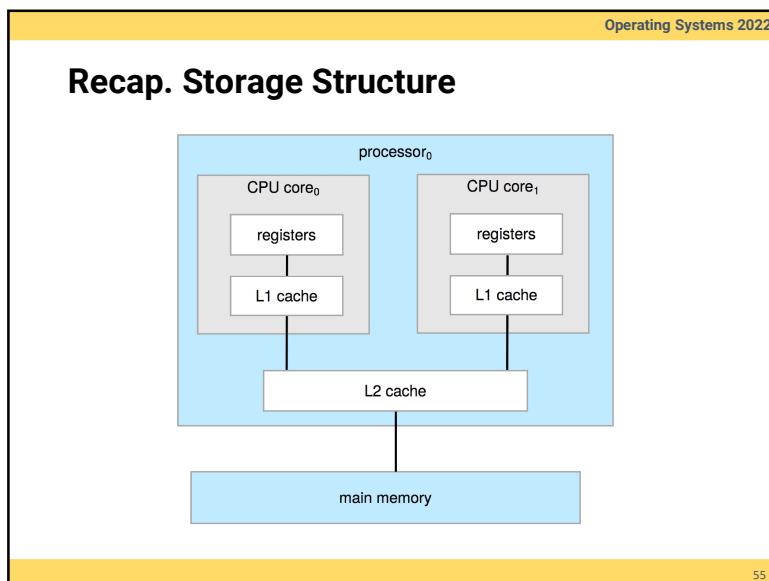


Operating Systems 2022

Recap. Storage Structure


Level	1	2	3	4	5
Name	registers	cache	main memory	solid-state disk	magnetic disk
Typical size	< 1 KB	< 16MB	< 64GB	< 1 TB	< 10 TB
Implementation technology	custom memory with multiple ports CMOS	on-chip or off-chip CMOS SRAM	CMOS SRAM	flash memory	magnetic disk
Access time (ns)	0.25-0.5	0.5-25	80-250	25,000-50,000	5,000,000
Bandwidth (MB/sec)	20,000-100,000	5,000-10,000	1,000-5,000	500	20-150
Managed by	compiler	hardware	operating system	operating system	operating system
Backed by	cache	main memory	disk	disk	disk or tape

54



Caching (cont.)


- Issues: data coherence and consistency among several storage levels
- Uni-tasking: no problem



57

Caching (cont.)

- Issues: data coherence and consistency among several storage levels
- Uni-tasking: no problem
- Multi-tasking: need handle



58

Protection and Security

59

Protection and Security

- Goal**
 - Prevent error and misuse
 - Resources are only allowed to be accessed by authorized processes
- Protection**
 - Any mechanism for controlling the access of processes or users to the resources defined by the computer system
- Security**
 - Defense of a system from external and internal attacks
 - Ex: viruses, denial of service, identity theft

60

Protection

• Dual-mode operations

- **User mode**
 - Executions except those after a trap or an interrupt occurs
- **Monitor mode (system mode, privileged mode)**
 - Can execute all instructions including privileged ones (machine instructions that may cause harm)
- Implemented by a **mode bit** and **system calls**


61

61

Protection (cont.)

• System calls

- Trap to OS for executing privileged instructions
- Protect hardware resources such as I/O devices, memory, and CPU



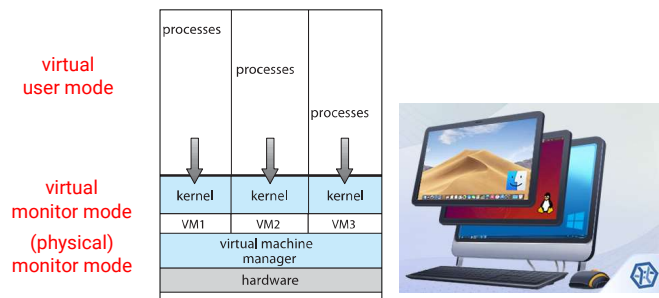
62

62

Protection (cont.)

• Virtual machine (more modes)

- Provide an interface that is identical to the underlying bare hardware



63

63

Protection (cont.)

• I/O protection

- I/O devices are scarce resources, user programs must issue I/O through OS
- Ex: fopen (open), gets (read), puts (write)

• Memory protection

- Prevent a user program from modifying the code or data structures of either the OS or other users
- Ex: instructions to modify the memory space

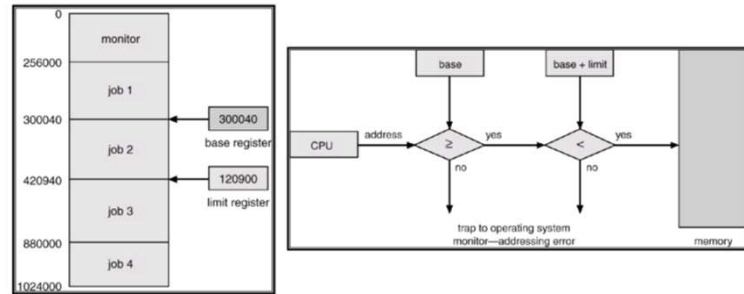
• CPU protection

- Prevent user programs from sucking up CPU power
- Implement by **timers** and time-sharing
 - Need context switch

64

64

Protection (cont.)




65

Kernel Data Structures

66

Kernel Data Structures

- Frequently used data structures
 - Array
 - List (singly, doubly, circular)
 - Stack
 - Queue
 - Tree
 - Hash
 - Bitmap: string of n binary digits representing the status of n items



67

Computer System Environments

68

65

66

67

68

Overview

- Today's computing environments have changed a lot
 - Traditional
 - Mobile
 - Client server
 - Peer-to-peer
 - Cloud computing
 - Real-time embedded
- Trends: network and mobility
- The relatively new mobile and cloud computing bring new issues and concerns on OS design

69

69

Traditional

- Stand-alone general-purpose machines
 - But blurred as most systems interconnect with others (i.e., the Internet)
- Portals provide web access to internal systems
- Mobile computers interconnect via wireless networks
- Networking becoming ubiquitous – even home systems use firewalls to protect home computers from Internet attacks

70

70

Mobile


- Handheld smartphones, tablets, etc.
- Has functional difference with a “traditional” laptop
- Extra feature – more OS features (GPS, gyroscope)
- Allows new types of apps like augmented reality
- Leaders are Apple iOS and Google Android

71

71

Client Server


- Dumb terminals supplanted by smart PCs
- Many systems are now servers, responding to requests generated by clients
 - Compute-server system provides an interface to client to request services (i.e., database)
 - File-server system provides interface for clients to store and retrieve files



72

72

Peer-to-Peer (P2P)



- Another model of distributed system
- P2P does not distinguish clients and servers
 - The role depends on who is requesting or providing a service
- Node must join P2P network
 - Registers its service with central lookup service on network
 - Broadcast request for service and respond to requests for service via discovery protocol
- Examples include Napster and Gnutella, Voice over IP (VoIP) such as Skype

73

73

Cloud Computing

- Deliver computing, storage, and even applications as a service across a network
- Many types
 - **Public cloud** – available via Internet to anyone willing to pay
 - **Private cloud** – run by a company for the company's own use
 - **Hybrid cloud** – includes both public and private cloud components
 - Software as a Service (**SaaS**) – one or more applications available via the Internet (i.e., online photo editor)
 - Platform as a Service (**PaaS**) – software stack ready for application use via the Internet (i.e., a database server)
 - Infrastructure as a Service (**IaaS**) – servers or storage available over Internet (i.e., storage available for backup use)

74

74

Cloud Computing (cont.)

- Why cloud computing becomes so popular?
 - Fast improvement of web technology
 - Portals, network computers ...
 - Network connectivity
 - New categories of devices
 - Embedded computing
 - Car engines, ETC, robots, home automation, AR glasses

75

75

Real-Time Embedded Systems

- Most prevalent form of computers
 - Vary considerable, special purpose, limited purpose OS, real-time OS
 - Use expanding
 - Some have OSes, some perform tasks without an OS
- Real-time OS has well-defined fixed time constraints
 - **Hard real-time system**
 - Processing **must** be done within constraint
 - Correct operation only if constraints met
 - **Soft real-time system**
 - Missing a timing is serious but does not necessary result in failure (ex: multimedia)
- Real-time means on time!

76

76

Free and Open-Source OSes

- Definition: OS with available source
(Otherwise: closed-source OS. E.g., MS Windows, iOS)
- Arguably issues on bugs, security, support
- Examples: GNU/Linux, BSD UNIX ...

77

77

Objectives Review

- Describe the general organization of a computer system and the role of interrupts
- Describe the components in a modern, multiprocessor computer system
- Illustrate the transition from user mode to kernel mode
- Discuss how operating systems are used in various computing environments

78

78