



Processes

Operating Systems

Yu-Ting Wu

(with slides borrowed from Prof. Jerry Chou)

1

Operating Systems 2022

Outline

- Process concept
- Process scheduling
- Operations on processes
- Inter-process communication

2

2

Operating Systems 2022

Process Concept

3

3

Operating Systems 2022

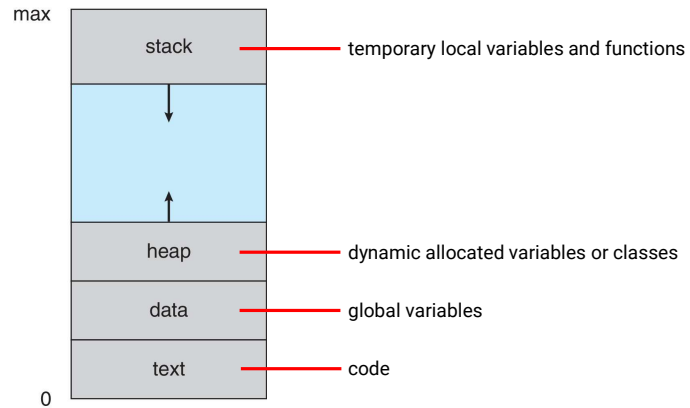
Process Concept

- An operating system concurrently executes a variety of programs
 - **Program: passive entity**, binary file stored **in disk**
 - **Process: active entity**, a running program **in memory**
- A process includes
 - **Code segment** (text section)
 - **Data section**: global variables
 - **Stack**: temporary local variables and functions
 - **Heap**: dynamic allocated variables or classes
 - **Current activity** (e.g., program counter, register contents)
 - **Associated resources** (e.g., handlers of open files)

4

4

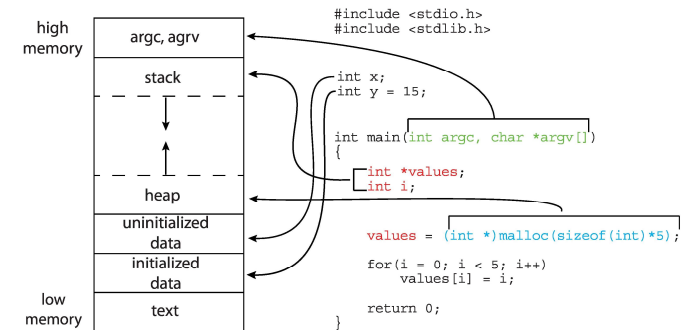
Process in Memory



5

Process in Memory (cont.)

- Example: memory layout of a C program



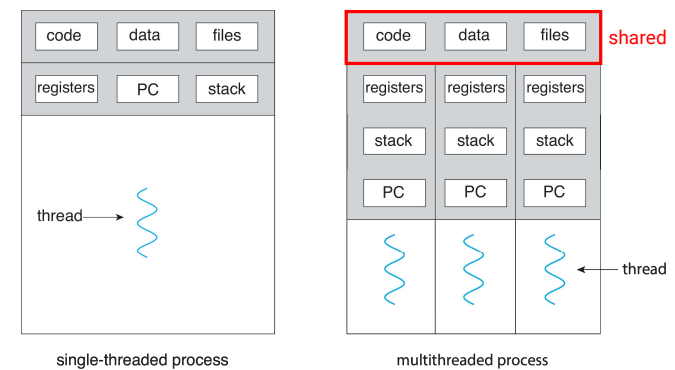
6

Thread

- A thread is a **lightweight process**
 - Basic unit of CPU utilization
- All threads belonging to the same process **share**
 - Code section
 - Data section
 - OS resource (open files, signals)
- But each thread has its own
 - Thread ID
 - Program counter
 - Register set
 - Stack

7

Thread (cont.)



8


Process State

• Types of states

- **New**
 - The process is being created
- **Ready**
 - The process is in the memory waiting to be assigned to a processor
- **Running**
 - The process whose instructions are being executed by CPU
- **Waiting**
 - The process is waiting for events to occur
- **Terminated**
 - The process has finished execution

9

Process State (cont.)



Only one process is running on any processor at any instant
However, many processes may be ready or waiting (put into a queue)

10

Process State (cont.)




11

Process Control Block (PCB)

• Store information of each process

- **Process state**
- **Program counter**
- **CPU register**
- **CPU scheduling information**
 - Priority
- **Memory management information**
 - base/limit register (loaded into registers while the program is going to the running state)
- **I/O state information**
- **Accounting information**




12

PCB (cont.)

• Process representation in Linux

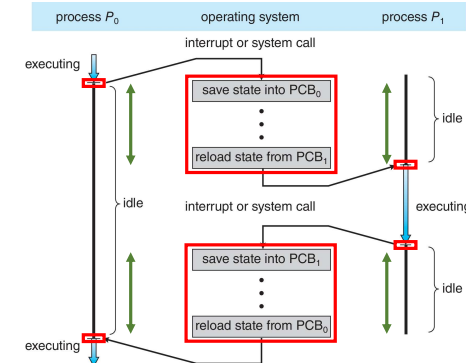
```
pid_t pid; /* process identifier */
long state; /* state of the process */
unsigned int time_slice; /* scheduling information */
struct task_struct *parent; /* this process's parent */
struct list_head children; /* this process's children */
struct files_struct *files; /* list of open files */
struct mm_struct *mm; /* address space of this process */
```



13

Context Switch

- Occurs when the CPU switches from one process to another



14

Context Switch (cont.)

- Context switch:** kernel saves the state of the old process and loads the saved state for the new process
 - The switched context is stored in the PCB
- Context switch time is **purely overhead**
- Switch time (about 1 ~ 1000 ms) depends on
 - Memory speed
 - Number of registers
 - Existence of special instructions
 - Example: a single instruction to save/load all registers
 - Hardware support
 - Example: multiple sets of registers per CPU (multiple contexts loaded at once)

15

Process Scheduling

16

Process Scheduling

- **Multi-programming**
 - CPU runs process at all times to maximize CPU utilization
- **Time sharing**
 - Switch CPU frequently such that users can interact with each program while it is running
- Process will have to wait until the CPU is free and can be re-scheduled

17

17

Process Scheduling Queues


- Maintain **scheduling queues** of processes
 - **Job queue (New state)**
 - Set of all processes in the system
 - **Ready queue (Ready state)**
 - Set of all processes residing in main memory
 - Ready and waiting to execute
 - **Waiting queue (Wait State)**
 - Set of processes waiting for an event (e.g., I/O)
- Processes migrate among the various queues

18

18

Process Scheduling Queues (cont.)

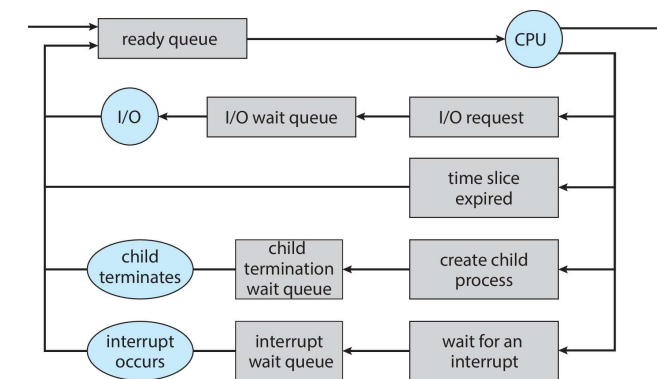
- Ready queue and waiting queue



19

19

Process Scheduling Queues (cont.)




20

20

Process Schedulers

- **Short-term scheduler (CPU scheduler)**
 - Select which process should be executed and allocated CPU
(Ready state → Running state)
- **Long-term scheduler (job scheduler)**
 - Select which processes should be loaded into memory and brought into the ready queue (New state → Ready state)
- **Medium-term scheduler**
 - Select which processes should be swapped in/out memory
(Ready state → Waiting state)



21

21

Long-Term Scheduler

- Control **degree of multi-programming**
- Execute less frequently
 - Invoke only when a process leaves the system or once several minutes
- Strategy
 - Select a **good mix of CPU-bound and I/O bound processes** to increase system overall performance
- New OSes might not contain long-term scheduler
 - The growing memory space
 - Virtual memory (by medium-term scheduler)

22

22

Short-Term Scheduler

- Execute quite frequently
 - Example: once per 100 ms.
- Must be efficient
 - If 10 ms for picking a job, 100 ms for such a pick
→ overhead = $10/110 = 9\%$
- Must ensure fairness

23

23

Medium-Term Scheduler

- **Swap out:**
 - Remove processes from memory (to virtual memory) to reduce the degree of multi-programming
- **Swap in:**
 - Reintroduce swap-out processes into memory
- Purpose: improve **process mix** and **free up memory**

24

24


Operations on Processes

25

25

Tree of Processes

- Each process is identified by a unique **processor identifier (pid)**



26

26

Process Creation

- Resource sharing** (three possibilities)
 - Parent and child processes share **all** resources
 - Child process shares **subset** of parent's resources
 - Parent and child share **no** resources
- Execution order** (two possibilities)
 - Parent and children execute **concurrently**
 - Parent **waits until children terminate**
- Address space** (two possibilities)
 - Children **duplicate** of parent, communication via sharing variables
 - Child **has a program loaded into it**, communication via message passing

27

27

UNIX / Linux Process Creation

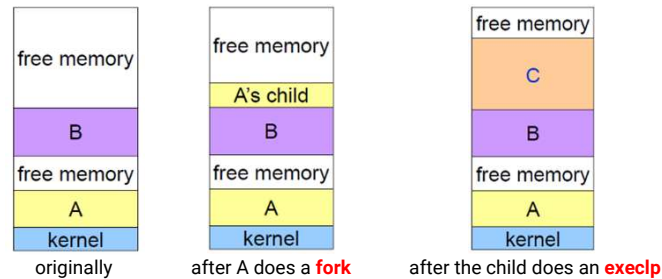
- fork system call**
 - Create a new (child) process
 - The new process **duplicates** the address space of its parent
 - Child and parent **execute concurrently** after fork
 - Child: return value of fork is 0
 - Parent: return value of fork is PID of the child process
- execvp system call**
 - Load a new binary file into memory
 - Destroy the old code
- wait system call**
 - The parent **waits** for one of its child processes to complete

28

28

UNIX / Linux Process Creation (cont.)

- Memory space of fork()
 - Old implementation: A's child is an exact copy of parent
 - Current implementation: use **copy-on-write** technique to store differences in A's child address space



29

29

UNIX / Linux Example

```
#include <sys/types.h>
#include <stdio.h>
#include <unistd.h>

int main()
{
    pid_t pid;

    /* fork a child process */
    pid = fork();

    if (pid < 0) { /* error occurred */
        fprintf(stderr, "Fork Failed");
        return 1;
    }
    else if (pid == 0) { /* child process */
        execvp("/bin/ls", "ls", NULL);
    }
    else { /* parent process */
        /* parent will wait for the child to complete */
        wait(NULL);
        printf("Child Complete");
    }

    return 0;
}
```

printf("Process End!");

Question:
How many times does
"Process End!" show?


30

30

UNIX / Linux Example Quiz

- How many processors are created?

```
#include <stdio.h>
#include <unistd.h>
int main()
{
    for (int i = 0; i < 3; i++)
        fork();
    return 0;
}
```



31

31

Process Termination

- Terminate when the **last statement is executed** or **exit()** is called
 - Return status data from child to parent
 - All resources of the process, including physical and virtual memory, open files, I/O buffers, are deallocated by the OS
- Parent may terminate execution of children processes by specifying its PID (**abort**)
 - Children has exceeded allocated resources
 - Task assigned to child is no longer required
 - The parent is exiting, and the OS does not allow a child to continue if its parent terminates

32

32

Inter-Process Communication

33

33

Inter-Process Communication (IPC)


- **Inter-process communication**
 - A set of methods for the exchange of data among multiple threads in one or more processes
- **Independent process**
 - Cannot affect or be affected by other processes
- **Cooperating process**
 - Otherwise
- **Purposes**
 - Information sharing
 - Computation speedup
 - Convenience (perform several tasks at one time)
 - Modularity

34

34

Communication Methods

- **Shared memory**
 - Require more careful **user synchronization**
 - Implemented by memory access (faster)
 - Use memory address to access data
- **Message passing**
 - No conflict: more efficient for small data
 - Use send/rcv message
 - Implement by system call (slower)




35

35

Message Passing Methods

- **Sockets**
 - A network connection identified by **IP** and **port**
 - Exchange **unstructured stream of bytes**
- **Remote Procedure Calls**
 - Cause a procedure to execute in another address space
 - **Parameters** and **return values** are passed by messages



36

36

Shared Memory

- Processes are responsible for
 - Establishing a region of shared memory** (ask OS for help)
 - Typically, the created shared-memory regions resides in the address of the process creating the shared memory segment
 - Participating processes must agree to remove memory access constraint from OS
 - Determining the form of the data and the location**
 - Synchronization:** ensuring data are not written simultaneously by processes

37


37

Consumer and Producer

- Producer** process produces information that is consumed by a **Consumer** process

- Buffer as a circular array with size B

- Next free: *in*
- First available: *out*
- Empty: $in = out$
- Full: $(in + 1) \% B = out$



- The solution allows at most $(B - 1)$ item in the buffer
 - Otherwise, cannot tell the buffer is empty or full

38

38

Consumer and Producer (cont.)

- Producer** process

```

item next_produced;

while (true) {
    /* produce an item in next produced */
    while (((in + 1) % BUFFER_SIZE) == out)
        ; /* do nothing */
    buffer[in] = next_produced;
    in = (in + 1) % BUFFER_SIZE;
}

```

39

39

Consumer and Producer (cont.)

- Consumer** process

```

item next_consumed;

while (true) {
    while (in == out)
        ; /* do nothing */
    next_consumed = buffer[out];
    out = (out + 1) % BUFFER_SIZE;

    /* consume the item in next consumed */
}

```

40

40

Consumer and Producer (cont.)

- Another solution for filling all the buffer
- Use an additional variable, **counter**, for keeping track of the number of items in the buffer
- Initially, counter is set to zero
- Counter is increased by one by the producer after it produces a new item
- Counter is decreased by one by the consumer after it consumes an item

41

Consumer and Producer (cont.)

- **Producer** process (new version)

```
while (true) {
    /* produce an item in next produced */

    while (counter == BUFFER_SIZE)
        ; /* do nothing */

    buffer[in] = next_produced;
    in = (in + 1) % BUFFER_SIZE;
    counter++;
}
```

42

Consumer and Producer (cont.)

- **Consumer** process (new version)

```
while (true) {
    while (counter == 0)
        ; /* do nothing */

    next_consumed = buffer[out];
    out = (out + 1) % BUFFER_SIZE;
    counter--;
    /* consume the item in next consumed */
}
```

43

Consumer and Producer (cont.)

- **Race condition**

- Counter++ can be implemented as


```
register1 = counter
register1 = register1 + 1
counter = register1
```
- Counter-- can be implemented as


```
register2 = counter
register2 = register2 - 1
counter = register2
```
- Example (initially counter = 5):

| | | |
|----------------------|---------------------------|-----------------|
| S0: producer execute | register1 = counter | {register1 = 5} |
| S1: producer execute | register1 = register1 + 1 | {register1 = 6} |
| S2: consumer execute | register2 = counter | {register2 = 5} |
| S3: consumer execute | register2 = register2 - 1 | {register2 = 4} |
| S4: producer execute | counter = register1 | {counter = 6} |
| S5: consumer execute | counter = register2 | {counter = 4} |
- Let's discussed this problem again in Chapter 6

44

Message Passing System

- Mechanism for processes to **communicate** and **synchronize** their actions
- IPC provides two operations
 - **Send** (message)
 - **Receive** (message)
- To communicate, processes need to
 - Establish a **communication link**
 - Exchange a message via send/receive

45

45

Message Passing System (cont.)

- Implementation of communication link
 - **Physical**
 - HW bus
 - Network
 - **Logical (properties of the link)**
 - **Direct or indirect communication**
 - Symmetric or asymmetric communication
 - **Blocking or non-blocking**
 - Automatic or explicit buffering
 - Send by copy or send by reference
 - Fixed-sized or variable-sized messages

46

46

Direct Communication

- Processes must **name each other explicitly**
 - **Send (P, message)**: send a message to process P
 - **Receive (Q, message)**: receive a message from process Q
- Properties of communication link
 - Links are **established automatically**
 - **One-to-one** relationship between links and processes
 - The link may be unidirectional, but is usually **bidirectional**
- **Limited modularity**: if the name of a process is changed, all old names should be found

47

47

Indirect Communication

- Messages are directed and received from **mailboxes** (also referred as **ports**)
 - Each mailbox has a unique ID
 - Processes can communicate if they share a mailbox
 - **Send (A, message)**: send a message to mailbox A
 - **Receive (A, message)**: receive a message from mailbox A
- Properties of communication link
 - Link established only **if processes share a common mailbox**
 - **Many-to-many** relationship between links and processes
 - Link may be unidirectional or bi-directional
 - Mailbox can be owned either by OS or processes

48

48

Indirect Communication (cont.)

- Mailbox sharing
 - P1, P2, and P3 share mailbox A
 - P1 sends, P2 and P3 receives
 - Who gets the message?
- Solutions
 - Allow a link to be associated with at most two processes
 - Allow only one process at a time to execute a receive operation (by locking and delay)
 - Allow the system to select arbitrary the receiver (sender is notified who the receiver was)

49

49

Synchronization


- Messages passing may be either **blocking** or **non-blocking**
- **Blocking (synchronous)**
 - **Blocking send**: sender is blocked until the message is received by receiver or by the mailbox
 - **Blocking receive**: receiver is blocking until the message is available
- **Non-blocking (asynchronous)**
 - **Non-blocking send**: sender sends the message and resumes operation
 - **Non-blocking receive**: receiver receives a valid message or a null

50

50

Synchronization (cont.)

- Buffer implementation for queue of messages attached to a link
 - **Zero capacity**
 - No messages are queued on a link
 - Sender must wait for receiver
 - **Bounded capacity**
 - Finite length of n messages
 - Sender must wait if the link is full
 - **Unbounded capacity**
 - Infinite length
 - Sender never waits

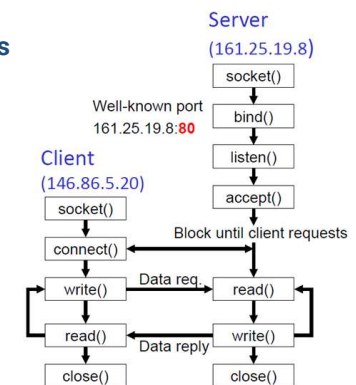


51

51

Sockets

- A socket is identified by a concatenation of **IP address** and a **port number**
- Communication consists between a pair of sockets
- Use **127.0.0.1** to refer itself




52

52

Socket (cont.)


- Consider as a low-level form of communication **unstructured stream of bytes** to be exchanged
- Data parsing responsibility falls upon the server and the client applications



53

Remote Procedure Calls (RPC)

- Remote procedure call abstracts procedure calls between processes on networked systems
 - Allow programs to call procedures located on other machines (and other processes)
- Stub/Skeleton**: client-side/server-side proxy for the actual procedure on the server



54

Remote Procedure Calls (cont.)

• Client stub

- Pack parameters into a message (parameter marshaling)
- Call OS to send directly to the server
- Wait for the results returned from the server



• Server skeleton

- Receive a call from a client
- Unpack the parameters
- Call the responding procedure
- Return results to the caller

55

Objectives Review

- Identify the separate components of a process and illustrate how they are represented and scheduled in an operating system
- Describe how processes are created and terminated in an operating system, including developing programs using the appropriate system calls that perform these operations
- Describe and contrast inter-process communication using shared memory and message passing

56