



Synchronization (I)

Operating Systems

Yu-Ting Wu

(with slides borrowed from Prof. Jerry Chou)

Outline

- Background
- Critical section
- Synchronization hardware
- Semaphores

Background

Background

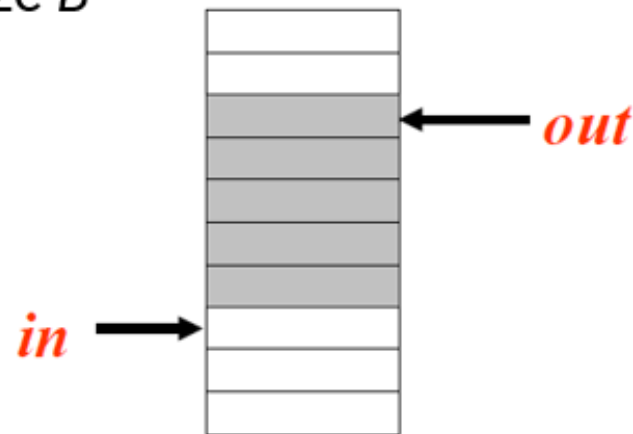
- **Concurrent access** to **shared data** may result in **data inconsistency**
- Maintaining data consistency requires mechanism to ensure the **orderly execution** of cooperating processes

Consumer & Producer Problem

- Determine whether buffer is empty or full
 - Use *in*, *out* position

- Buffer as a circular array with size B

- Next free: *in*
- First available: *out*
- Empty: $in = out$
- Full: $(in + 1) \% B = out$



- The solution allows at most $(B - 1)$ item in the buffer
 - Otherwise, cannot tell the buffer is empty or full

Consumer & Producer Problem (cont.)

- Determine whether buffer is empty or full
 - Use **count** value

```
/* Producer */  
while (true) {  
    // produce an item in next produced.  
    while (counter == BUFFER_SIZE);  
    // do nothing.  
    buffer[in] = next_produced;  
    in = (in + 1) % BUFFER_SIZE;  
    counter++;  
}
```

```
/* Consumer */  
while (true) {  
    while (counter == 0);  
    // do nothing.  
    next_consumed = buffer[out];  
    out = (out + 1) % BUFFER_SIZE;  
    counter--;  
    // consume the item in next consumed.  
}
```

Concurrent Operations on Counter

- The statement “counter++” may be implemented in machine language as

move ax, counter

add ax, 1

move counter, ax

- The statement “counter--” may be implemented as

move bx, counter

sub bx, 1

move counter, bx

Instruction Interleaving

- Assume counter is initially 5. One interleaving of statement is

producer: move ax, counter

→ ax = 5

producer: add ax, 1

→ ax = 6

context switch

consumer: move bx, counter

→ bx = 5

consumer: sub bx, 1

→ bx = 4

context switch

producer: move counter, ax

→ counter = 6

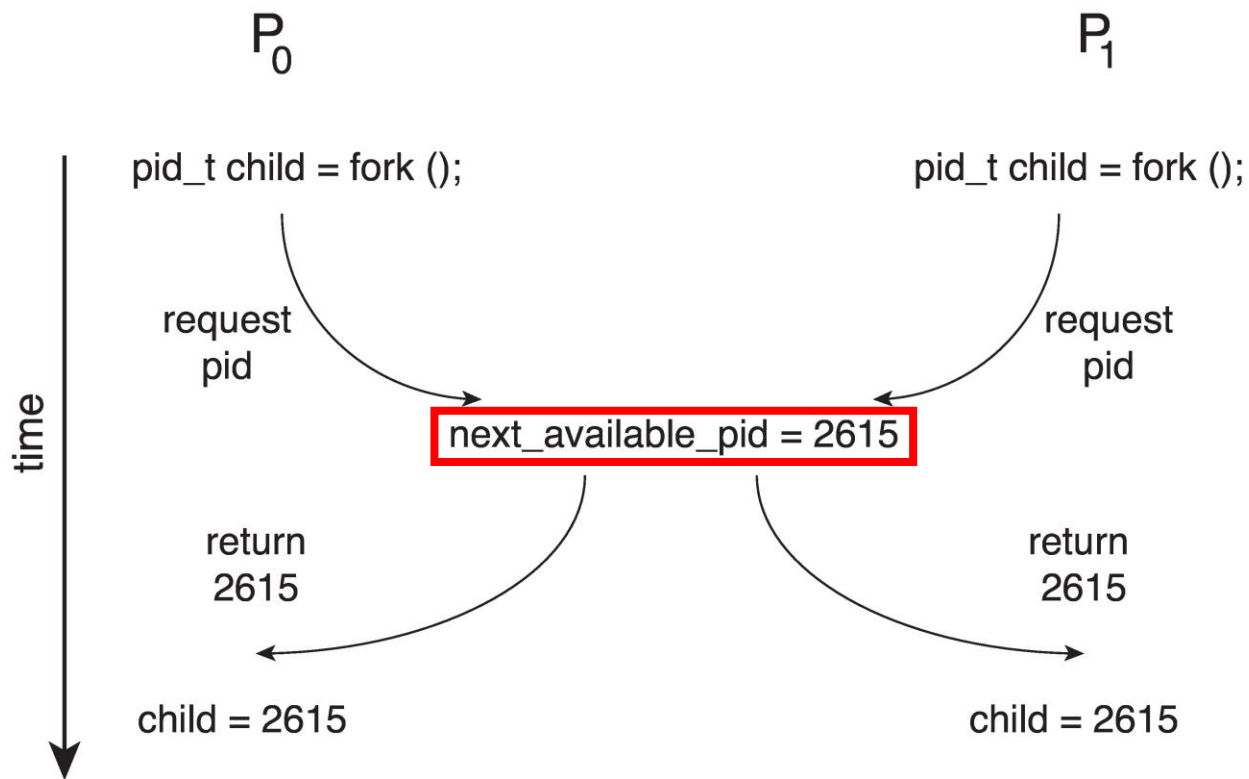
context switch

consumer: move counter, bx

→ counter = 4

Another Example

- An example in the kernel



Race Condition

- The situation where several processes access and manipulate **shared** data concurrently.
- **The final value of the shared data depends upon which process finishes last**
- To prevent race condition, concurrent processes must be **synchronized**
 - On a single-processor machine, we could disable interrupt or use non-preemptive CPU scheduling
 - But how about on **multi-processor** machines and **preemptive** scheduling?
- We need a mechanism to solve the synchronization issue, commonly described as **critical section problem**

Critical Section

The Critical-Section Problem

- **Purpose**

- A **protocol** for processes to cooperate

- **Problem description**

- N processes are competing to use some **shared** data
 - Each process has a **code segment**, called **critical section**, in which the shared data is accessed
 - Ensure that when one process is executing in its critical section, **no other process is allowed** to execute in its critical section

➔ **mutually exclusive !**

The Critical-Section Problem (cont.)

- **General code section structure**

- Only one process can be in a critical section

do {

entry section



get **entry permission**

critical section



modified **shared data**

exit section



release **entry permission**

remainder section

} while (1);

Critical-Section Requirements

- **Mutual exclusion**

- If a process P is executing in its critical section (CS), no other processes can be executing in their CS

- **Progress**

- If no process is executing in its CS and there exist some processes that wish to enter their CS, these processes cannot be postponed indefinitely

- **Bounded Waiting**

- A **bound** must exist on the number of times that other processes are allowed to enter their CS after a process has made a request to enter its CS

- How to design **entry** and **exit section** to satisfy the above requirement?

CS Solutions and Synchronization Tools

- **Software solution**
- **Synchronization hardware**
- **Semaphore**
- **Monitor**

Algorithm for Two Processes

- Only 2 processes P_0 and P_1
- Shared variables
 - `int turn;` // initially `turn = 0`
 - `turn == i` \rightarrow P_i can enter its critical section

/ Process 0 */*

do {

`while (turn != 0);`

critical section

`turn = 1;`

remainder section

} while (1);

/ Process 1 */*

do {

`while (turn != 1);`

critical section

`turn = 0;`

remainder section

} while (1);

mutual exclusion? **Y** progress? **N** bounded-wait? **Y**

Peterson's Solution for Two Processes

- Shared variables
 - `int turn;` // initially `turn = 0`
 - `turn == i` \rightarrow P_i can enter its critical section
 - `boolean flag[2];` // initially `flag[0] = flag[1] = false`
 - `flag[i] == true` \rightarrow P_i is ready to enter its critical section

/ Process i */*

do {

```
flag[i] = true;
turn = j;
while (flag[j] && turn == j);
```

critical section

```
flag[i] = false;
```

remainder section

} while (1);

— **entry section**

— **exit section**

Proof of Peterson's Solution

• Mutual exclusion

- If P_0 in CS $\rightarrow flag[1] == false \parallel turn == 0$
- If P_1 in CS $\rightarrow flag[0] == false \parallel turn == 1$
- Assume both processes in CS $\rightarrow flag[0] == flag[1] == true$
 - $\rightarrow turn == 0$ for P_0 to enter, $turn == 1$ for P_1 to enter
 - $\rightarrow turn$ will be either 0 or 1, so P_0, P_1 cannot in CS at the same time

/ Process 0 */*

do {

```
flag[0] = true;
turn = 1;
while (flag[1] && turn == 1);
```

 *critical section*

```
flag[0] = false;
```

remainder section

} while (1);

/ Process 1 */*

do {

```
flag[1] = true;
turn = 0;
while (flag[0] && turn == 0);
```

 *critical section*

```
flag[1] = false;
```

remainder section

} while (1);

Proof of Peterson's Solution

• Progress (e.g., P_0 wishes to enter its CS)

- (1) If P_1 is not ready $\rightarrow flag[1] = false \rightarrow P_0$ can enter
- (2) If both are ready $\rightarrow flag[0] == flag[1] == true$
 \rightarrow If $turn == 0$ then P_0 enters, otherwise P_1 enters
- Either cases, some waiting process can enter CS

```

/* Process 0 */
do {
    flag[0] = true;
    turn = 1;
    → while (flag[1] && turn == 1);
        critical section
    flag[0] = false;
        remainder section
} while (1);
  
```

```

/* Process 1 */
do {
    flag[1] = true;
    (2) turn = 0;
    → while (flag[0] && turn == 0);
        critical section
    (1) flag[1] = false;
        remainder section
    → } while (1);
  
```

Proof of Peterson's Solution

- **Bounded waiting (e.g., P_0 wishes to enter its CS)**

- (1) Once P_1 exits CS $\rightarrow flag[1] == false \rightarrow P_0$ can enter
- (2) If P_1 exits CS and reset $flag[1] = true$
 $\rightarrow turn == 0$ (overwrite P_0 setting) $\rightarrow P_0$ can enter
- P_0 won't wait infinitely

```

/* Process 0 */
do {
    flag[0] = true;
    turn = 1;
    → while (flag[1] && turn == 1);
        critical section
    flag[0] = false;
        remainder section
} while (1);
  
```

```

/* Process 1 */
do {
    (2) flag[1] = true;
    → turn = 0;
        while (flag[0] && turn == 0);
            critical section
    (1) flag[1] = false;
        remainder section
} while (1);
  
```

Peterson's Solution and Modern Architecture

- Peterson's solution is **not guaranteed** to work on modern architectures
 - To improve performance, processors and/or compilers may **reorder operations** that have no **dependencies**
- For **single-threaded process** this is **OK** as the result will always be the same
- For **multi-threaded process** the reordering may produce inconsistent or unexpected results

Peterson's Solution and Modern Architecture (cont.)

- Example:
 - Two threads share the data:

```
bool flag = true;  
int x = 0;
```

- Thread1 performs

```
while (!flag);  
print x;
```

- Thread2 performs

```
x = 100;  
flag = true;
```

- Expected output will be 100

Peterson's Solution and Modern Architecture (cont.)

- Example (cont.):
 - Because the variables *flag* and *x* are independent of each other, the instructions:

```
x = 100;  
flag = true;
```

for Thread2 may be reordered

- If this occurs, the output may be 0!

Peterson's Solution and Modern Architecture (cont.)

```
/* Process 0 */
```

```
do {
```

```
    flag[0] = true;
    turn = 1;
    while (flag[1] && turn == 1);
```

critical section

```
    flag[0] = false;
```

remainder section

```
} while (1);
```

```
/* Process 1 */
```

```
do {
```

```
    flag[1] = true;
    turn = 0;
    while (flag[0] && turn == 0);
```

critical section

```
    flag[1] = false;
```

remainder section

```
} while (1);
```

The variables *flag[]* and *turn* are independent, so they might be reordered

Peterson's Solution and Modern Architecture (cont.)

/ Process 0 */*

do {

```
turn = 1;
flag[0] = true;
while (flag[1] && turn == 1);
```

critical section

```
flag[0] = false;
```

remainder section

} while (1);

/ Process 1 */*

do {

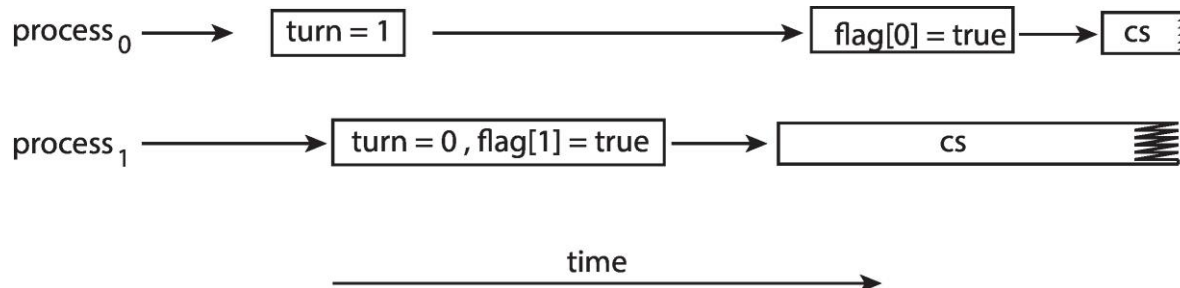
```
turn = 0;
flag[1] = true;
while (flag[0] && turn == 0);
```

critical section

```
flag[1] = false;
```

remainder section

} while (1);



Both processes will enter their CS !

We can use **Memory Barrier** to ensure the correctness

Memory Barrier

- When a memory barrier instruction is performed, the system ensures that **all loads and stores are completed before any subsequent loads or stores operations are performed**
- Recall previous example:

```
/* Thread 1 */  
while (! flag) ;   load  
print x
```

```
/* Thread 2 */  
x = 100;          store  
flag = true;
```

Memory Barrier (cont.)

- When a memory barrier instruction is performed, the system ensures that **all loads and stores are completed before any subsequent loads or stores operations are performed**
- Modification:

```
/* Thread 1 */  
while (! flag) ;  
memory_barrier();  
print x
```

```
/* Thread 2 */  
x = 100;  
memory_barrier();  
flag = true;
```

For Thread 1, we are guaranteed that the value of *flag* is loaded before the value of *x*

Memory Barrier (cont.)

- When a memory barrier instruction is performed, the system ensures that **all loads and stores are completed before any subsequent load or store operations are performed**
- Modification:

```
/* Thread 1 */  
while (! flag) ;  
memory_barrier();  
print x
```

```
/* Thread 2 */  
x = 100;  
memory_barrier();  
flag = true;
```

For Thread 2, we are guaranteed that the assignment to *x* occurs before the assignment to *flag*

Producer & Consumer Problem

```
/* Producer process 0 */
```

```
while (true) {
```

entry section

```
nextItem = getItem();  
while (counter == BUFFER_SIZE);  
buffer[in] = nextItem;  
in = (in + 1) % BUFFER_SIZE;  
counter++;  
computing();
```

exit section

```
}
```

```
/* Consumer process 0 */
```

```
while (true) {
```

entry section

```
while (counter == 0);  
item = buffer[out];  
out = (out + 1) % BUFFER_SIZE;  
counter--;  
computing();
```

exit section

```
}
```

Incorrect. Deadlock if consumer enters the CS first

Producer & Consumer Problem (cont.)

```
/* Producer process 0 */
```

```
while (true) {  
    nextItem = getItem();  
    while (counter == BUFFER_SIZE);  
    buffer[in] = nextItem;  
    in = (in + 1) % BUFFER_SIZE;
```

entry section

```
    counter++;  
    computing();
```

exit section

```
}
```

```
/* Consumer process 0 */
```

```
while (true) {  
    while (counter == 0);  
    item = buffer[out];  
    out = (out + 1) % BUFFER_SIZE;
```

entry section

```
    counter--;  
    computing();
```

exit section

```
}
```

Correct but **poor performance**

Producer & Consumer Problem (cont.)

/ **Producer** process 0 */*

```
while (true) {
    nextItem = getItem();
    while (counter == BUFFER_SIZE);
    buffer[in] = nextItem;
    in = (in + 1) % BUFFER_SIZE;
```

entry section

```
counter++;
```

exit section

```
computing();
```

```
}
```

/ **Consumer** process 0 */*

```
while (true) {
    while (counter == 0);
    item = buffer[out];
    out = (out + 1) % BUFFER_SIZE;
```

entry section

```
counter--;
```

exit section

```
computing();
```

```
}
```

Correct and **maximize concurrent performance**

Bakery Algorithm (n processes)

- Before entering its CS, each process receives a **number (#)**
- **Holder of the smallest # enters CS**
- The numbering scheme always generates # in **non-decreasing order**; i.e., 1, 2, 3, 3, 4, 5, 5, 5 ...
- If processes P_i and P_j receive the same #, if $i < j$, then P_i is served first
- Notation:
 - $(a, b) < (c, d)$ if
 - $a < c$ or
 - $a == c \ \&\& \ b < d$

Bakery Algorithm (n processes) (cont.)

```

// Process i:
do {
    get ticket      → choosing[i] = true;
                    num[i] = max (num[0], num[1], ...num[n-1]) + 1;
                    choosing[i] = false;
    FCFS           → for (j = 0; j < n; j++) {
                    while (choosing[j]);
                    while ((num[j] != 0) && ((num[j], j) < (num[i], i)));
                    }
    release ticket → critical section
                    num[i] = 0;
                    remainder section
}

```

Bounded waiting because processes enter CS on a **first come, first served** basis

Bakery Algorithm (n processes) (cont.)

- Why cannot compare when num is being modified?
- Without locking
 - Let 5 be the current maximum number
 - If P_1 and P_4 take number together, but P_4 finishes before P_1
 - $P_1 = 0, P_4 = 6 \rightarrow P_4$ will enter the CS
 - After P_1 takes the number
 - $P_1 = P_4 = 6 \rightarrow P_1$ will enter the CS as well !
- With locking
 - P_4 will have to wait until P_1 finish taking the number
 - Both P_1 & P_4 will have the new number “6” before comparison

Pthread Lock/Mutex Routines

- To use mutex, it must be declared as of type `pthread_mutex_t` and initialized with `pthread_mutex_init()`
- A mutex is destroyed with `pthread_mutex_destroy()`
- A critical section can then be protected using `pthread_mutex_lock()` and `pthread_mutex_unlock()`

```
#include "pthread.h"
pthread_mutex_t mutex;
pthread_mutex_init(&mutex, NULL);
pthread_mutex_lock(&mutex);
critical section
pthread_mutex_unlock(&mutex);
pthread_mutex_destroy(&mutex);
```

← enter critical section

← exit critical section

Condition Variables

- **Condition variables (CV)** represent some condition that a **thread** can
 - **Wait on**, until the condition occurs; or
 - **Notify** other waiting threads that the condition has occurred
- Three operations on condition variables
 - **wait()** – block until another thread calls **signal()** or **broadcast()** on the CV
pthread_cond_wait(&theCV, &someLock)
 - **signal()** – wake up one thread waiting on the CV
pthread_cond_signal(&theCV)
 - **broadcast()** – wake up all threads waiting on the CV
pthread_cond_broadcast(&theCV)

Condition Variables (cont.)

- Example

- A thread is designed to take action when $x == 0$
- Another thread is responsible for decrementing the counter

```
pthread_cond_t  cond;
pthread_cond_init(&cond, NULL);
pthread_mutex_t mutex;
pthread_mutex_init(&mutex, NULL);
```

```
action() {
    pthread_mutex_lock(&mutex);
    if (x != 0)
        pthread_cond_wait(&cond, &mutex);
    pthread_mutex_unlock(&mutex);
    take_action();
}
```

```
counter() {
    pthread_mutex_lock(&mutex);
    x--;
    if (x == 0)
        pthread_cond_signal(&cond);
    pthread_mutex_unlock(&mutex);
}
```

Using Condition Variables

```
action() {  
→ pthread_mutex_lock(&mutex);  
  if (x != 0)  
    pthread_cond_wait(cond, mutex);  
  pthread_mutex_unlock(&mutex);  
  take_action();  
}
```

```
→ counter() {  
  pthread_mutex_lock(&mutex);  
  x--;  
  if (x == 0)  
    pthread_cond_signal(cond);  
  pthread_mutex_unlock(&mutex);  
}
```

Lock mutex

Using Condition Variables

```

action() {
    pthread_mutex_lock(&mutex);
    if (x != 0)
        pthread_cond_wait(cond, mutex);
    pthread_mutex_unlock(&mutex);
    take_action();
}

```

```

counter() {
    pthread_mutex_lock(&mutex);
    x--;
    if (x == 0)
        pthread_cond_signal(cond);
    pthread_mutex_unlock(&mutex);
}

```

Lock mutex

Wait()

- Put the thread into **sleep** and **releases the lock**

Lock mutex

Using Condition Variables

```

action() {
    pthread_mutex_lock(&mutex);
    if (x != 0)
        pthread_cond_wait(cond, mutex);
    pthread_mutex_unlock(&mutex);
    take_action();
}

```

```

counter() {
    pthread_mutex_lock(&mutex);
    x--;
    if (x == 0)
        pthread_cond_signal(cond);
    pthread_mutex_unlock(&mutex);
}

```

Lock mutex

Wait()

- Put the thread into **sleep** and **releases the lock**
- **Waked up**, but the thread is locked

Lock mutex

Signal()

Using Condition Variables

```

action() {
    pthread_mutex_lock(&mutex);
    if (x != 0)
        pthread_cond_wait(cond, mutex);
    pthread_mutex_unlock(&mutex);
    take_action();
}

```

```

counter() {
    pthread_mutex_lock(&mutex);
    x--;
    if (x == 0)
        pthread_cond_signal(cond);
    pthread_mutex_unlock(&mutex);
}

```

Lock mutex

Wait()

- Put the thread into **sleep** and **releases the lock**
- **Waked up**, but the thread is locked
- **Re-acquire lock** and resume execution

Lock mutex

Signal()

Release the lock

Using Condition Variables

```

action() {
    pthread_mutex_lock(&mutex);
    if (x != 0)
        pthread_cond_wait(cond, mutex);
    pthread_mutex_unlock(&mutex);
    take_action();
}

```

```

counter() {
    pthread_mutex_lock(&mutex);
    x--;
    if (x == 0)
        pthread_cond_signal(cond);
    pthread_mutex_unlock(&mutex);
}

```

Lock mutex

Wait()

- Put the thread into **sleep** and **releases the lock**
- **Waked up**, but the thread is locked
- **Re-acquire lock** and resume execution

Release the lock

Lock mutex

Signal()

Release the lock

ThreadPool Implementation

- Task structure

```
typedef struct {
    void (*function)(void *);
    void *argument;
} threadpool_task_t;
```

- ThreadPool structure

```
struct threadpool_t {
    pthread_mutex_t lock;
    pthread_cond_t notify;
    pthread_t *threads;
    threadpool_task_t *queue;
    int thread_count;
    int queue_size;
    int head;
    int tail;
    int count;
    int shutdown;
    int started;
};
```

- Allocate thread and task queue

```
/* Allocate thread and task queue */
pool->threads = (pthread_t *) malloc(sizeof(pthread_t) * thread_count);
pool->queue = (threadpool_task_t *) malloc(sizeof(threadpool_task_t) * queue_size);
```

ThreadPool Implementation (cont.)

```
static void *threadpool_thread(void *threadpool)
{
    threadpool_t *pool = (threadpool_t *)threadpool;
    threadpool_task_t task;

    for(;;) {
        /* Lock must be taken to wait on conditional variable */
        pthread_mutex_lock(&(pool->lock));

        /* Wait on condition variable, check for spurious wakeups.
           When returning from pthread_cond_wait(), we own the lock. */
        while((pool->count == 0) && (!pool->shutdown)) {
            pthread_cond_wait(&(pool->notify), &(pool->lock));
        }
    }
}
```

ThreadPool Implementation (cont.)

```
/* Grab our task */  
task.function = pool->queue[pool->head].function;  
task.argument = pool->queue[pool->head].argument;  
pool->head += 1;  
pool->head = (pool->head == pool->queue_size) ? 0 : pool->head;  
pool->count -= 1;  
  
/* Unlock */  
pthread_mutex_unlock(&(pool->lock));  
  
/* Get to work */  
(*task.function)(task.argument);  
}
```

Synchronization Hardware

Hardware Support

- The CS problem occurs because the modification of a shared variable may be **interrupted**
- If disable interrupts when in CS
 - Not feasible in multiprocessor machine
 - Clock interrupts cannot fire in any machine
- **HW support solution: atomic instructions**
 - atomic: as one **uninterruptible** unit
 - Example: **TestAndSet(var)** and **Swap(a, b)**

Atomic TestAndSet()

```
bool TestAndSet (bool &lock) {
    bool value = lock;
    lock = true;
    return value;
}
```

execute atomically:
return the value of “lock” and
set “lock” to true

```
shared data: bool lock; // initially lock = false
// P0
do {
    while (TestAndSet (lock));
    critical section
    lock = false;
    remainder section
} while (1);
```

```
// P1
do {
    while (TestAndSet (lock));
    critical section
    lock = false;
    remainder section
} while (1);
```

mutual exclusion? **Y** progress? **Y** bounded-wait? **N**

Atomic Swap()

Enter CS if lock == false

shared data: bool lock; // initially lock = false

// P₀

do {

```
key0 = true;
while (key0 == true)
    Swap(lock, key0);
```

critical section

```
lock = false;
```

remainder section

} while (1);

// P₁

do {

```
key1 = true;
while (key1 == true)
    Swap(lock, key1);
```

critical section

```
lock = false;
```

remainder section

} while (1);

mutual exclusion? **Y** progress? **Y** bounded-wait? **N**

Atomic CompareAndSwap()

```
int CompareAndSwap (int &value, int expected, int new_value) {
    int temp = value;
    if (value == expected)
        value = new_value;
    return temp;
}
```

shared data: int lock; // initially lock = 0

// P_0

do {

while (CompareAndSwap (lock, 0, 1) != 0);

critical section

lock = 0;

remainder section

} while (1);

mutual exclusion? **Y** progress? **Y** bounded-wait? **N**

Atomic Variables

- **Atomic variable** is another tool that provides **atomic (uninterruptible)** updates on basic data types such as integers and Booleans
- Usually built with atomic instructions such as ***CompareAndSwap***
- Example:
 - Let ***sequence*** be an atomic variable
 - Let ***increment()*** be an operation for incrementing the atomic variable ***sequence***
 - The command ***increment(&sequence)*** ensures ***sequence*** is incremented without interruption

Atomic Variables (cont.)

- The *increment()* function can be implemented as follows

```

void increment (atomic_int &v) {
    int temp;
    do {
        temp = v;
    }
    while (temp != (CompareAndSwap (v, temp, temp+1)));
}

```

```

int CompareAndSwap (int &value, int expected, int new_value) {
    int temp = value;
    if (value == expected)
        value = new_value;
    return temp;
}

```

Atomic Variables (cont.)

- The *increment()* function can be implemented as follows

```

        5
    bool increment (atomic_int &v) {
        int temp;
        do {
            temp = v;

```

v is modified



```

    }
    5
    while (temp != (CompareAndSwap (v, temp, temp+1)));
    3 5 6
}

```

```

int CompareAndSwap (int &value, int expected, int new_value) {
    int temp = value;
    if (value == expected)
        value = new_value;
    return temp;
}
    3      5      6

```

Atomic Variables (cont.)

- The *increment()* function can be implemented as follows

```
bool increment (atomic_int &v) {
    int temp;
    do {
        temp = v;
```

v is modified



```
    } while (temp != (CompareAndSwap (v, temp, temp+1)));
}
```

```
int CompareAndSwap (int &value, int expected, int new_value) {
    int temp = value;
    if (value == expected)
        value = new_value;
    return temp;
}
```

Semaphores

Semaphores

- A tool to generalize the synchronization problem
- More specifically
 - A **record** of **how many units of a particular resource is available**
 - If # record = 1 → binary semaphore, **mutex lock**
 - If # record > 1 → counting semaphore
 - Accessed only through 2 **atomic** operations: **wait** & **signal**
- **Spinlock** implementation
 - Semaphore **S** is an integer variable

```
wait (S) {  
    while (S <= 0);  
    S--;  
}  
signal (S) {  
    S++;  
}
```

busy waiting

POSIX Semaphore

- Semaphore is part of **POSIX standard** BUT it is not belonged to pthread
 - It can be used with or without thread
- POSIX Semaphore routines

```
#include <semaphore.h>
```

```
sem_t sem;
```

```
sem_init(&sem);
```

```
sem_wait(&sem);
```

```
critical section
```

```
sem_signal(&sem);
```

```
sem_destroy(&sem);
```

n-Process CS Problem Revisit

- Shared data:

semaphore mutex; // initially mutex = 1

- Process P_i :

```
do {  
    wait(mutex);           // pthread_mutex_lock(&mutex)  
    critical section  
    signal(mutex);        // pthread_mutex_unlock(&mutex)  
    remainder section  
} while (1);
```

progress? **Y**

bounded-wait? **depends on the implementation of wait()**

Semaphores with Non-busy Waiting

- Semaphore is a **data structure with queue**

- May use any queuing strategy (FIFO, FILO, etc)

```
typedef struct {
    int value; // init to # resource
    struct process *L; // PCB queue
} semaphore;
```



- wait()** and **signal()**

- Use system calls: **block()** and **wakeup()**
- Must be **executed atomically**

```
void wait (semaphore S) {
    S.value--; // subtract first
    if (S.value < 0) {
        add this process to S.L;
        sleep();
    }
}
```

```
void signal (semaphore S) {
    S.value++;
    if (S.value <= 0) {
        remove this process from S.L;
        wakeup(P);
    }
}
```

How to Ensure Atomic Wait & Signal Ops?

- Hardware support
 - TestAndSet
 - Swap
- Software solution
 - Peterson's solution
 - Bakery algorithm

Semaphore with Critical Section

```
void wait (semaphore S) {
```

entry section

```
S.value--;
```

```
if (S.value < 0) {
```

add this process to S.L;

exit section

```
sleep();
```

```
} else {
```

exit section

```
}
```

```
}
```

```
void signal (semaphore S) {
```

entry section

```
S.value++;
```

```
if (S.value <= 0) {
```

remove this process from S.L;

exit section

```
wakeup(P);
```

```
} else {
```

exit section

```
}
```

```
}
```

Cooperation Synchronization

- P1 executes S1; P2 executes S2
 - S2 will be executed only after ...

- Implementation

Shared variable:

semaphore sync; // initially sync = 0

P1:

S1;

signal (sync);

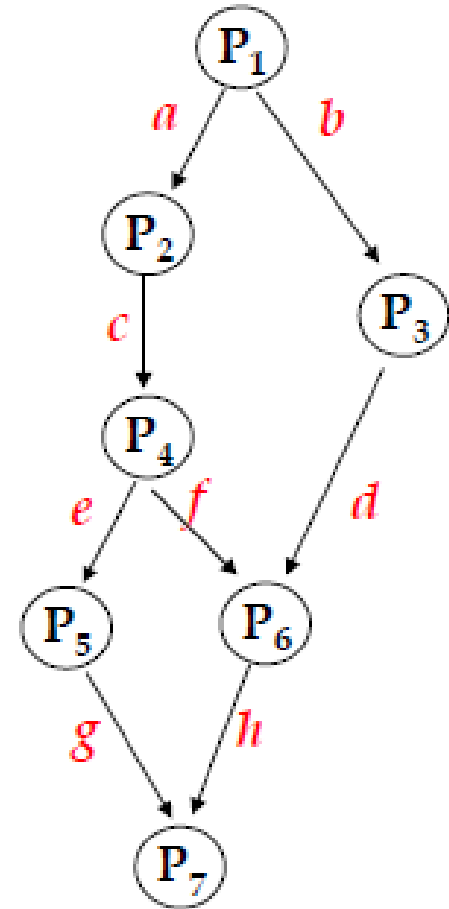
P2:

wait (sync);

S2;

A More Complicated Example

- Initially, all semaphores are 0
- Begin
 - P1: S1; signal(a); signal(b);
 - P2: wait(a); S2; signal(c);
 - P3: wait(b); S3; signal(d);
 - P4: wait(c); S4; signal(e); signal(f);
 - P5: wait(e); S5; signal(g);
 - P6: wait(f); wait(d); S6; signal(h);
 - P7: wait(g); wait(h); S7;
- End



Deadlocks and Starvation

- **Deadlock**

- Two processes are waiting indefinitely for each other to release resources

- **Starvation**

- Some processes (threads) wait infinitely

