



Deadlock

Operating Systems

Yu-Ting Wu

(with slides borrowed from Prof. Jerry Chou)

1

Operating Systems 2022

Outline

- System model
- Deadlock characterization
- Deadlock prevention
- Deadlock avoidance
- Deadlock detection
- Deadlock recovery

2

2

Operating Systems 2022

System Model

3

3

Operating Systems 2022

Deadlock Problem

- A set of blocked processes each **holding** some resources and **waiting** to acquire a resource held by another process in the set
- Example:
 - 2 processes and semaphores A and B
 - P_1 (hold B, wait A): **wait (A)**, signal (B)
 - P_2 (hold A, wait B): **wait (B)**, signal (A)
- Example:
 - Dining philosophers' problem

4


4

Necessary Conditions

- **Mutual exclusion**
 - Only 1 process at a time can use a resource
- **Hold and wait**
 - A process holding some resources and is waiting for another resource
- **No preemption**
 - A resource can be only released by a process **voluntarily**
- **Circular wait**
 - There exists a set $\{P_0, P_1, \dots, P_n\}$ of waiting processes such that $P_0 \rightarrow P_1 \rightarrow P_2 \rightarrow \dots \rightarrow P_n \rightarrow P_0$

5

Necessary Conditions (cont.)



6

System Model

- **Resources types** R_1, R_2, \dots, R_m
 - E.g. CPU, memory pages, I/O devices
- Each resource type R_i has W_i **instances**
 - E.g. a computer has 2 CPUs
- Each process utilizes a resource as follows:
 - **Request** \rightarrow **use** \rightarrow **release**

7

Resource-Allocation Graph

- 3 processes, $P_1 \sim P_3$
- 4 resources, $R_1 \sim R_4$
 - R_1 and R_3 each has one instance
 - R_2 has two instances
 - R_4 has three instances


Request edges

- $P_1 \rightarrow R_1$: P_1 requests R_1

Assignment edges

- $R_2 \rightarrow P_1$: one instance of R_2 is allocated to P_1

$\rightarrow P_1$ is **holding on** an instance of R_2 and **waiting for** an instance of R_1



8

Operating Systems 2022

Resource-Allocation Graph w/ Deadlock

- If the graph contains a **cycle**, a deadlock **may** exist
- In the example
 - P_1 is waiting for P_2
 - P_2 is waiting for P_3
 - P_1 is also waiting for P_3
 - Since P_3 is waiting for P_1 or P_2 , and they both waiting for P_3
 - **Deadlock !**

9

Operating Systems 2022

RA Graph w/ Cycle but NO Deadlock

- If the graph contains a **cycle**, a deadlock **may** exist
- In the example
 - P_1 is waiting for P_2 or P_3
 - P_3 is waiting for P_1 or P_4
 - Since P_2 and P_4 wait for no one
 - **No Deadlock between P_1 and P_3**

10

Operating Systems 2022

Deadlock Detection

- If the graph contains **no cycle** → **no deadlock**
 - Circular wait cannot be held**
- If the graph contains a cycle
 - If **one instance** per resource type → **deadlock**
 - If **multiple instances** per resource type → **possibility** of deadlock

11

Operating Systems 2022

Handling Deadlocks

- Ensure the system will **never** enter a deadlock state
 - Deadlock prevention**: ensure that at least one of the **4 necessary conditions** cannot hold
 - Deadlock avoidance**: **dynamically** examines the resource-allocation state before allocation
- Allow to **enter a deadlock state** and then **recover**
 - Deadlock detection**
 - Deadlock recovery**
- Ignore the problem** and pretend that deadlocks never occur in the system
 - Used by most operating systems, including UNIX**

12

Deadlock Prevention

13

13

Deadlock Prevention

- **Mutual exclusion (ME):** do not require ME on sharable resources
 - E.g. there is no need to ensure ME on read-only files
 - However, some resources are not shareable (e.g. printer)
- **Hold and wait:**
 - When a process requests a resource, it does not hold any resource
 - Pre-allocate all resources before executing
 - Resource utilization is low; starvation is possible

14

14

Deadlock Prevention (cont.)

- **No preemption:**
 - When a process is waiting on a resource, all its holding resources are preempted
 - E.g. P_1 request R_1 , which is allocated to P_2 , which in turn is waiting on R_2 ($P_1 \rightarrow R_1 \rightarrow P_2 \rightarrow R_2$)
 - **R_1 can be preempted and reallocated to P_1**
- **Applied to resources whose states can be easily saved and restored later**
 - E.g. CPU registers and memory
- It cannot easily be applied to other resources
 - E.g. printers and tape drives

15

15

Deadlock Prevention (cont.)

- **Circular wait:**
 - Impose a **total ordering** of all resource types
 - A process requests resources in an increasing order
 - Let $R = \{R_0, R_1, \dots, R_n\}$ be the set of resource types
 - **When request R_k , should release all $R_i, i \geq k$**
- **Example**
 - $F(\text{disk drive}) = 5, F(\text{printer}) = 12$
 - A process must request disk drive before printer
- **Proof: counter-example does not exist**
 - $P_0(R_0) \rightarrow R_1, P_1(R_1) \rightarrow R_2, \dots, P_n(R_n) \rightarrow R_0$ ← P_n holds on R_n , waiting for R_0
 - **Conflict: $R_0 < R_1 < R_2 < \dots < R_n < R_0$**

16

16

Deadlock Avoidance

17

17

Avoidance Algorithms


- **Single instance of a resource type**
 - **Resource-allocation graph (RAG) algorithm** based on **circle detection**
- **Multiple instance of a resource type**
 - **banker's algorithm** based on safe **sequence detection**

18

18

Resource-Allocation Graph Algorithm

- **Request edges**
 - $P_i \rightarrow R_j$: P_i **is waiting** for resource R_j
- **Assignment edges**
 - $R_i \rightarrow P_j$: Resource R_i **is allocated** and held by P_j
- **Claim edge**
 - Process P_i **may** request R_j in the future
- **Claim edge** converts to **request edge**
 - When a resource **is requested** by process
- **Assignment edge** converts back to a **claim edge**
 - When a resource is released by a process




19

19

Resource-Allocation Graph Algorithm (cont.)

- Resource **must be claimed a priori** in the system
- **Grant a request** only if **NO cycle created**
- Check for safety using a **cycle-detection algorithm**, $O(n^2)$
- Example: R_2 cannot be allocated to P_2



20

20

Avoidance Algorithms


- **Single instance of a resource type**
 - **Resource-allocation graph (RAG) algorithm** based on **circle detection**
- **Multiple instance of a resource type**
 - **banker's algorithm** based on safe **sequence detection**

21

21

Deadlock Avoidance

- **Safe state:** a system is in a safe state if there exists **a sequence of allocations** to satisfy requests by all processes
 - This sequence of allocations is called **safe sequence**
- **Safe state** → **no deadlock**
- **Unsafe state** → **possibility of deadlock**
- **Deadlock avoidance** → **ensure that a system never enters an unsafe state**



22

22

Safe State with Safe Sequence

- There are 12 tape drives
- Assuming at t_0 : **hints from processes**

	Max Needs	Current Holding
P0	10	5
P1	4	2
P2	9	2

→ $\langle P_1, P_0, P_2 \rangle$ is a safe sequence

23

23

Safe State with Safe Sequence

- There are 12 tape drives
- Assuming at t_0 :

	Max Needs	Current Holding	Available
P0	10	5	3
P1	4	2	
P2	9	2	

→ $\langle P_1, P_0, P_2 \rangle$ is a safe sequence

1. P_1 satisfies its allocation with 3 available resources

24

24

Safe State with Safe Sequence

- There are 12 tape drives
- Assuming at t_0 :

	Max Needs	Current Holding	Available
P0	10	5	5
P1	4	0	
P2	9	2	

→ $\langle P_1, P_0, P_2 \rangle$ is a safe sequence

1. P_1 satisfies its allocation with 3 available resources
2. P_0 satisfies its allocation with 5 available resources

25

25

Safe State with Safe Sequence

- There are 12 tape drives
- Assuming at t_0 :

	Max Needs	Current Holding	Available
P0	10	5	
P1	4	0	
P2	9	2	10

→ $\langle P_1, P_0, P_2 \rangle$ is a safe sequence

1. P_1 satisfies its allocation with 3 available resources
2. P_0 satisfies its allocation with 5 available resources
3. P_2 satisfies its allocation with 10 available resources

26

26

Safe State with Safe Sequence

- There are 12 tape drives
- Assuming at t_1 :

	Max Needs	Current Holding	Available
P0	10	5	
P1	4	2	2
P2	9	3	

- If P_2 requests and is allocated 1 more resource

- No safe sequence exist ...
- This allocation makes the system enter an unsafe state

- **A request is only granted if the allocation leaves the system in a safe state**

27

27

Banker's Algorithm

- Use for **multiple instances** of each resource type

• Banker's Algorithm

- Use a general safety algorithm to **pre-determine** if any **safe sequence** exists after allocation
- **Only proceed the allocation if safe sequence exists**

• Safety algorithm

1. Assume processes need **maximum** resources
2. Find a process that can be satisfied by free resources
3. Free the resource usage of the process
4. Repeat to step 2 until all processes are satisfied

28

28

Banker's Algorithm Example

- Total instances: A: 10, B: 5, C: 7
- Available instances: A: 3, B: 3, C: 2

	Max			Allocation			Need (Max - Alloc.)		
	A	B	C	A	B	C	A	B	C
P ₀	7	5	3	0	1	0	7	4	3
P ₁	3	2	2	2	0	0	1	2	2
P ₂	9	0	2	3	0	2	6	0	0
P ₃	2	2	2	2	1	1	0	1	1
P ₄	4	3	3	0	0	2	4	3	1

- Safe sequence: P₁

29

Banker's Algorithm Example

- Total instances: A: 10, B: 5, C: 7
- Available instances: A: **5**, B: 3, C: 2

	Max			Allocation			Need (Max - Alloc.)		
	A	B	C	A	B	C	A	B	C
P ₀	7	5	3	0	1	0	7	4	3
P ₁	3	2	2	2	0	0	1	2	2
P ₂	9	0	2	3	0	2	6	0	0
P ₃	2	2	2	2	1	1	0	1	1
P ₄	4	3	3	0	0	2	4	3	1

- Safe sequence: P₁, P₃

30

Banker's Algorithm Example

- Total instances: A: 10, B: 5, C: 7
- Available instances: A: **7**, B: **4**, C: **3**

	Max			Allocation			Need (Max - Alloc.)		
	A	B	C	A	B	C	A	B	C
P ₀	7	5	3	0	1	0	7	4	3
P ₁	3	2	2	2	0	0	1	2	2
P ₂	9	0	2	3	0	2	6	0	0
P ₃	2	2	2	2	1	1	0	1	1
P ₄	4	3	3	0	0	2	4	3	1

- Safe sequence: P₁, P₃, P₄

31

Banker's Algorithm Example

- Total instances: A: 10, B: 5, C: 7
- Available instances: A: 7, B: 4, C: **5**

	Max			Allocation			Need (Max - Alloc.)		
	A	B	C	A	B	C	A	B	C
P ₀	7	5	3	0	1	0	7	4	3
P ₁	3	2	2	2	0	0	1	2	2
P ₂	9	0	2	3	0	2	6	0	0
P ₃	2	2	2	2	1	1	0	1	1
P ₄	4	3	3	0	0	2	4	3	1

- Safe sequence: P₁, P₃, P₄, P₂

32

Banker's Algorithm Example

- Total instances: A: 10, B: 5, C: 7
- Available instances: A: **10**, B: 4, C: **7**

	Max			Allocation			Need (Max - Alloc.)		
	A	B	C	A	B	C	A	B	C
P ₀	7	5	3	0	1	0	7	4	3
P ₁	3	2	2	2	0	0	1	2	2
P ₂	9	0	2	3	0	2	6	0	0
P ₃	2	2	2	2	1	1	0	1	1
P ₄	4	3	3	0	0	2	4	3	1

- Safe sequence: P₁, P₃, P₄, P₂, P₀

33

Banker's Algorithm Example

- Total instances: A: 10, B: 5, C: 7
- Available instances: A: 3, B: 3, C: 2

	Max			Allocation			Need (Max - Alloc.)		
	A	B	C	A	B	C	A	B	C
P ₀	7	5	3	0	1	0	7	4	3
P ₁	3	2	2	2	0	0	1	2	2
P ₂	9	0	2	3	0	2	6	0	0
P ₃	2	2	2	2	1	1	0	1	1
P ₄	4	3	3	0	0	2	4	3	1

- If Request (P₁) = (1, 0, 2) ...

34

Banker's Algorithm Example

- Total instances: A: 10, B: 5, C: 7
- Available instances: A: **2**, B: 3, C: **0**

	Max			Allocation			Need (Max - Alloc.)		
	A	B	C	A	B	C	A	B	C
P ₀	7	5	3	0	1	0	7	4	3
P ₁	3	2	2	3	0	2	0	2	0
P ₂	9	0	2	3	0	2	6	0	0
P ₃	2	2	2	2	1	1	0	1	1
P ₄	4	3	3	0	0	2	4	3	1

- If Request (P₁) = (1, 0, 2): P₁ allocation → (3, 0, 2)
 - Enter another safe state (Safe sequence: P₁, P₃, P₄, P₀, P₂)

35

Banker's Algorithm Example

- Total instances: A: 10, B: 5, C: 7
- Available instances: A: **0**, B: **0**, C: 2

	Max			Allocation			Need (Max - Alloc.)		
	A	B	C	A	B	C	A	B	C
P ₀	7	5	3	0	1	0	7	4	3
P ₁	3	2	2	2	0	0	1	2	2
P ₂	9	0	2	3	0	2	6	0	0
P ₃	2	2	2	2	1	1	0	1	1
P ₄	4	3	3	3	3	2	1	0	1

- If Request (P₄) = (3, 3, 0): P₄ allocation → (3, 3, 2)
 - Enter into an unsafe state (no safe sequence can be found)

36

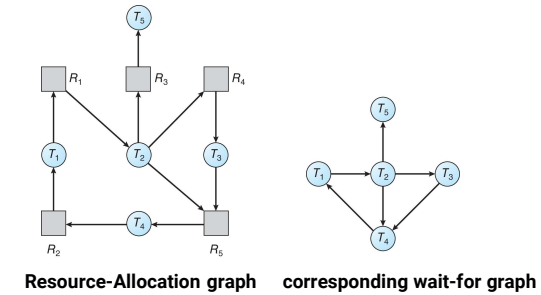
Deadlock Detection

37

37

Deadlock Detection

- **Single instance** of each resource type
 - Convert request/assignment edges into **wait-for graph**
 - Deadlock exists if there is a cycle in the wait-for graph



38

38

Multiple Instance for Each Resource Type

- Total instances: A: 7, B: 2, C: 6
- Available instances: A: 0, B: 0, C: 0

	Allocation			Request		
	A	B	C	A	B	C
P ₀	0	1	0	0	0	0
P ₁	2	0	0	2	0	2
P ₂	3	0	3	0	0	0
P ₃	2	1	1	1	0	0
P ₄	0	0	2	0	0	2

- The system is in a safe state → <P₀, P₂, P₃, P₁, P₄>
→ No deadlock

39

39

Multiple Instance for Each Resource Type

- Total instances: A: 7, B: 2, C: 6
- Available instances: A: 0, B: 0, C: 0

	Allocation			Request		
	A	B	C	A	B	C
P ₀	0	1	0	0	0	0
P ₁	2	0	0	2	0	2
P ₂	3	0	3	0	0	1
P ₃	2	1	1	1	0	0
P ₄	0	0	2	0	0	2

- If P₂ requests (0, 0, 1) → no safe sequence can be found
→ The system is deadlocked

40

40

Deadlock Recovery

41

41

Deadlock Recovery

• Process termination

- Abort all deadlocked processes
- Abort 1 process at a time until the deadlock cycle is eliminated
 - Which process should we abort first?

• Resource preemption

- Select a victim: which one to preempt?
- Rollback: partial rollback or total rollback?
- Starvation: can the same process be preempted always?

42

42

Objective Review

- Illustrate how deadlock can occur
- Define the four necessary conditions that characterize deadlock
- Identify a deadlock situation in a resource allocation graph
- Evaluate the four different approaches for preventing deadlocks
- Apply the banker's algorithm for deadlock avoidance
- Apply the deadlock detection algorithm

43

43