



Main Memory (I)

Operating Systems

Yu-Ting Wu

(with slides borrowed from Prof. Jerry Chou)

Outline

- Background
- Swapping
- Contiguous allocation
- Paging

Background

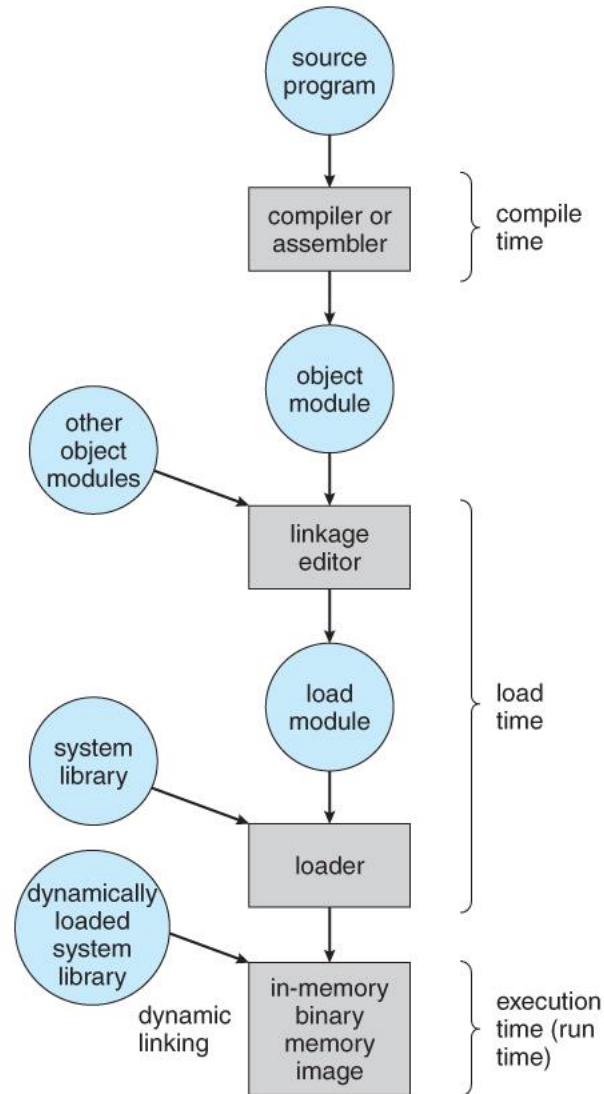
Background

- **Main memory** and **registers** are the only storage CPU can access directly
- Collection of **processes** are waiting on disk to be brought into memory and be executed
- **Multiple programs** are brought into memory to improve resource utilization and response time to users
- A process may be **moved between disk and memory** during run time

Questions

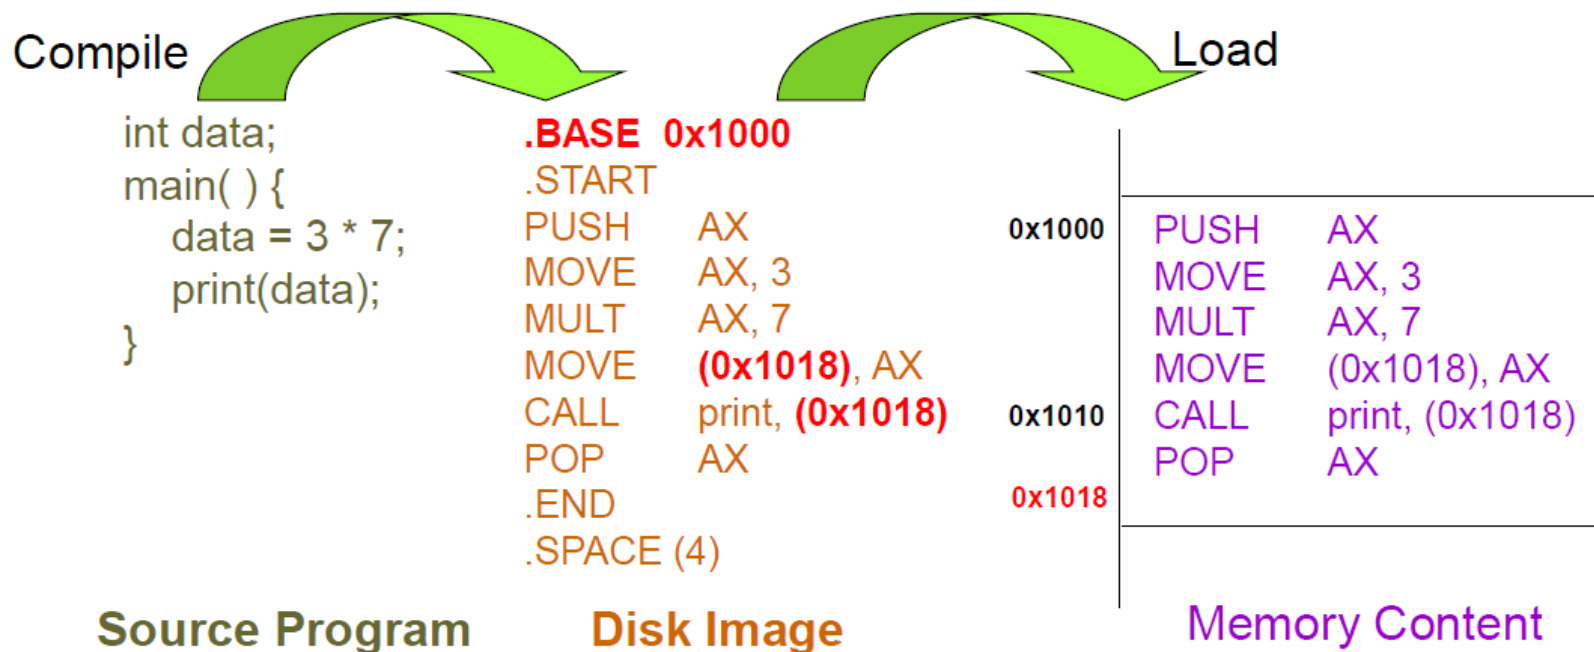
- How to refer memory in a program?
 - Address binding
- How to load a program into memory?
 - Static / dynamic loading and linking
- How to move a program between memory and disk?
 - Swap
- How to allocate memory
 - Paging, segment

Steps of Processing a Program



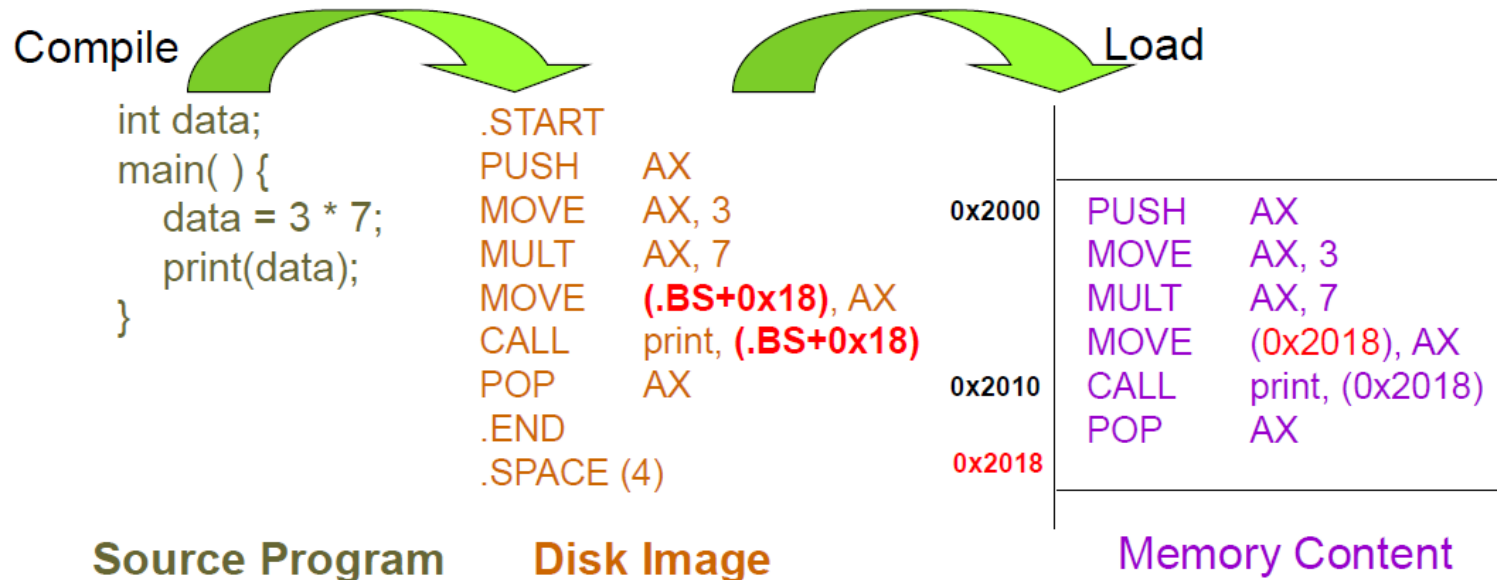
Address Binding: Compile Time

- Program is written as symbolic code
- Compiler translates symbolic code into **absolute code**
- If starting location changes → **recompile**



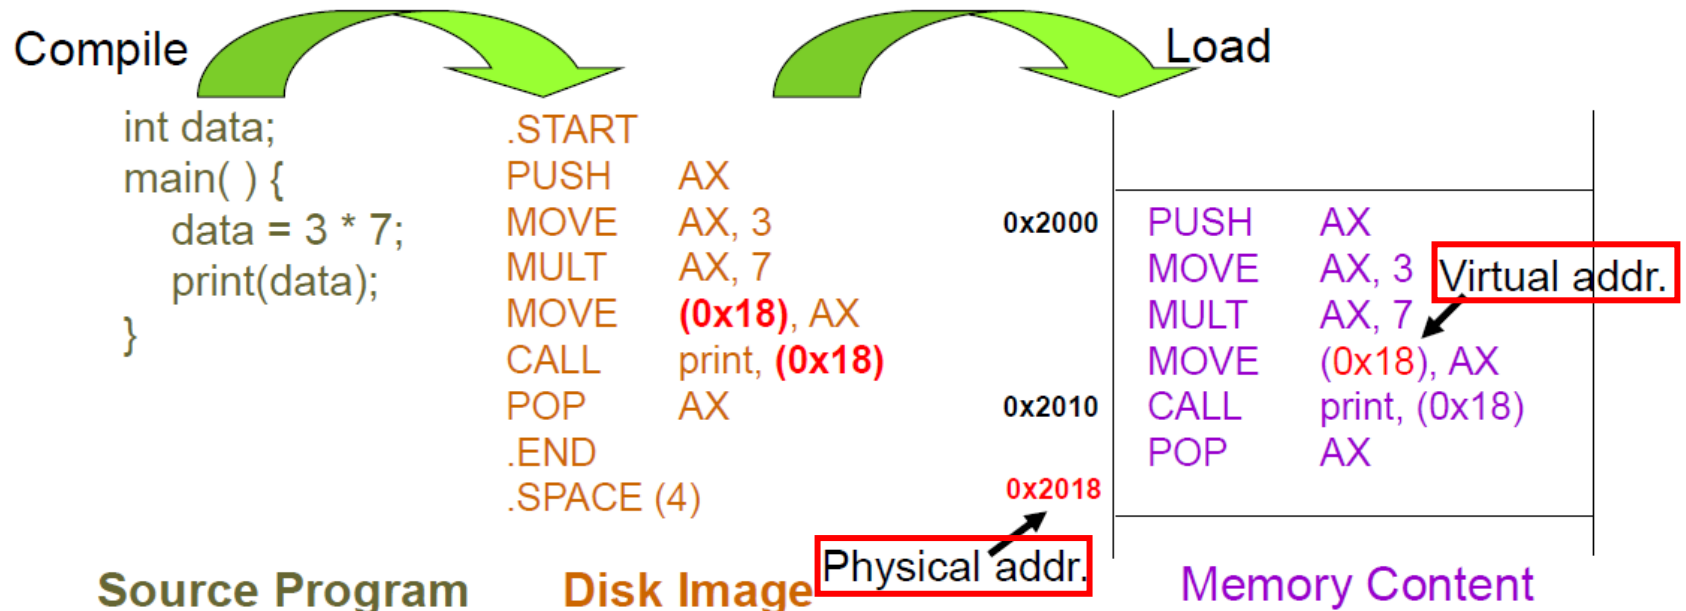
Address Binding: Load Time

- Compiler translates symbolic code into **relocatable code**
- **Relocatable code**
 - Machine language that can be run from any memory location
 - If starting location changes → **reload the code**



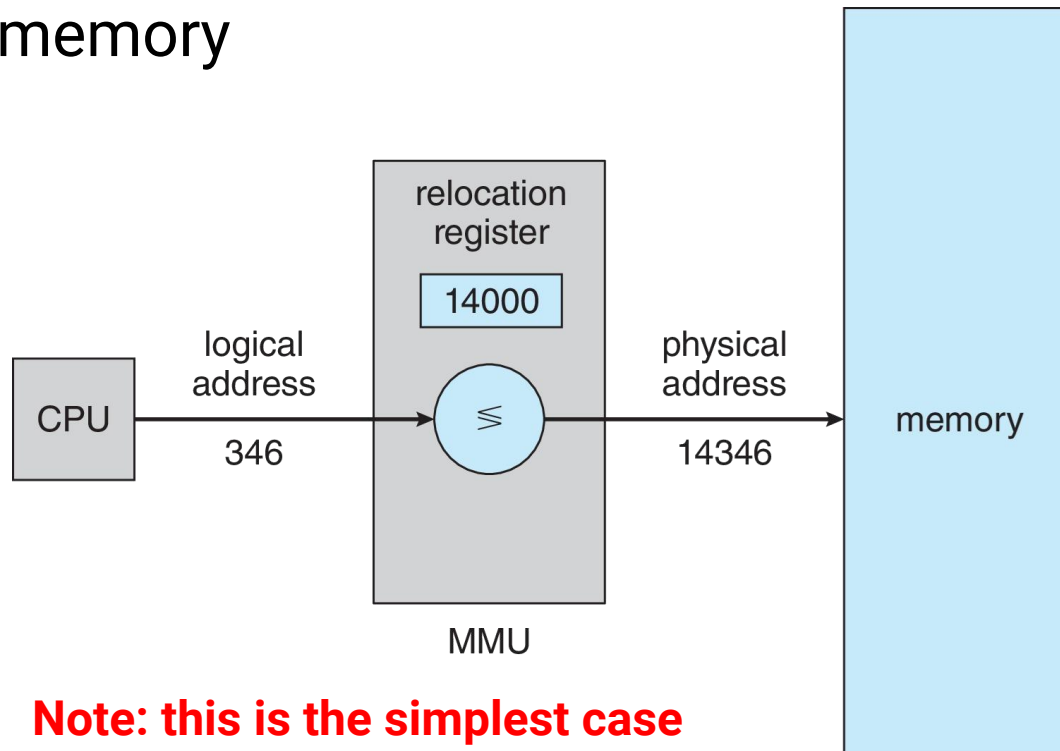
Address Binding: Execution Time

- Compiler translates symbolic code into **logical-address** (i.e. **virtual-address**) code
- **Special hardware** (i.e. **MMU**) is needed for this scheme
- Most general-purpose OS use this method



Memory-Management Unit (MMU)

- **Hardware** device that **maps virtual to physical** address
- The value in the **relocation register** is **added to every address** generated by a user process at the time it is sent to memory



Logical v.s. Physical Address

- **Logical address** – generated by CPU
 - a.k.a virtual address
- **Physical address** – seen by the memory module
- **Compile-time** and **load-time** address binding
 - Logical address = physical address
- **Execution-time** address binding
 - Logical address \neq physical address
- The user program deals with **logical** addresses; it never sees the real physical addresses

Questions

- How to refer memory in a program?
 - Address binding
- How to load a program into memory?
 - Static / dynamic loading and linking
- How to move a program between memory and disk?
 - Swap
- How to allocate memory
 - Paging, segment

Dynamic Loading

- The entire program must be in memory for it to execute?
- No, we can use **dynamic loading**
 - A routine is loaded into memory when it is called
- **Better memory-space utilization**
 - Unused routine is never loaded
 - Particularly useful when large amounts of code are infrequently used (e.g., error handling code)
- **No special support from OS** is required, implemented through programs (library, API calls)

Dynamic Loading Example in C

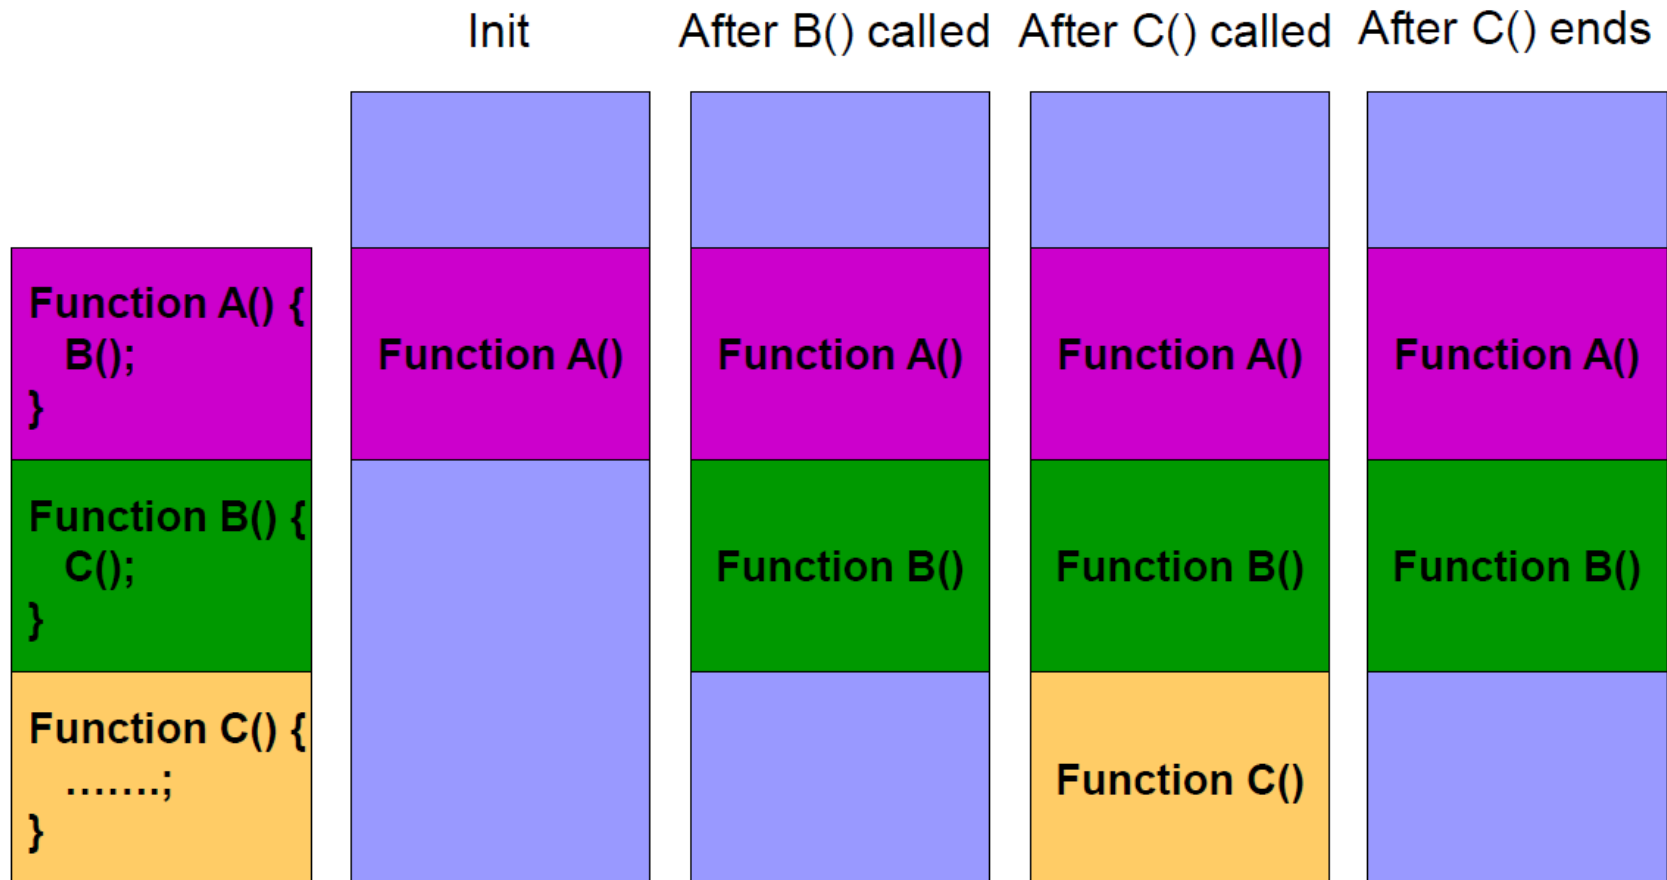
- ***dlopen()***: opens a library and prepares it for use
- ***dlsym()***: looks up the value of a symbol in a given (opened) library
- ***dlclose()***: closes a DL library

```
#include <dlfcn.h>
int main() {
    double (*cosine)(double);
    void* handle = dlopen ("/lib/libm.so.6", RTLD_LAZY);
    cosine = dlsym(handle, "cos");
    printf ("%f\n", (*cosine)(2.0));
    dlclose(handle);
}
```

Dynamic Loading

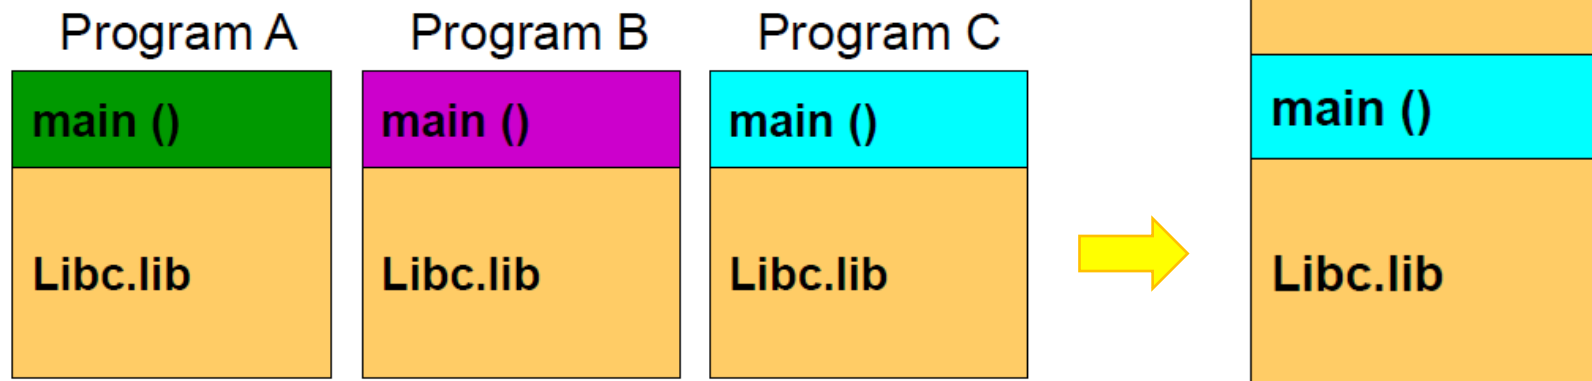
Disk image

Memory content



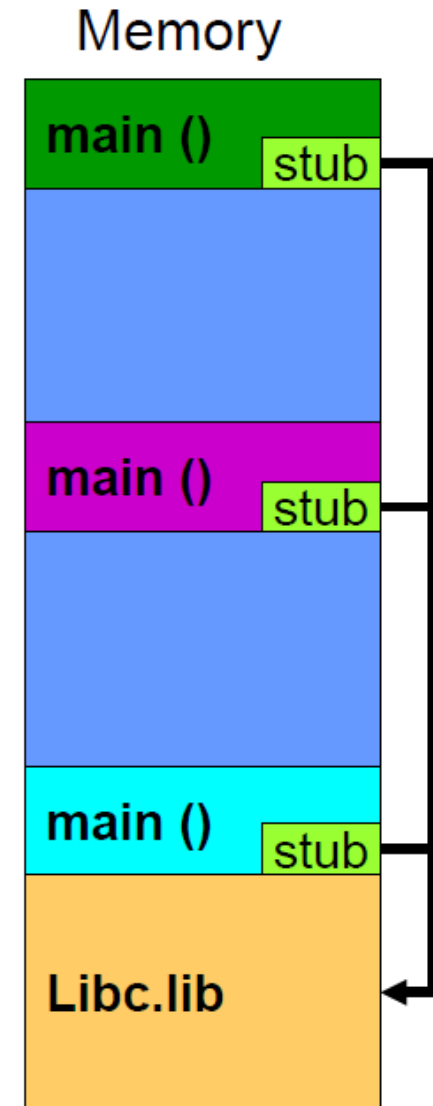
Static Linking

- **Static linking:** libraries are combined by the loader into the program in-memory image
 - Waste memory: **duplicated code**
 - Faster during execution time
- **Static linking + dynamic loading ?**
 - Still can't prevent duplicate code



Dynamic Linking

- **Dynamic linking:** linking postponed until execution time
 - **Only one code copy** in memory and shared by everyone
 - **A stub** is included in the program in-memory image for each lib reference
 - **Stub call**
 - ➔ check if the referred lib is in memory
 - ➔ if not, load the lib
 - ➔ execute the lib
 - DLL (dynamic link library) on Windows



Swapping

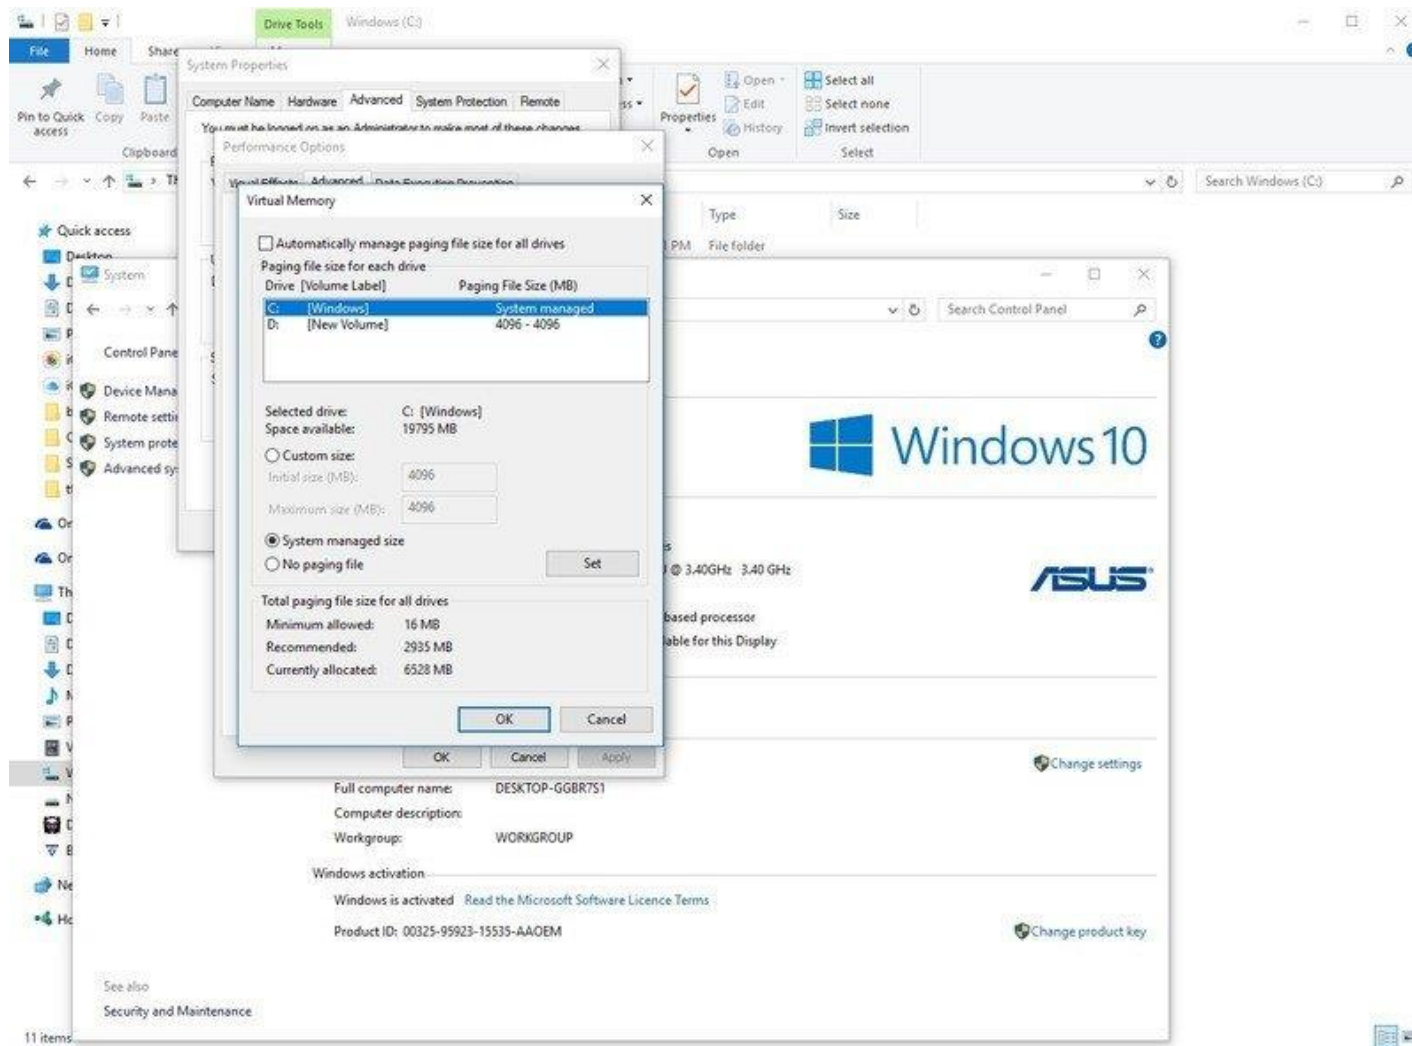
Questions

- How to refer memory in a program?
 - Address binding
- How to load a program into memory?
 - Static / dynamic loading and linking
- How to move a program between memory and disk?
 - Swap
- How to allocate memory
 - Paging, segment

Swapping

- A process can be swapped out of memory to a **backing store**, and later brought back into memory for continuous execution
 - Also used by **midterm scheduling**, different from context switch
- **Backing store** – a chunk of the disk, **separated from the file system**, to provide direct access to these memory images
- Why swap a process?
 - Free up memory
 - **Roll out, roll in**: swap lower-priority process with a higher one

Swapping (cont.)




Swapping (cont.)

- Swap back memory location
 - If binding is done at compile / load time
 - Swap back memory address must be the **same**
 - If binding is done at execution time
 - Swap back memory address can be **different**
- A process to be swapped → **must be idle**
 - Imagine a process that is waiting for I/O is swapped
 - Solutions:
 - Never swap a process with pending I/O
 - I/O operations are done through **OS buffers** (i.e. a memory space not belongs to any user processes)

Process Swapping to Backing Store

- **Major part of swap time is transfer time**; total transfer time is directly proportional to the amount of memory swapped



Contiguous Allocation

Memory Allocation

- **Fixed-partition allocation**

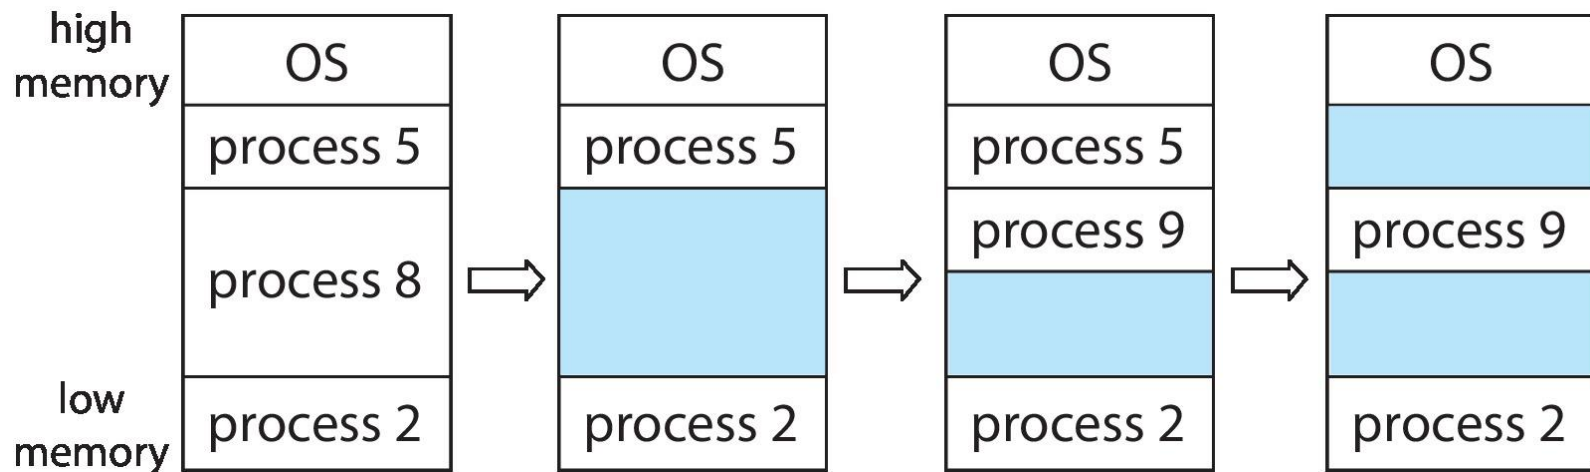
- Each process loads into one partition of fixed-size
- **Degree of multi-programming** is bounded by the number of partitions

- **Variable-size partition**

- Hole: block of contiguous free memory
- Holes of various sizes are scattered in memory

Multiple Partition (Variable-Size) Method

- When a process arrives, it is allocated a hole **large enough** to accommodate it
- The OS maintains info. of each in-use and free hole
- A freed hole can be merged with another hole to form a larger hole




Dynamic Storage Allocation Problem

- How to satisfy a request of size n from a list of free holes
- **First-fit** – allocate the 1st hole that fits
- **Best-fit** – allocate the smallest hole that fits
 - Must search through the whole list
- **Worst-fit** – allocate the largest hole
 - Must also search through the whole list
- **First-fit** and **best-fit** are better than worst-fit in terms of **speed** and **storage utilization**

Fragmentation


• External fragmentation

- Total free memory space is big enough to satisfy a request but is not contiguous
- Occur in **variable-size allocation**
- Solution: **compaction**
 - Shuffle the memory contents to place all free memory together in one large block at execution time
 - Only if the binding is done at execution time



• Internal fragmentation

- Memory that is internal to a partition but is not being used
- Occur in **fixed-partition allocation**



Paging

(Non-Contiguous Memory Allocation)


Paging Concept

- Method
 - Divide **physical memory** into fixed-size blocks called **frames**
 - Divide **logical address** space into blocks of the **same size** called **pages**
 - To run a program of n pages, need to find n free frames and load the program
 - **Must keep track of free frames**
 - Set up a **page table** to translate logical to physical addresses
- Benefit
 - Allow the physical-address space of a process to be **noncontiguous**
 - Avoid external fragmentation
 - Limited internal fragmentation
 - Provide **shared memory / pages**

Paging Example

- **Page table**

- Each entry maps to the **base address of a page** in physical memory
- A structure maintained by OS **for each process**
 - Page table includes only pages owned by a process
 - A process cannot access memory outside its space




Address Translation Scheme

- Logical address is divided into two parts
 - **Page number (p)**
 - Used as an **index into a page table** which contains **base address** of each page in physical memory
 - N bits means a process can allocate at most 2^N pages
 - **$2^N \times$ page size memory size**
 - **Page offset (d)**
 - Combined with base address to define the physical memory address that is sent to the memory unit
 - **N bits means the page size is 2^N**
- **Physical address = page base address + page offset**

Address Translation Architecture

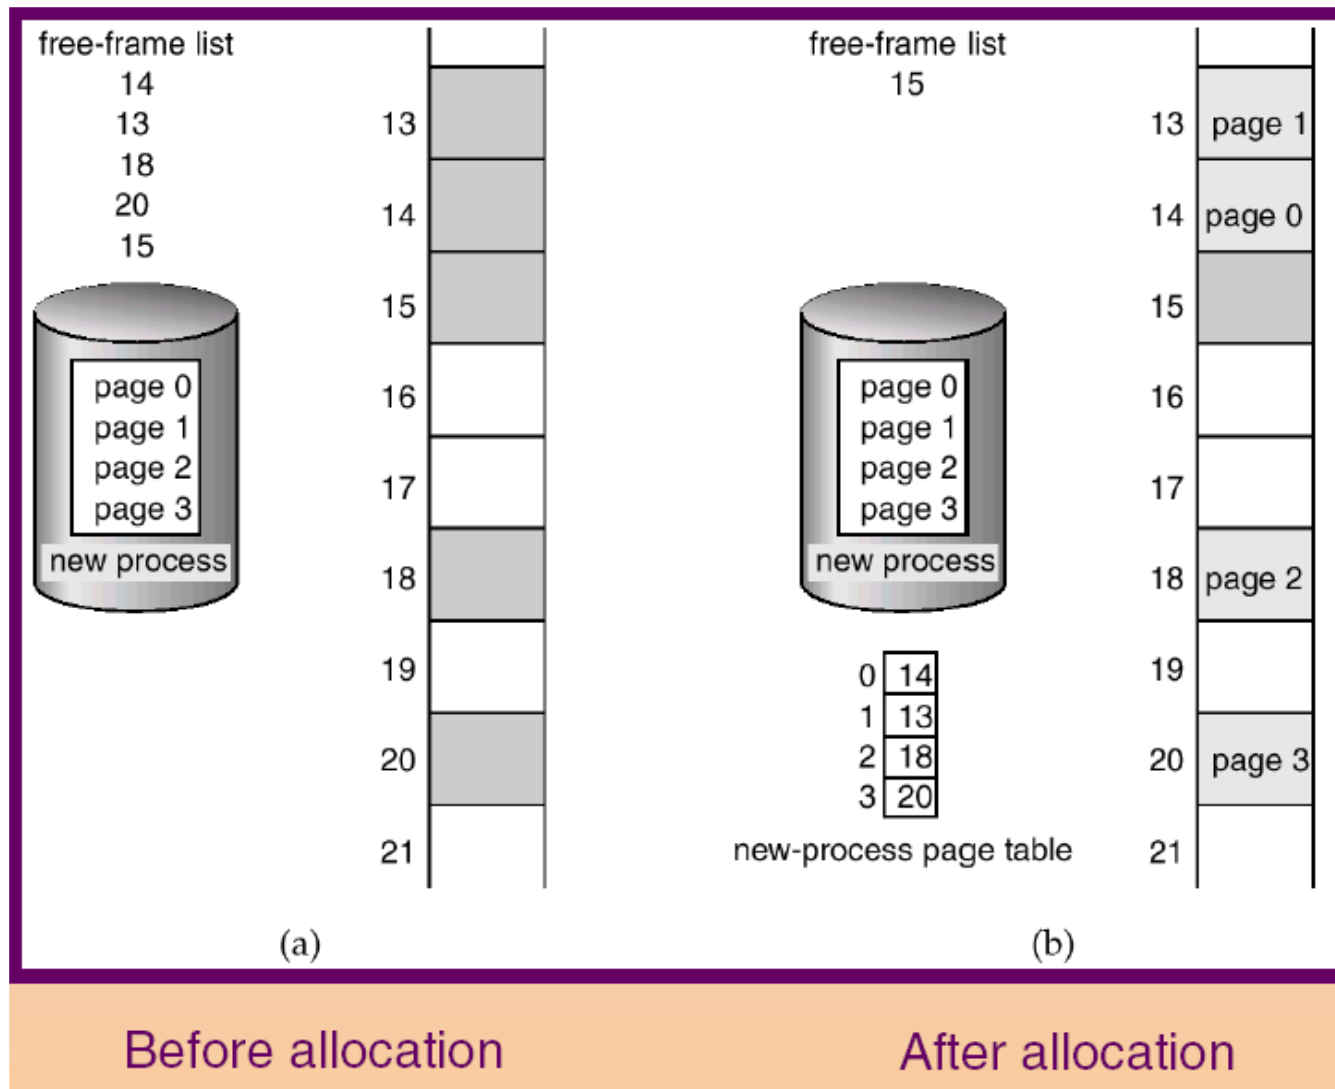
- If page size is 1KB (2^{10}) and page 3 maps to frame 5
- Given 13 bits logical address ($p = 3$, $d = 20$), what is the physical address?
 - $5 * (1\text{KB}) + 20 = \underline{101}0000000000 + 0000010100 =$
 $\underline{101}0000010100$



Address Translation

- Total number of pages does not need to be the same as the total number of frames
 - **Total # pages** determines the logical memory size of a process
 - **Total # frames** depending on the size of physical memory
- E.g.: Given 32 bits logical address, 36 bits physical address, and 4KB page size, what does it mean?
 - Number of bits for page offset: 4KB page size = 2^{12} bytes \rightarrow 12
 - Number of bits for page number: 2^{20} pages \rightarrow 20 bits
 - Page table size: $2^{32} / 2^{12} = 2^{20}$ entries
 - Max program memory: $2^{32} = 4\text{GB}$
 - Number of bits for frame number: 2^{24} frames \rightarrow 24 bits
 - Total physical memory size: $2^{36} = 64\text{GB}$

Free Frames



Page / Frame Size

- The page (frame) size is defined by hardware
 - **Typically, a power of 2**
 - Ranging from 512 bytes to 16 MB / page
 - **4KB / 8KB** page is commonly used
- Internal fragmentation?
 - Larger page size → More space waste
- But **page sizes cannot be too small**
 - Memory, process, and data sets have become larger
 - Need to keep page table small
 - Fewer access means better I/O performance

Paging Summary

- Paging helps separate **user's view** of memory and the actual **physical memory**
- User view's memory: one single contiguous space
 - Actually, user's memory is scattered out in physical memory
- OS maintains a copy of the **page table** for each process
- OS maintains a **frame table** for managing physical memory
 - One entry for each physical frame
 - Indicate whether a frame is free or allocated
 - If allocated, to which page of which process or processes

Implementation of Page Table

- Page table is kept **in memory**
- **Page-table base register (PTBR)**
 - The **physical memory address** of the page table
 - The PTBR value is stored in **PCB** (Process Control Block)
 - Changing the value of PTBR during the context switch
- With PTBR, each memory reference results in **2 memory reads**
 - One for the page table and one for the real address
- The 2-access problem can be solved by
 - **Translate Look-aside Buffers (TLB)** (HW) which is implemented by **Associative memory** (HW)

Associative Memory

- All memory entries can be accessed at the same time
 - Each entry corresponds to an associative register
- But **the number of entries are limited**
 - Typical number of entries: 64 ~ 1024

■ Associative memory – parallel search

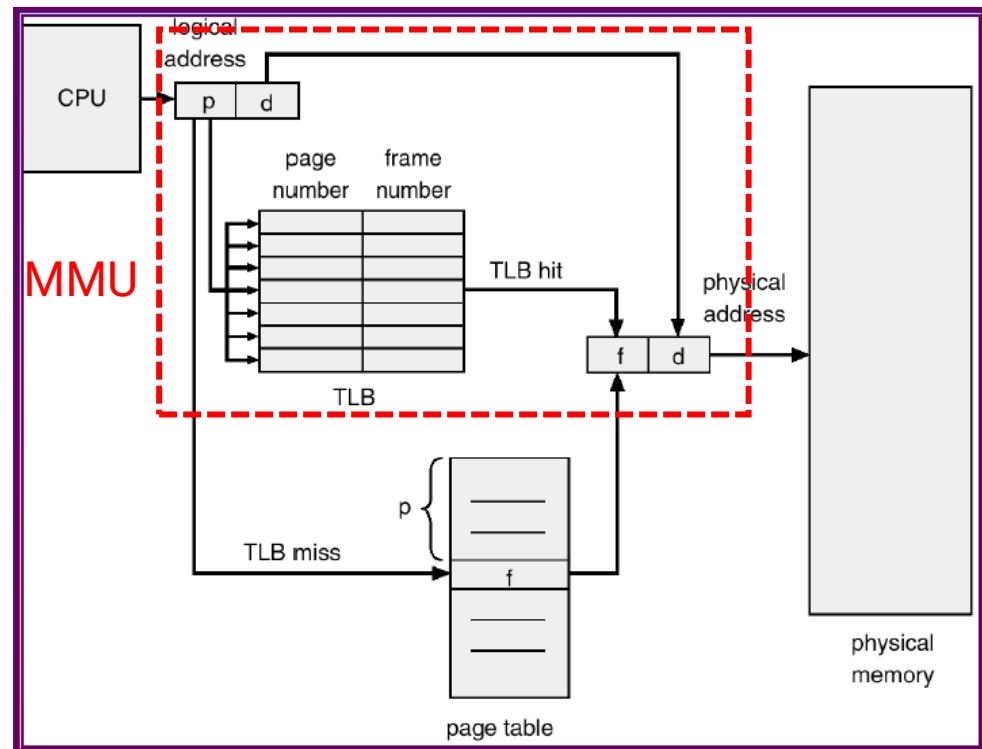
Page #	Frame #

Address translation (A' , A'')

- ◆ If A' is in associative register, get frame # out.
- ◆ Otherwise get frame # from page table in memory

Translation Look-aside Buffer (TLB)

- A cache for page table shared by all processes
- TLB must be **flushed** after a context switch
 - Otherwise, TLB entry must have a PID field (address-space identifiers (ASIDs))



Effective Memory-Access Time

- 20 ns for TLB search
- 100 ns for memory access
- Effective Memory-Access Time (**EMAT**)
 - 70% TLB hit-ratio:
→ $EMAT = 0.70 \times (20 + 100) + (1 - 0.70) \times (20 + 100 + 100) = 150 \text{ ns}$
 - 98% TLB hit-ratio:
→ $EMAT = 0.98 \times 120 + 0.02 \times 220 = 122 \text{ ns}$