



Virtual Memory

Operating Systems

Yu-Ting Wu

(with slides borrowed from Prof. Jerry Chou)

Outline

- Background
- Demand Paging
- Process Creation
- Page Replacement
- Allocation of Frames
- Thrashing

Background

Background


- Why we don't want to run a program that is entirely in memory
 - Many code for handling unusual errors or conditions
 - Certain program routines or features are rarely used
 - The same library code used by many programs
 - Arrays, lists, and tables allocated but not used

→ We want better utilization

Virtual Memory

- Separation of user logical memory from physical memory
 - To run an **extremely large process**
 - Logical address space can be much larger than physical address space
 - To increase **CPU/resource utilization**
 - Higher degree of multiprogramming degree
 - To simplify programming (compiler) tasks
 - Free programmer from memory limitation
 - To **launch** programs **faster**
 - Less I/O would be needed to load or swap
- Can be implemented via
 - **Demand paging**
 - Demand segmentation (more complicated due to variable sizes)

Virtual Memory v.s. Physical Memory



Demand Paging


Demand Paging

- **A page rather than the whole process** is brought into memory only when it is needed
 - Less I/O needed → fast response
 - Less memory needed → more users
- Page is needed when there is a reference to the page
 - Invalid reference → abort
 - Not-in-memory → bring to memory via **paging**
- **Pure demand paging**
 - Start a process with no page
 - Never bring a page into memory until it is required

Demand Paging (cont.)

- A swapper (midterm scheduler) manipulates the **entire process**, whereas a **pager** is concerned with the **individual pages of a process**
- Hardware support
 - **Page table: a valid-invalid bit**
 - 1 → page in memory
 - 0 → page not in memory
 - Initially, all such bits are set to 0
 - Secondary memory (swap space, backing store): usually a high-speed disk (swap device) is used

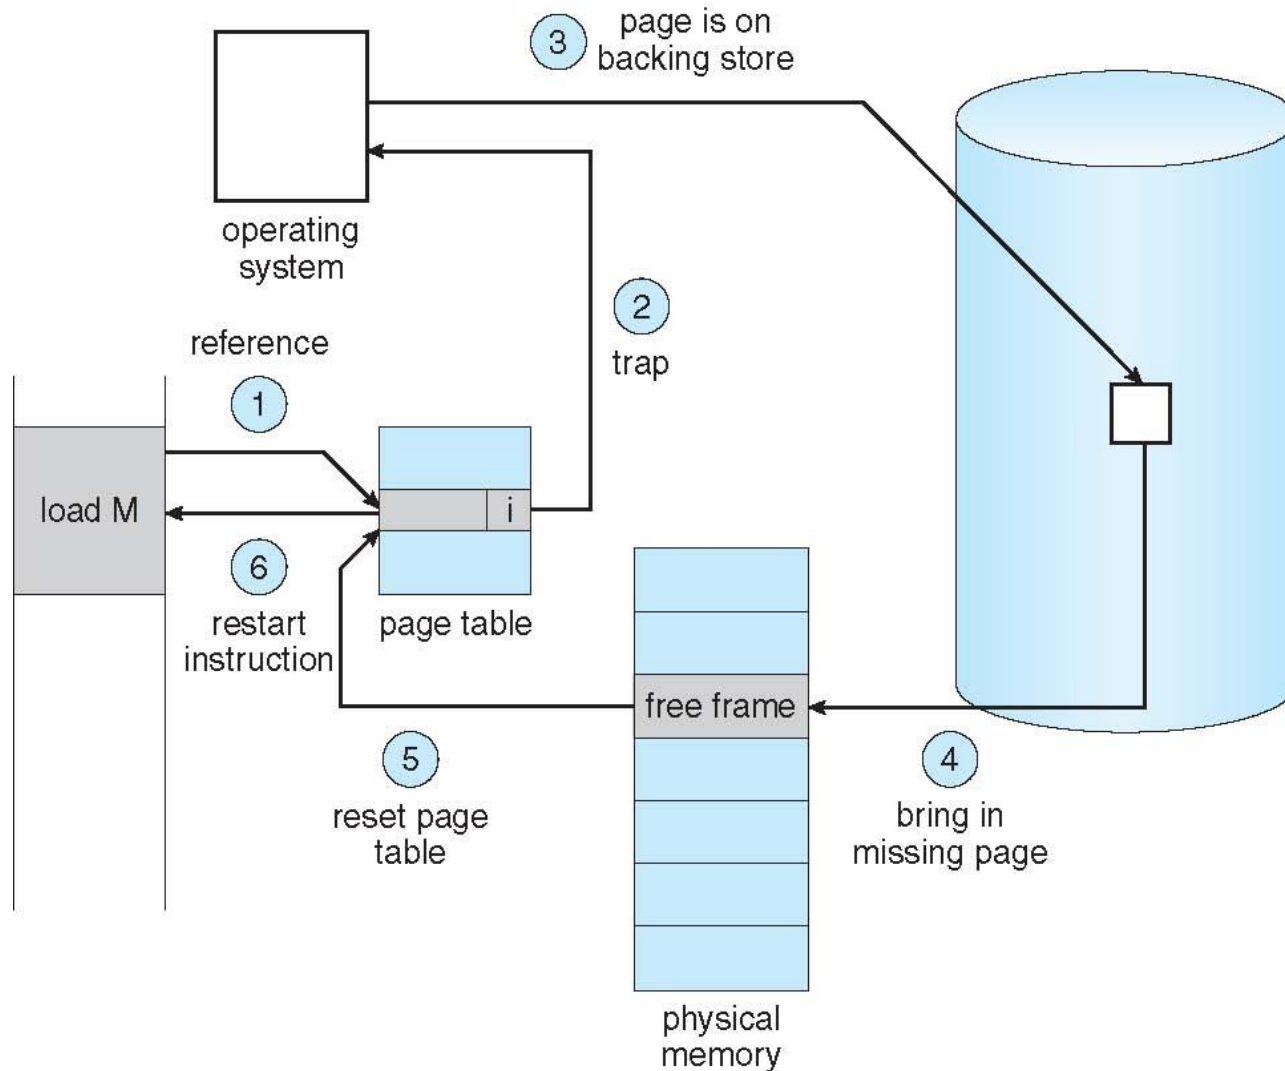
Demand Paging (cont.)



Page Fault

- First reference to a page will trap to OS
 - ➔ **page-fault trap**
- OS looks at the internal table (in PCB) to decide
 - Invalid reference ➔ abort
 - Just not in memory ➔ continue
- Get an empty frame
- Swap the page from disk (swap space) into the frame
- Reset page table, valid-invalid bit = 1
- **Restart instruction**

Page Fault Handling



Page Replacement

- If there is no free frame when a page fault occurs
 - Swap a frame to backing store
 - Swap a page from backing store into the frame
 - Different **page replacement algorithms** pick different frames for replacement

Demand Paging Performance

- **Effective Access Time (EAT):** $(1-p) \times ma + p \times pft$
 - p : page frame rate; ma : memory access time; pft : page fault time
- Example: $ma = 200ns$, $pft = 8ms$
 - $EAT = (1 - p) \times 200ns + p \times 8ms$
 $= 200ns + 7,999,800ns \times p$
- **Access time is proportional to the page fault rate**
 - If one access out of 1,000 causes a page fault, then
 $EAT = 8.2$ microseconds (slowdown by a factor of 40!)
 - For degradation less than 10%:
 $220 > 200 + 7,999,800 \times p$
 $p < 0.0000025$ (one access out of 399,990 to page fault)

Demand Paging Performance (cont.)

- Programs tend to have **locality** of reference
- Locality means program often accesses memory addresses that are close together
 - A single page fault can bring in 4KB memory content
 - Greatly reduce the occurrence of page fault
- Major components of page fault time (about 8 ms)
 - Serve the page-fault interrupt
 - **Read in the page from disk (most expensive)**
 - Restart the instruction
 - ➔ The 1st and 3rd can be reduced to several hundred instructions
 - ➔ The page switch time is close to 8 ms

Process Creation

Process and Virtual Memory

- **Demand Paging**

- Only bring in the page containing the first instruction

- **Copy-on-Write**

- The parent and the child process share the same frames initially, and frame-copy when a page is written


Copy-on-Write

- Allow both the parent and the child process to **share the same frames in memory**
- If either process modifies a frame, then a frame is **copied**
- Copy-on-write allows efficient process creation
- Free frames are allocated from a pool of **zeroed-out** frames (security reason)
 - The content of a frame is erased to 0

When a Child Process is Forked

```
#include <stdio.h>
void main( )
{
    int A;
    /* fork child process */
    A = fork( );

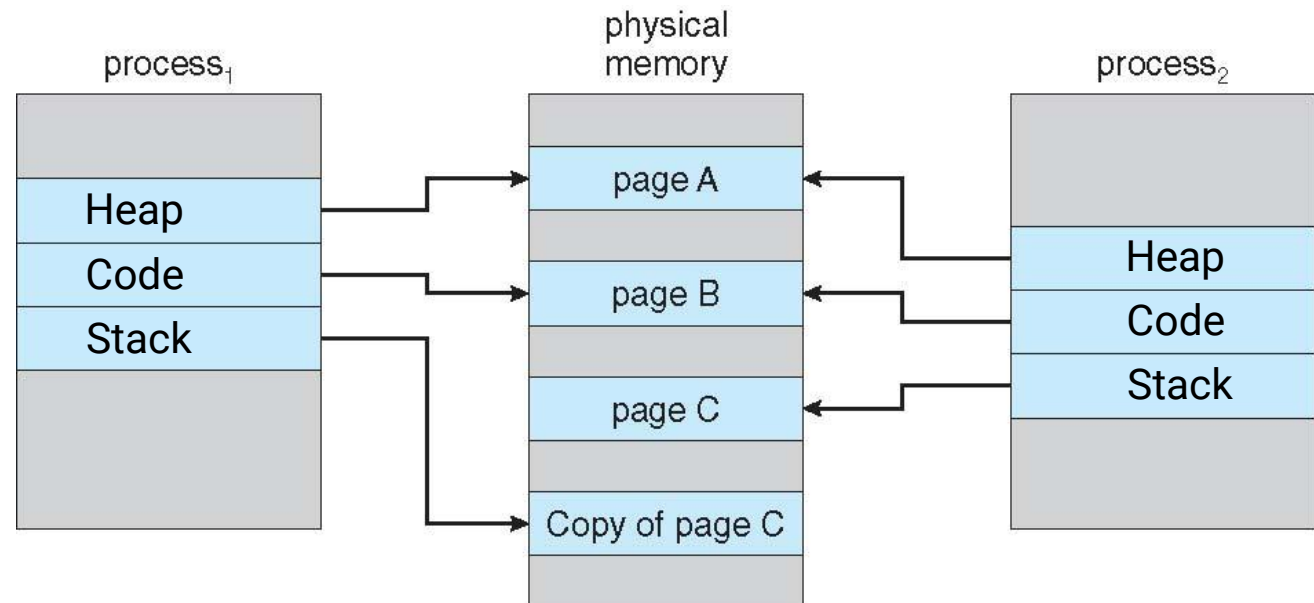
    if (A != 0) {
        /* parent process */
        int test1=0;
    }
    printf("process ends");
}
```



After a Page is Modified

```
#include <stdio.h>
void main( )
{
    int A;
    /* fork child process */
    A = fork( );

    if (A != 0) {
        /* parent process */
        int test1=0;
    }
    printf("process ends");
}
```



Page Replacement


Page Replacement Concept

- When a page fault occurs with no free frame
 - **Swap out a process**, freeing all its frames, or
 - **Page replacement**: find one not currently used and free it
- Solve two major problems for demand paging
 - **Frame-allocation algorithm**
 - Determine how many frames to be allocated to a process
 - **Page-replacement algorithm**
 - Select which frame to be replaced

Page Replacement (Page Fault) Steps

- Find the location of the desired page on disk
- Find a free frame
 - If there is a free frame, use it
 - If there is no free frame, use a page replacement algorithm to select a victim
- Read the desired page into the (newly) free frame
- Update the page and frame tables
- Restart the instruction

Page Replacement (Page Fault) Example



Page Replacement Algorithms


- **Goal: lowest page-fault rate**
- Evaluation: running against a string of memory references (**reference string**) and computing the number of page faults
- Reference string example:
1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5

Replacement Algorithms

- FIFO algorithm
- Optimal algorithm
- LRU algorithm
- Counting algorithm
 - LFU
 - MFU

First-In-First-Out (FIFO) Algorithm

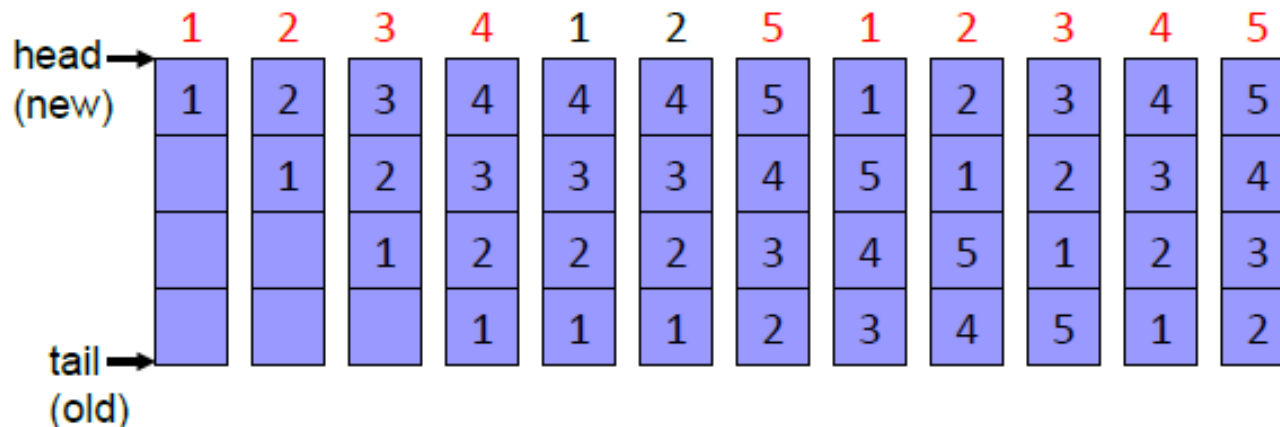
- The oldest page in a FIFO queue is replaced
 - Reference string: 1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5
 - 3 frames (available memory frames = 3)
- ➔ 9 page faults



(example borrowed from Prof. Jerry Chou's slides)


FIFO Illustrating Belady's Anomaly

- Does more allocated frames guarantee less page fault?
 - Reference string: 1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5
 - 4 frames (available memory frames = 4)
- **10** page faults !
- Belady's anomaly**
 - More allocated frames could result in more page faults




(example borrowed from Prof. Jerry Chou's slides)

FIFO Illustrating Belady's Anomaly (cont.)



Optimal (Belady) Algorithm

- Replace the page that will not be used for the **longest period of time**
 - Need future knowledge
- 4 frames: 1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5 → 6 page faults !
- In practice, we don't have future knowledge
 - Only used for reference and comparison




(example borrowed from Prof. Jerry Chou's slides)

LRU (Least Recently Used) Algorithm

- An approximation of optimal algorithm
 - **Looking backward** rather than forward
- It replaces the page that has **not been used for the longest period of time**
- It is often used, and is considered as **quite good**

LRU Algorithm Implementations

- Time stamp implementation
 - Page referenced: **time stamp** is copied into the counter
 - Replacement: remove the one with oldest counter
 - Linear search is required
- **Stack implementation**
 - Page referenced: move to top of the double-linked list
 - Replacement: remove the page at the bottom
 - 4 frames: 1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5 → 8 page faults !



(example borrowed from Prof. Jerry Chou's slides)

Stack Algorithm

- A **property** of algorithms
- **Stack algorithm**
 - The set of pages in memory for n frames is always a subset of the set of pages that would be in memory with $n+1$ frames
- **Stack algorithms do not suffer from Belady's anomaly**
- Both **optimal** algorithm and **LRU** algorithm are stack algorithm

Counting Algorithm

- **LFU (Least Frequently Used) Algorithm**
 - Keep a counter for each page
 - Idea: an actively used page should have a large reference count
- **MFU (Most Frequently Used) Algorithm**
 - Idea: the page with the smallest count was probably just brought in and has yet to be used
- Both counting algorithms are not common
 - Implementation is expensive
 - Do not approximate OPT algorithm very well

Allocation of Frames

Frame Allocation

- **Fixed allocation**

- **Equal allocation**

- Example: 100 frames, 5 processes → 20 frame / process

- **Proportional allocation**

- Allocate frames according to the size of the process

- **Priority allocation**

- Using proportional allocation based on priority instead of size
 - If process P generates a page fault
 - Select one of its frame for replacement
 - Select from a process with lower priority for replacement

Frame Allocation (cont.)

- **Local allocation**

- Each process select from its own set of allocated frames


- **Global allocation**

- Process selects a replacement frame from the set of all frames
- One process can take away a frame of another process
 - E.g., allow a high-priority process to take frames from a low-priority process
- Good system performance and thus is commonly used
- **Need to prevent thrashing**

Thrashing

Definition of Thrashing

- If a process does not have enough **frames**
 - The process does not have # frames it needs to support pages in active use
 - ➔ Very high paging activity
- **A process is thrashing if it is spending more time paging than executing**



Thrashing

- Performance problem caused by thrashing (assume global replacement is used)
 - Processes queued for I/O to swap (page fault)
 - ➔ Low CPU utilization
 - ➔ OS increases the degree of multi-programming
 - ➔ New processes take frames from old processes
 - ➔ More page faults and thus more I/O
 - ➔ CPU utilization drops even further
- To prevent thrashing, must provide enough frames for each process
 - **Working-set model**
 - **Page-fault frequency**

Working-Set Model


- **Locality**: a set of pages that are actively used together
- Locality model: as a process executes, it moves from locality to locality
 - Program structure (subroutine, loop, stack)
 - Data structure (array, table)
- Working-set model (based on locality model)
 - Working-set **window**: a parameter Δ (delta)
 - Working-set: set of pages in most recent Δ page references (an approximation locality)

Working-Set Example


- If Δ (delta) = 10

page reference table

... 2 6 1 5 7 7 7 7 5 1 6 2 3 4 1 2 3 4 4 4 3 4 3 4 4 4 4 1 3 2 3 4 4 4 3 4 4 4 ...



$$WS(t_1) = \{1, 2, 5, 6, 7\}$$



$$WS(t_2) = \{3, 4\}$$


Working-Set Model (cont.)

- Prevent thrashing using the working-set size
 - WSS: working-set size for process i
 - $D = \sum WSS_i$ (total demand frames)
 - if $D > m$ (available frames) \rightarrow thrashing
 - The OS monitors the WSS_i of each process and allocates to the process enough frames
 - if $D \ll m$, increase degree of MP
 - If $D > m$, suspend a process
- Prevent thrashing while keeping the degree of multiprogramming as high as possible
- Optimize CPU utilization
- **However, too expensive for tracking**

Page Fault Frequency Scheme

- **Page fault frequency** directly measures and controls the page-fault rate to prevent thrashing
 - Establish **upper** and **lower** bounds on the desired page-fault rate of a process
 - If page fault rate exceeds the upper limit
 - Allocate another frame to the process
 - If page rate falls below the lower limit
 - Remove a frame from the process

Page Fault Frequency Scheme (cont.)



Objective Review

- Define virtual memory and describe its benefits
- Illustrate how pages are loaded into memory using demand paging
- Apply the FIFO, optimal, and LRU page-replacement algorithms
- Describe the working set of a process and explain how it is related to program locality