



# Final Exam Review

Operating Systems  
Yu-Ting Wu

1

Operating Systems 2022

## Scope of the Final Exam

- Chapter 6/7: Synchronization
- Chapter 8: Deadlocks
- Chapter 9: Main memory
- Chapter 10: Virtual memory

2

2

Operating Systems 2022

## Synchronization Overview

3

3

Operating Systems 2022

## The Goals of Synchronization

- **Concurrent access** to **shared data** may result in **data inconsistency**
  - Prevent **race condition**: The final value of the shared data depends upon which process finishes last
  - Solve the **critical section problem**
- Maintaining data consistency requires mechanism to ensure the **orderly execution** of cooperating processes

4

4

## The Critical-Section Problem

- A **protocol** for processes to cooperate
- **General code section structure**
  - Only one process can be in a critical section

do {

**entry section**



get **entry permission**

**critical section**



modified **shared data**

**exit section**



release **entry permission**

*remainder section*

} while (1);

5

5

## Critical-Section Requirements

- **Mutual exclusion**
  - If a process *P* is executing in its critical section (CS), no other processes can be executing in their CS
- **Progress**
  - If no process is executing in its CS and there exist some processes that wish to enter their CS, there processes cannot be postponed indefinitely
- **Bounded Waiting**
  - A **bound** must exist on the number of times that other processes are allowed to enter their CS after a process has made a request to enter its CS
- Goal: design **entry** and **exit section** to satisfy the above requirement

6

6

## CS Solutions and Synchronization Tools

- **Software solution**
- **Synchronization hardware**
- **Semaphore**
- **Monitor**

7

7

## Synchronization Software Solution

8

8

## Algorithm for Two Processes

- Only 2 processes  $P_0$  and  $P_1$
- Shared variables
  - int *turn*; // initially *turn* = 0
  - turn* == *i* →  $P_i$  can enter its critical section

```

/* Process 0 */
do {
    while (turn != 0); — entry section
    critical section
    turn = 1; — exit section
    remainder section
} while (1);

/* Process 1 */
do {
    while (turn != 1); — entry section
    critical section
    turn = 0; — exit section
    remainder section
} while (1);

```

mutual exclusion? **Y**   progress? **N**   bounded-wait? **Y**

9

## Peterson's Solution for Two Processes

- Shared variables
  - int *turn*; // initially *turn* = 0
  - turn* == *i* →  $P_i$  can enter its critical section
  - boolean *flag*[2]; // initially *flag*[0] = *flag*[1] = false
  - flag*[*i*] == true →  $P_i$  is ready to enter its critical section

```

/* Process i */
do {
    flag[i] = true;
    turn = i; — entry section
    while (flag[j] && turn == j);
    critical section
    flag[i] = false; — exit section
    remainder section
} while (1);

```

mutual exclusion? **Y**   progress? **Y**   bounded-wait? **Y**

10

## Bakery Algorithm for $n$ Processes

- Before entering its CS, each process receives a **number (#)**
- Holder of the smallest # enters CS**
- The numbering scheme always generates # in **non-decreasing order**; i.e., 1, 2, 3, 3, 4, 5, 5, 5 ...
- If processes  $P_i$  and  $P_j$  receive the same #, if  $i < j$ , then  $P_i$  is served first
- Notation:
  - $(a, b) < (c, d)$  if
    - $a < c$  or
    - $a == c$  &&  $b < d$

11

## Bakery Algorithm for $n$ Processes (cont.)

```

// Process i:
do {
    get ticket → choosing[i] = true;
    num[i] = max (num[0], num[1], ...num[n-1]) + 1;
    choosing[i] = false;
    FCFS → for (j = 0; j < n; j++) {
        while (choosing[j]);
        while ((num[j] != 0) && ((num[j], j) < (num[i], i)));
    }
    release ticket → critical section
    num[i] = 0;
    remainder section
}

```

**Bounded waiting** because processes enter CS on a **first come, first served** basis

12

## Condition Variables

- **Condition variables (CV)** represent some condition that a **thread** can
  - **Wait on**, until the condition occurs; or
  - **Notify** other waiting threads that the condition has occurred
- Three operations on condition variables
  - **wait()** – block until another thread calls **signal()** or **broadcast()** on the CV  
`pthread_cond_wait(&theCV, &someLock)`
  - **signal()** – wake up one thread waiting on the CV  
`pthread_cond_signal(&theCV)`
  - **broadcast()** – wake up all threads waiting on the CV  
`pthread_cond_broadcast(&theCV)`

13

13

## Synchronization Hardware Solution

14

14

## Atomic TestAndSet()

```
bool TestAndSet (bool &lock) {
    bool value = lock;
    lock = true;
    return value;
}
```

execute atomically:  
**return the value of "lock" and  
 set "lock" to true**

```
shared data: bool lock; // initially lock = false
// P0
do {
    while (TestAndSet (lock));
    critical section
    lock = false;
    remainder section
} while (1);
```

```
// P1
do {
    while (TestAndSet (lock));
    critical section
    lock = false;
    remainder section
} while (1);
```

mutual exclusion? **Y**   progress? **Y**   bounded-wait? **N**

15

15

## Atomic Swap()

Enter CS if lock == false

```
shared data: bool lock; // initially lock = false
// P0
do {
    key0 = true;
    while (key0 == true)
        Swap(lock, key0);
    critical section
    lock = false;
    remainder section
} while (1);
```

```
// P1
do {
    key1 = true;
    while (key1 == true)
        Swap(lock, key1);
    critical section
    lock = false;
    remainder section
} while (1);
```

mutual exclusion? **Y**   progress? **Y**   bounded-wait? **N**

16

16

## Synchronization Semaphores

17

17

## Semaphores


- A tool to generalize the synchronization problem
- More specifically
  - A **record** of **how many units of a particular resource** is available
    - If # record = 1 → binary semaphore, **mutex lock**
    - If # record > 1 → counting semaphore
  - Accessed only through 2 **atomic** operations: **wait** & **signal**
- Implementation
  - Spinlock
  - Data structure with queue

18

18

## Semaphore Example

- Initially, all semaphores are 0
- Begin
  - P1: S1; signal(a); signal(b);
  - P2: wait(a); S2; signal(c);
  - P3: wait(b); S3; signal(d);
  - P4: wait(c); S4; signal(e); signal(f);
  - P5: wait(e); S5; signal(g);
  - P6: wait(f); wait(d); S6; signal(h);
  - P7: wait(g); wait(h); S7;
- End



19

19

## Synchronization Monitors

20

20

## Monitor

- A **high-level language** construct
- The representation of a monitor type consists of
  - Declaration of **variables** whose values define the state of an instance of the type
  - **Procedures/functions** that implement operations on the type
- The monitor type is similar to a **class in O.O language**
  - A procedure within a monitor can access only **local variable** and the **formal parameters**
  - The local variables of a monitor can be used only by the local procedures
- But, the monitor ensures that **only one process at a time can be active within the monitor**

21

21

## Monitor Condition Variables

- To allow a process to **wait within** the monitor, a condition variable must be declared, as  
**condition x, y;**
- Condition variable can only be used with the operations **wait()** and **signal()**

```
x.wait();
```

 means that the process invoking this operation is suspended until another process invokes it
 

```
x.signal();
```


 resumes exactly one suspended process. If no suspended, then the signal operation **has no effects** (in contrast, signal always change the state of a semaphore)

22

22

## Dining-Philosophers Problem

- **5 persons** sitting on 5 chairs with **5 chopsticks**
- A person is either thinking or eating
  - **thinking**: no interaction with the rest 4 persons
  - **eating**: need 2 chopsticks at hand
  - a person **picks up 1 chopstick at a time**
  - done eating: **put down both chopsticks**



23

23

## Dining Philosophers Example

```
monitor dp {
    enum { thinking, hungry, eating } state[5]; // current state
    condition self[5]; // delay eating if can't obtain chopsticks

    void pickup(int i); // pickup chopsticks
    void putdown(int i); // putdown chopsticks
    void test(int i); // try to eat

    void init() {
        for (int i = 0; i < 5; i++)
            state[i] = thinking;
    }
}
```

24

24

## Deadlocks Overview

25

25

## Deadlock Problem

- A set of blocked processes each **holding** some resources and **waiting** to acquire a resource held by another process in the set
- Example:
  - 2 processes and semaphores A and B
    - $P_1$  (hold B, wait A): **wait** (A), signal (B)
    - $P_2$  (hold A, wait B): **wait** (B), signal (A)
- Example:
  - Dining philosophers' problem

26

26

## Necessary Conditions

- **Mutual exclusion**
  - Only 1 process at a time can use a resource
- **Hold and wait**
  - A process holding some resources and is waiting for another resource
- **No preemption**
  - A resource can be only released by a process **voluntarily**
- **Circular wait**
  - There exists a set  $\{P_0, P_1, \dots, P_n\}$  of waiting processes such that  $P_0 \rightarrow P_1 \rightarrow P_2 \rightarrow \dots \rightarrow P_n \rightarrow P_0$


27

27

## Resource-Allocation Graph

- 3 processes,  $P_1 \sim P_3$
- 4 resources,  $R_1 \sim R_4$ 
  - $R_1$  and  $R_3$  each has one instance
  - $R_2$  has two instances
  - $R_4$  has three instances
- **Request edges**
  - $P_1 \rightarrow R_1$ :  $P_1$  requests  $R_1$
- **Assignment edges**
  - $R_2 \rightarrow P_1$ : one instance of  $R_2$  is allocated to  $P_1$

→  $P_1$  is **holding on** an instance of  $R_2$  and **waiting for** an instance of  $R_1$




28

28

## Resource-Allocation Graph w/ Deadlock


- If the graph contains a **cycle**, a deadlock **may** exist
- In the example
  - $P_1$  is waiting for  $P_2$
  - $P_2$  is waiting for  $P_3$
  - $P_1$  is also waiting for  $P_3$
  - Since  $P_3$  is waiting for  $P_1$  or  $P_2$ , and they both waiting for  $P_3$
  - **Deadlock !**



29

## RA Graph w/ Cycle but NO Deadlock

- If the graph contains a **cycle**, a deadlock **may** exist
- In the example
  - $P_1$  is waiting for  $P_2$  or  $P_3$
  - $P_3$  is waiting for  $P_1$  or  $P_4$
  - Since  $P_2$  and  $P_4$  wait for no one
  - **No Deadlock between  $P_1$  and  $P_3$**



30

## Handling Deadlocks

- Ensure the system will **never** enter a deadlock state
  - **Deadlock prevention**: ensure that at least one of the **4 necessary conditions** cannot hold
  - **Deadlock avoidance**: **dynamically** examines the resource-allocation state before allocation
- Allow to **enter a deadlock state** and then **recover**
  - **Deadlock detection**
  - **Deadlock recovery**
- **Ignore the problem** and pretend that deadlocks never occur in the system
  - **Used by most operating systems, including UNIX**

31

## Deadlocks Prevention

32



## Deadlock Prevention

- **Mutual exclusion (ME):** do not require ME on sharable resources
  - E.g. there is no need to ensure ME on read-only files
  - However, some resources are not shareable (e.g. printer)
- **Hold and wait:**
  - When a process requests a resource, it does not hold any resource
  - Pre-allocate all resources before executing
  - Resource utilization is low; starvation is possible

33

33

## Deadlock Prevention (cont.)

- **No preemption:**
  - When a process is waiting on a resource, all its holding resources are preempted
    - E.g.  $P_1$  request  $R_1$ , which is allocated to  $P_2$ , which in turn is waiting on  $R_2$  ( $P_1 \rightarrow R_1 \rightarrow P_2 \rightarrow R_2$ )
    - **$R_1$  can be preempted and reallocated to  $P_1$**
- **Applied to resources whose states can be easily saved and restored later**
  - E.g. CPU registers and memory
- It cannot easily be applied to other resources
  - E.g. printers and tape drives

34

34

## Deadlock Prevention (cont.)

- **Circular wait:**
  - Impose a **total ordering** of all resource types
  - A process requests resources in an increasing order
    - Let  $R = \{R_0, R_1, \dots, R_n\}$  be the set of resource types
    - **When request  $R_k$ , should release all  $R_i, i \geq k$**
- Example
  - $F(\text{disk drive}) = 5, F(\text{printer}) = 12$
  - A process must request disk drive before printer
- Proof: counter-example does not exist
  - $P_0(R_0) \rightarrow R_1, P_1(R_1) \rightarrow R_2, \dots, P_n(R_n) \rightarrow R_0$  ←  $P_n$  holds on  $R_n$ , waiting for  $R_0$
  - Conflict:  $R_0 < R_1 < R_2 < \dots < R_n < R_0$

35

35


## Deadlocks Avoidance

36

36

## Resource-Allocation Graph Algorithm

- **Request edges**
  - $P_i \rightarrow R_j$ :  $P_i$  is waiting for resource  $R_j$
- **Assignment edges**
  - $R_j \rightarrow P_i$ : Resource  $R_j$  is allocated and held by  $P_i$
- **Claim edge**
  - Process  $P_i$  may request  $R_j$  in the future
- **Claim edge** converts to **request edge**
  - When a resource is requested by process
- **Assignment edge** converts back to a **claim edge**
  - When a resource is released by a process




37

37

## Safe / Unsafe State

- **Safe state**: a system is in a safe state if there exists a **sequence of allocations** to satisfy requests by all processes
  - This sequence of allocations is called **safe sequence**
- **Safe state**  $\rightarrow$  **no deadlock**
- **Unsafe state**  $\rightarrow$  **possibility of deadlock**
- **Deadlock avoidance**  $\rightarrow$  **ensure that a system never enters an unsafe state**



38

38

## Banker's Algorithm

- Use for **multiple instances** of each resource type
- **Banker's Algorithm**
  - Use a general safety algorithm to **pre-determine** if any **safe sequence** exists after allocation
  - **Only proceed the allocation if safe sequence exists**
- **Safety algorithm**
  1. Assume processes need **maximum** resources
  2. Find a process that can be satisfied by free resources
  3. Free the resource usage of the process
  4. Repeat to step 2 until all processes are satisfied

39

39

## Banker's Algorithm Example

- Total instances: A: 10, B: 5, C: 7
- Available instances: A: 3, B: 3, C: 2

	Max			Allocation			Need (Max - Alloc.)		
	A	B	C	A	B	C	A	B	C
P <sub>0</sub>	7	5	3	0	1	0	7	4	3
P <sub>1</sub>	3	2	2	2	0	0	1	2	2
P <sub>2</sub>	9	0	2	3	0	2	6	0	0
P <sub>3</sub>	2	2	2	2	1	1	0	1	1
P <sub>4</sub>	4	3	3	0	0	2	4	3	1

- Safe sequence ?

40

40

## Deadlocks Detection

41

41

## Multiple Instance for Each Resource Type

- Total instances: A: 7, B: 2, C: 6
- Available instances: A: 0, B: 0, C: 0

	Allocation			Request		
	A	B	C	A	B	C
P <sub>0</sub>	0	1	0	0	0	0
P <sub>1</sub>	2	0	0	2	0	2
P <sub>2</sub>	3	0	3	0	0	0
P <sub>3</sub>	2	1	1	1	0	0
P <sub>4</sub>	0	0	2	0	0	2

- The system is in a safe state → <P<sub>0</sub>, P<sub>2</sub>, P<sub>3</sub>, P<sub>1</sub>, P<sub>4</sub>>  
→ No deadlock

42

42

## Multiple Instance for Each Resource Type

- Total instances: A: 7, B: 2, C: 6
- Available instances: A: 0, B: 0, C: 0

	Allocation			Request		
	A	B	C	A	B	C
P <sub>0</sub>	0	1	0	0	0	0
P <sub>1</sub>	2	0	0	2	0	2
P <sub>2</sub>	3	0	3	0	0	1
P <sub>3</sub>	2	1	1	1	0	0
P <sub>4</sub>	0	0	2	0	0	2

- If P<sub>2</sub> requests (0, 0, 1) → no safe sequence can be found  
→ The system is deadlocked

43

43

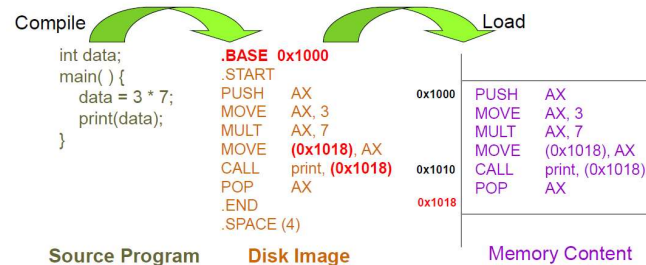
## Main Memory Address Binding

44

44

## Address Binding: Compile Time


- Program is written as symbolic code
- Compiler translates symbolic code into **absolute code**
- If starting location changes → **recompile**



45

## Address Binding: Load Time


- Compiler translates symbolic code into **relocatable code**
- **Relocatable code**
  - Machine language that can be run from any memory location
  - If starting location changes → **reload the code**



46

## Address Binding: Execution Time

- Compiler translates symbolic code into **logical-address (i.e. virtual-address) code**
- **Special hardware (i.e. MMU)** is needed for this scheme
- Most general-purpose OS use this method



47

## Logical v.s. Physical Address

- **Logical address** – generated by CPU
  - a.k.a virtual address
- **Physical address** – seen by the memory module
- **Compile-time and load-time address binding**
  - Logical address = physical address
- **Execution-time address binding**
  - Logical address ≠ physical address

48

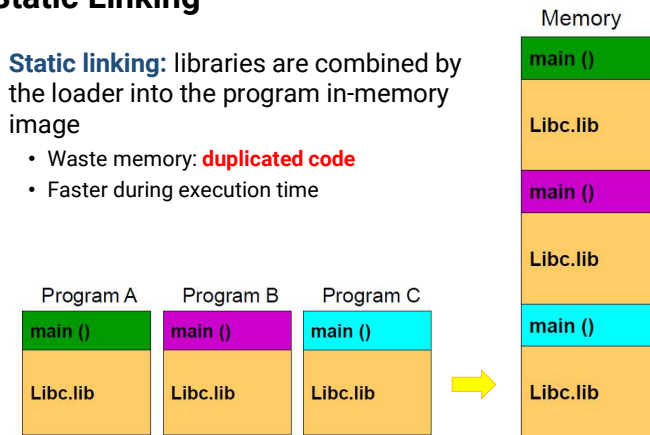
## Main Memory Static / Dynamic Linking

49

49

## Static Linking

- **Static linking:** libraries are combined by the loader into the program in-memory image
  - Waste memory: **duplicated code**
  - Faster during execution time




50

50

## Dynamic Linking

- **Dynamic linking:** linking postponed until execution time
  - **Only one code copy** in memory and shared by everyone
  - A **stub** is included in the program in-memory image for each lib reference
  - **Stub call**
    - check if the referred lib is in memory
    - if not, load the lib
    - execute the lib



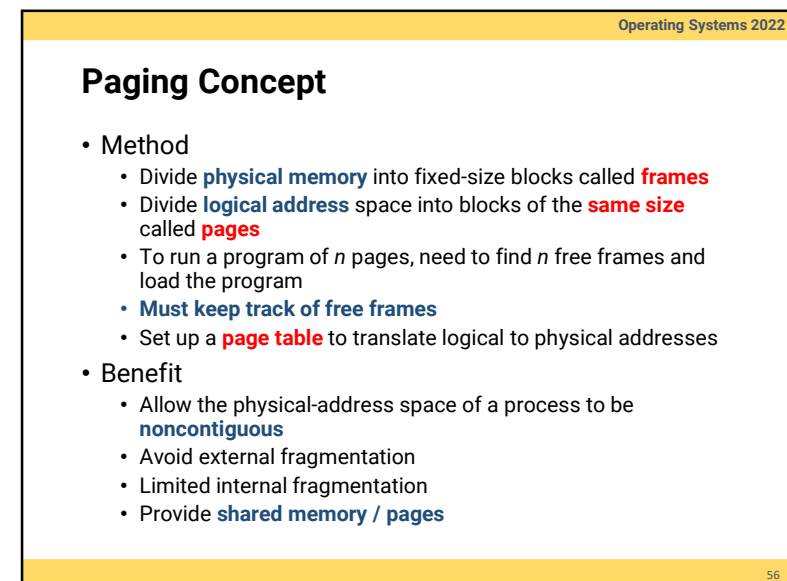
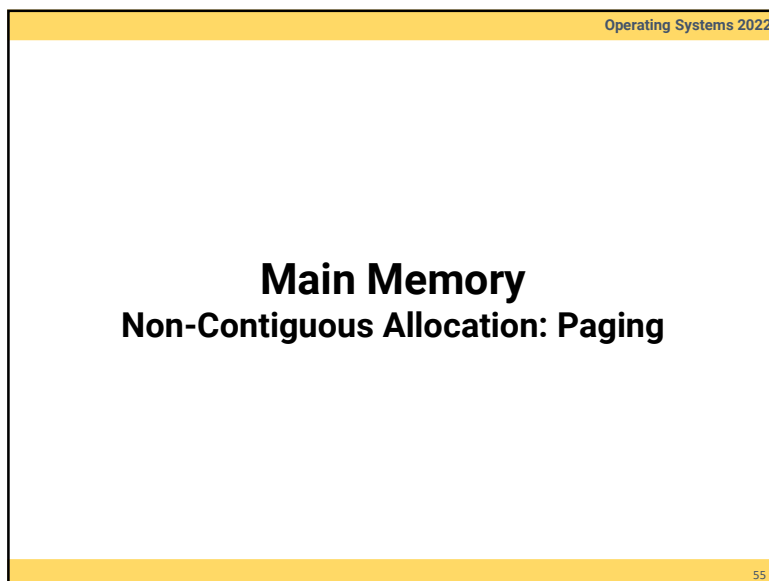
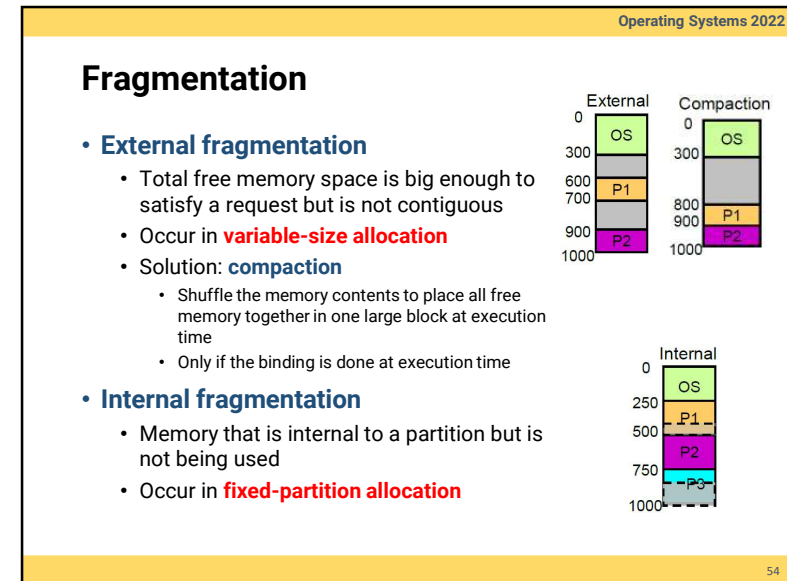
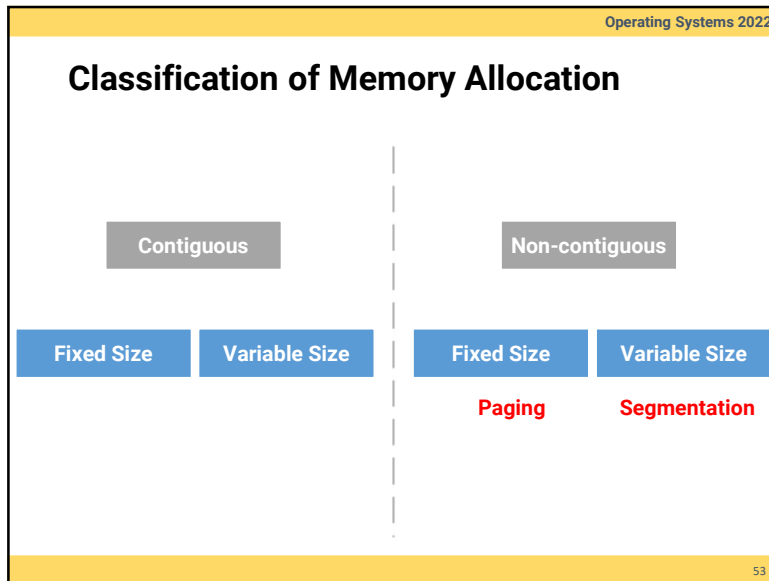
51

51

## Main Memory Contiguous Allocation

52

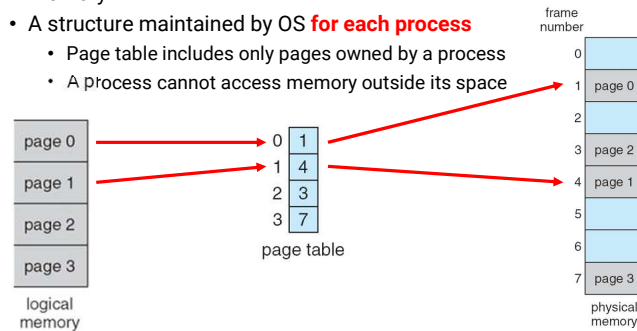
52



## Paging Example

### • Page table

- Each entry maps to the **base address of a page** in physical memory
- A structure maintained by OS **for each process**
  - Page table includes only pages owned by a process
  - A process cannot access memory outside its space



57

57

## Address Translation Scheme

- Logical address is divided into two parts

### • Page number (p)

- Used as an **index into a page table** which contains **base address** of each page in physical memory
- $N$  bits means a process can allocate at most  $2^N$  pages  
 $\rightarrow 2^N \times \text{page size} = \text{memory size}$

### • Page offset (d)

- Combined with base address to define the physical memory address that is sent to the memory unit
- $N$  bits means the **page size is  $2^N$**

- Physical address = page base address + page offset**


58

58

## Address Translation Architecture

- If page size is 1KB ( $2^{10}$ ) and page 3 maps to frame 5
- Given 13 bits logical address ( $p = 3, d = 20$ ), what is the physical address?

- $5 * (1\text{KB}) + 20 = 101000000000 + 0000010100 = 1010000010100$



59

59

## Address Translation

- Total number of pages does not need to be the same as the total number of frames
  - Total # pages** determines the logical memory size of a process
  - Total # frames** depending on the size of physical memory
- E.g.: Given 32 bits logical address, 36 bits physical address, and 4KB page size, what does it mean?
  - Number of bits for page offset: 4KB page size =  $2^{12}$  bytes  $\rightarrow 12$
  - Number of bits for page number:  $2^{20}$  pages  $\rightarrow 20$  bits
  - Page table size:  $2^{32} / 2^{12} = 2^{20}$  entries
  - Max program memory:  $2^{32} = 4\text{GB}$
  - Number of bits for frame number:  $2^{24}$  frames  $\rightarrow 24$  bits
  - Total physical memory size:  $2^{36} = 64\text{GB}$

60

60

## Page / Frame Size

- The page (frame) size is defined by hardware
  - Typically, a power of 2
  - Ranging from 512 bytes to 16 MB / page
  - 4KB / 8KB page is commonly used
- Internal fragmentation?
  - Larger page size → More space waste
- But **page sizes cannot be too small**
  - Memory, process, and data sets have become larger
  - Need to keep page table small
  - Fewer access means better I/O performance

61

61

## Implementation of Page Table

- Page table is kept **in memory**
- **Page-table base register (PTBR)**
  - The **physical memory address** of the page table
  - The PTBR value is stored in **PCB** (Process Control Block)
  - Changing the value of PTBR during the context switch
- With PTBR, each memory reference results in **2 memory reads**
  - One for the page table and one for the real address
- The 2-access problem can be solved by
  - **Translate Look-aside Buffers (TLB)** (HW) which is implemented by **Associative memory** (HW)

62

62

## Effective Memory-Access Time

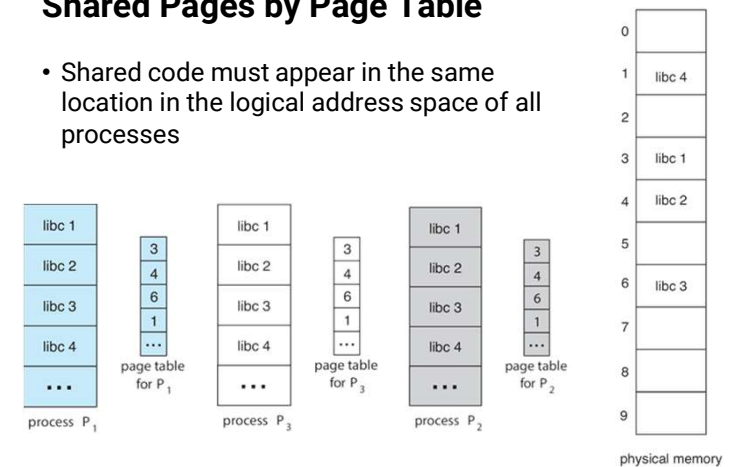
- 20 ns for TLB search
- 100 ns for memory access
- Effective Memory-Access Time (**EMAT**)
  - 70% TLB hit-ratio:
    - $EMAT = 0.70 \times (20 + 100) + (1 - 0.70) \times (20 + 100 + 100) = 150 \text{ ns}$
  - 98% TLB hit-ratio:
    - $EMAT = 0.98 \times 120 + 0.02 \times 220 = 122 \text{ ns}$

63

63

## Shared Pages by Page Table

- Shared code must appear in the same location in the logical address space of all processes



64

64



## Page Table Memory Structure


- Page table could be huge and difficult to be loaded
  - 4GB ( $2^{32}$ ) logical address space with 4KB ( $2^{12}$ ) page
    - 1 million ( $2^{20}$ ) page table entry
  - Assume each entry need 4 bytes (32 bits)
    - Total size = **4MB**
  - Need to break it into several smaller page tables, better within a single page size (i.e. 4KB)
  - Or reduce the total size of page table
- Solutions
  - Hierarchical paging
  - Hash page table
  - Inverted page table

65

65

## Hierarchical Paging


- Break up the logical address space into **multiple page tables**
  - Paged the page table
  - i.e.  $n$ -level page table
- Two-level paging (32-bit address with 4KB ( $2^{12}$ ) page size)
  - 12-bit offset ( $d$ ) → 4KB ( $2^{12}$ ) page size
  - 10-bit **outer** page number → 1K ( $2^{10}$ ) page table entries
  - 10-bit **inner** page number → 1K ( $2^{10}$ ) page table entries
  - 3 memory accesses



66

66

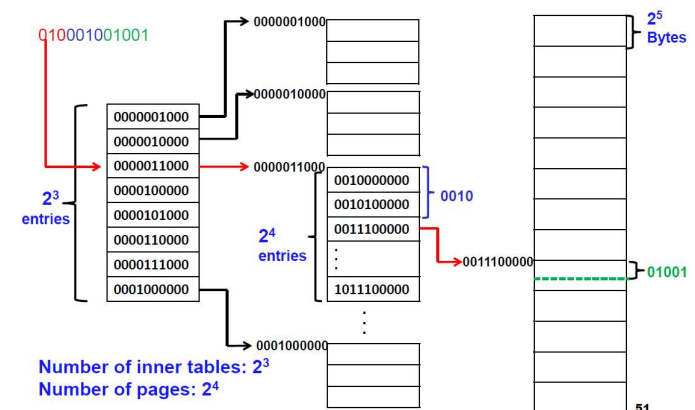
## Two-Level Page Table Example



67

67

## Two-Level Address Translation Example




68

68

## Hashed Page Table

- **Commonly-used for address > 32 bits**
- Virtual page number is hashed into a hash table
- The size of the hash table varies
  - Larger hash table → smaller chains in each entry
- Each entry in the hashed table contains
  - (Virtual Page Number, Frame Number, Next Pointer)
  - Pointers waste memory
  - Traverse linked list waste time and cause additional memory references



69

69


## Inverted Page Table (Frame Table)

- Maintains **no** page table for each process
- Maintains a **frame table** for the whole memory
  - One entry for each real frame of memory
- Each entry in the hashed table contains
  - (PID, Page Number)
- Eliminate the memory needed for page tables but **increase memory access time**
  - Each access needs to search the whole frame table
  - Solution: use hashing for the frame table
- **Hard to support shared page / memory**

70

70

## Inverted Page Table Address Translation



71

71


## Main Memory Non-Contiguous Allocation: Segmentation

72

72


## Segmentation

- Memory-management scheme that supports **user view of memory**
- A program is a collection of segments
- A segment is a logical unit such as
  - Main program
  - Function
  - Object
  - Local/global variables
  - Stack
  - Symbol table
  - Arrays



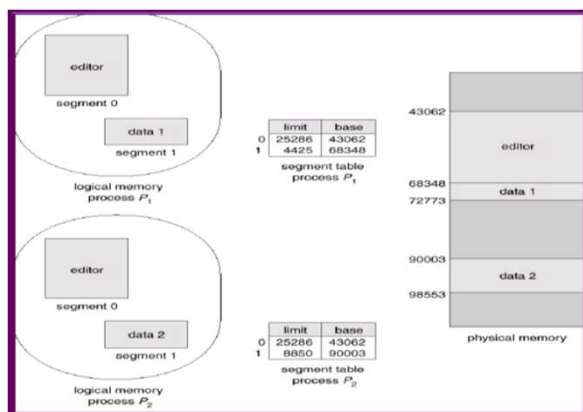
73

## Logical View of Segmentation



74

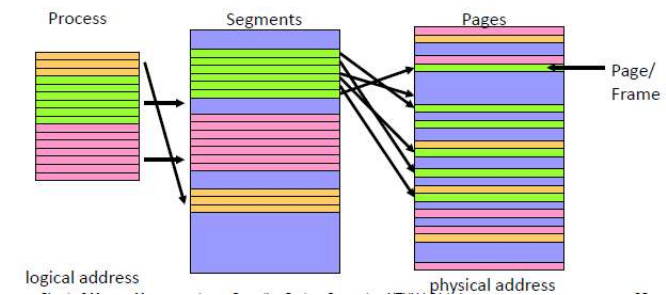
## Sharing of Segments



75

## Basic Concept

- Apply **segmentation** in **logical** address space
- Apply **paging** in **physical** address space



76

## Virtual Memory Overview

77

77


## Virtual Memory

- Separation of user logical memory from physical memory
  - To run an **extremely large process**
    - Logical address space can be much larger than physical address space
  - To increase **CPU/resource utilization**
    - Higher degree of multiprogramming degree
  - To simplify programming (compiler) tasks
    - Free programmer from memory limitation
  - To **launch** programs **faster**
    - Less I/O would be needed to load or swap
- Can be implemented via
  - **Demand paging**
  - Demand segmentation (more complicated due to variable sizes)

78

78

## Page Fault Handling



79

79

## Virtual Memory Process Creation

80

80

## Process and Virtual Memory

- **Demand Paging**

- Only bring in the page containing the first instruction

- **Copy-on-Write**

- The parent and the child process share the same frames initially, and frame-copy when a page is written

81

81

## Copy-on-Write

- Allow both the parent and the child process to **share the same frames in memory**
- If either process modifies a frame, then a frame is **copied**
- Copy-on-write allows efficient process creation
- Free frames are allocated from a pool of **zeroed-out** frames (security reason)
  - The content of a frame is erased to 0


82

82

## When a Child Process is Forked

```
#include <stdio.h>
void main( )
{
    int A;
    /* fork child process */
    A = fork( );

    if (A != 0) {
        /* parent process */
        int test1=0;
    }
    printf("process ends");
}
```




83

83

## After a Page is Modified

```
#include <stdio.h>
void main( )
{
    int A;
    /* fork child process */
    A = fork( );

    if (A != 0) {
        /* parent process */
        int test1=0;
    }
    printf("process ends");
}
```



84

84

## Virtual Memory Page Replacement

85

85

## Page Replacement Algorithms

- **Goal: lowest page-fault rate**
- Evaluation: running against a string of memory references (**reference string**) and computing the number of page faults
- Reference string example:  
1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5

86

86

## Replacement Algorithms


- FIFO algorithm
- Optimal algorithm
- LRU algorithm
- Counting algorithm
  - LFU
  - MFU

87

87

## First-In-First-Out (FIFO) Algorithm

- The oldest page in a FIFO queue is replaced
- Reference string: 1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5
- 3 frames (available memory frames = 3)  
→ 9 page faults



(example borrowed from Prof. Jerry Chou's slides)

88


88

## FIFO Illustrating Belady's Anomaly

- Does more allocated frames guarantee less page fault?
    - Reference string: 1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5
    - 4 frames (available memory frames = 4)
- **10** page faults !

### • Belady's anomaly

- More allocated frames could result in more page faults




(example borrowed from Prof. Jerry Chou's slides)

89

89

## Optimal (Belady) Algorithm

- Replace the page that will not be used for the **longest period of time**
  - Need future knowledge
- 4 frames: 1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5 → 6 page faults !
- In practice, we don't have future knowledge
  - Only used for reference and comparison




(example borrowed from Prof. Jerry Chou's slides)

90

90

## LRU Algorithm Implementations

- Time stamp implementation
  - Page referenced: **time stamp** is copied into the counter
  - Replacement: remove the one with oldest counter
    - Linear search is required
- Stack implementation
  - Page referenced: move to top of the double-linked list
  - Replacement: remove the page at the bottom
  - 4 frames: 1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5 → 8 page faults !



(example borrowed from Prof. Jerry Chou's slides)

91

91


## Virtual Memory Thrashing

92

92

## Definition of Thrashing

- If a process does not have enough **frames**
  - The process does not have # frames it needs to support pages in active use
    - Very high paging activity
- **A process is thrashing if it is spending more time paging than executing**



93

93

## Thrashing

- Performance problem caused by thrashing (assume global replacement is used)
  - Processes queued for I/O to swap (page fault)
    - Low CPU utilization
    - OS increases the degree of multi-programming
    - New processes take frames from old processes
    - More page faults and thus more I/O
    - CPU utilization drops even further
- To prevent thrashing, must provide enough frames for each process
  - **Working-set model**
  - **Page-fault frequency**

94

94

## Working-Set Model


- **Locality**: a set of pages that are actively used together
- Locality model: as a process executes, it moves from locality to locality
  - Program structure (subroutine, loop, stack)
  - Data structure (array, table)
- Working-set model (based on locality model)
  - Working-set **window**: a parameter  $\Delta$  (delta)
  - Working-set: set of pages in most recent  $\Delta$  page references (an approximation locality)

95

95

## Working-Set Example

- If  $\Delta$  (delta) = 10



96

96



## Working-Set Model (cont.)

- Prevent thrashing using the working-set size
  - WSS: working-set size for process  $i$
  - $D = \sum WSS_i$  (total demand frames)
  - if  $D > m$  (available frames)  $\rightarrow$  thrashing
  - The OS monitors the  $WSS_i$  of each process and allocates to the process enough frames
    - if  $D \ll m$ , increase degree of MP
    - If  $D > m$ , suspend a process
- Prevent thrashing while keeping the degree of multiprogramming as high as possible
- Optimize CPU utilization
- **However, too expensive for tracking**

97

97


## Page Fault Frequency Scheme

- **Page fault frequency** directly measures and controls the page-fault rate to prevent thrashing
  - Establish **upper** and **lower** bounds on the desired page-fault rate of a process
  - If page fault rate exceeds the upper limit
    - Allocate another frame to the process
  - If page rate falls below the lower limit
    - Remove a frame from the process

98

98

## Page Fault Frequency Scheme (cont.)



99

99

That's All !



100

100