



# Midterm Review

**Operating Systems**

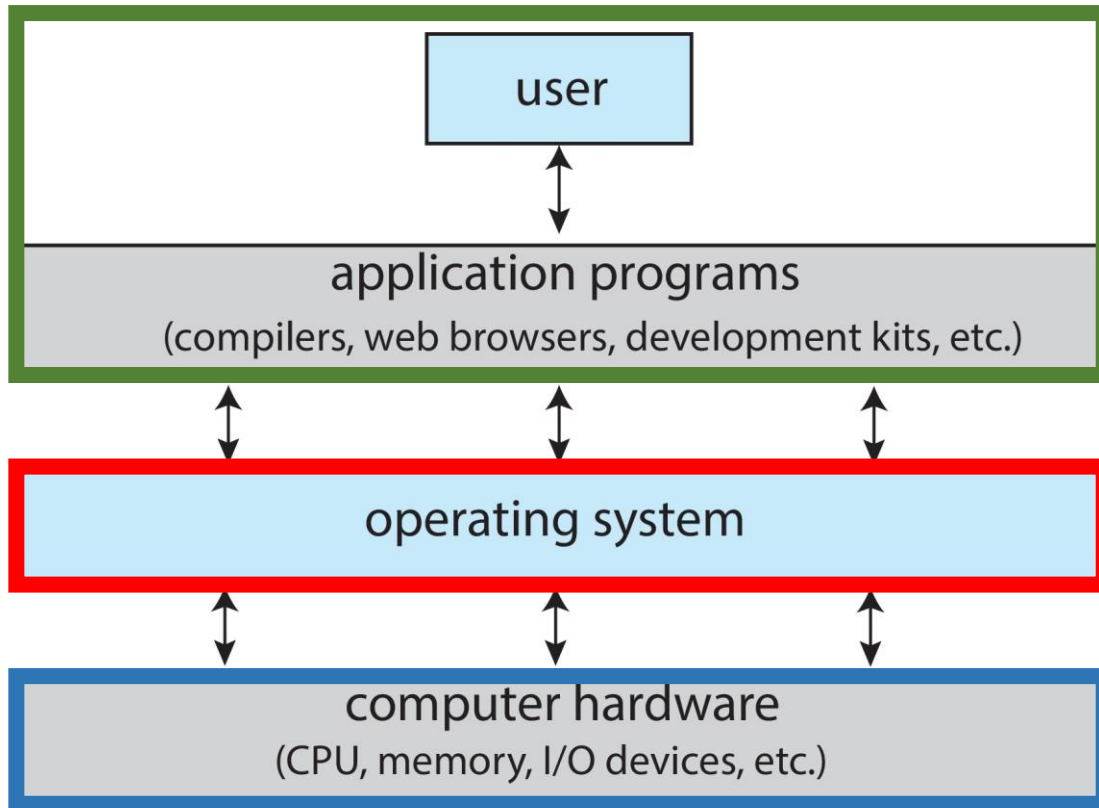
**Yu-Ting Wu**

# Scope of the Midterm

- Chapter 1: Introduction
- Chapter 2: Operating system structure
- Chapter 3: Processes
- Chapter 4: Threads and concurrency
- Chapter 5: CPU scheduling

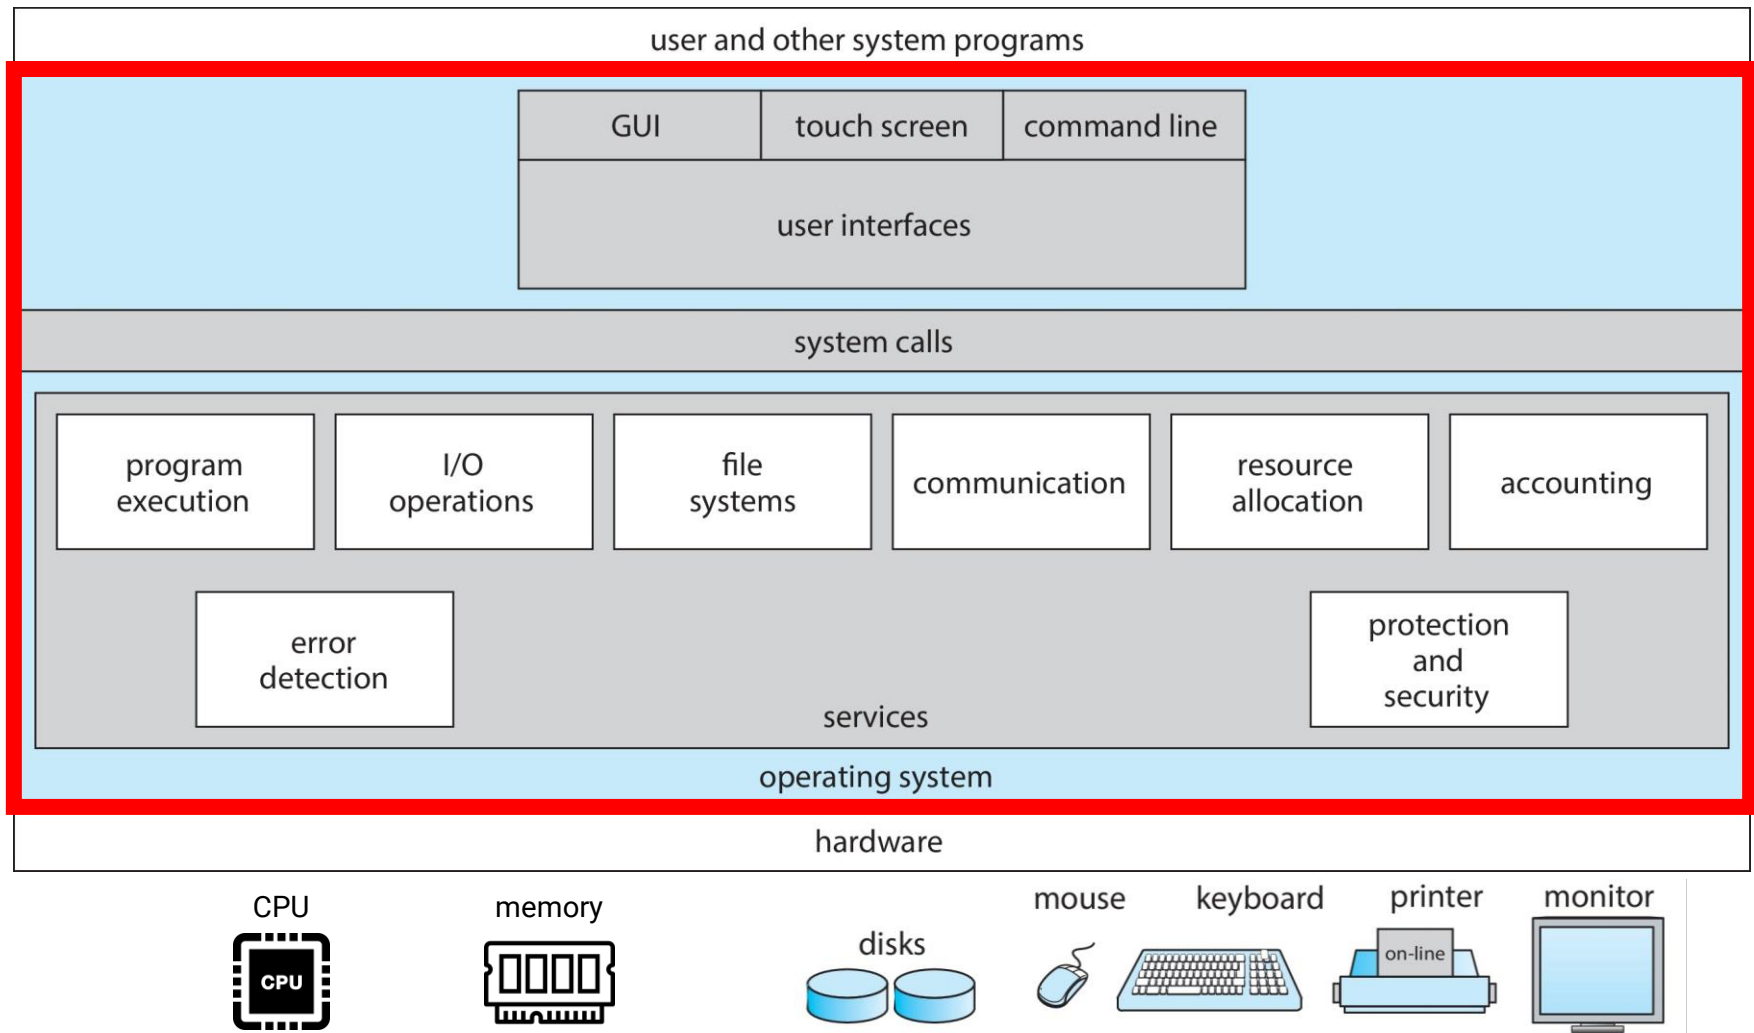
# Overview

# Goals of an OS

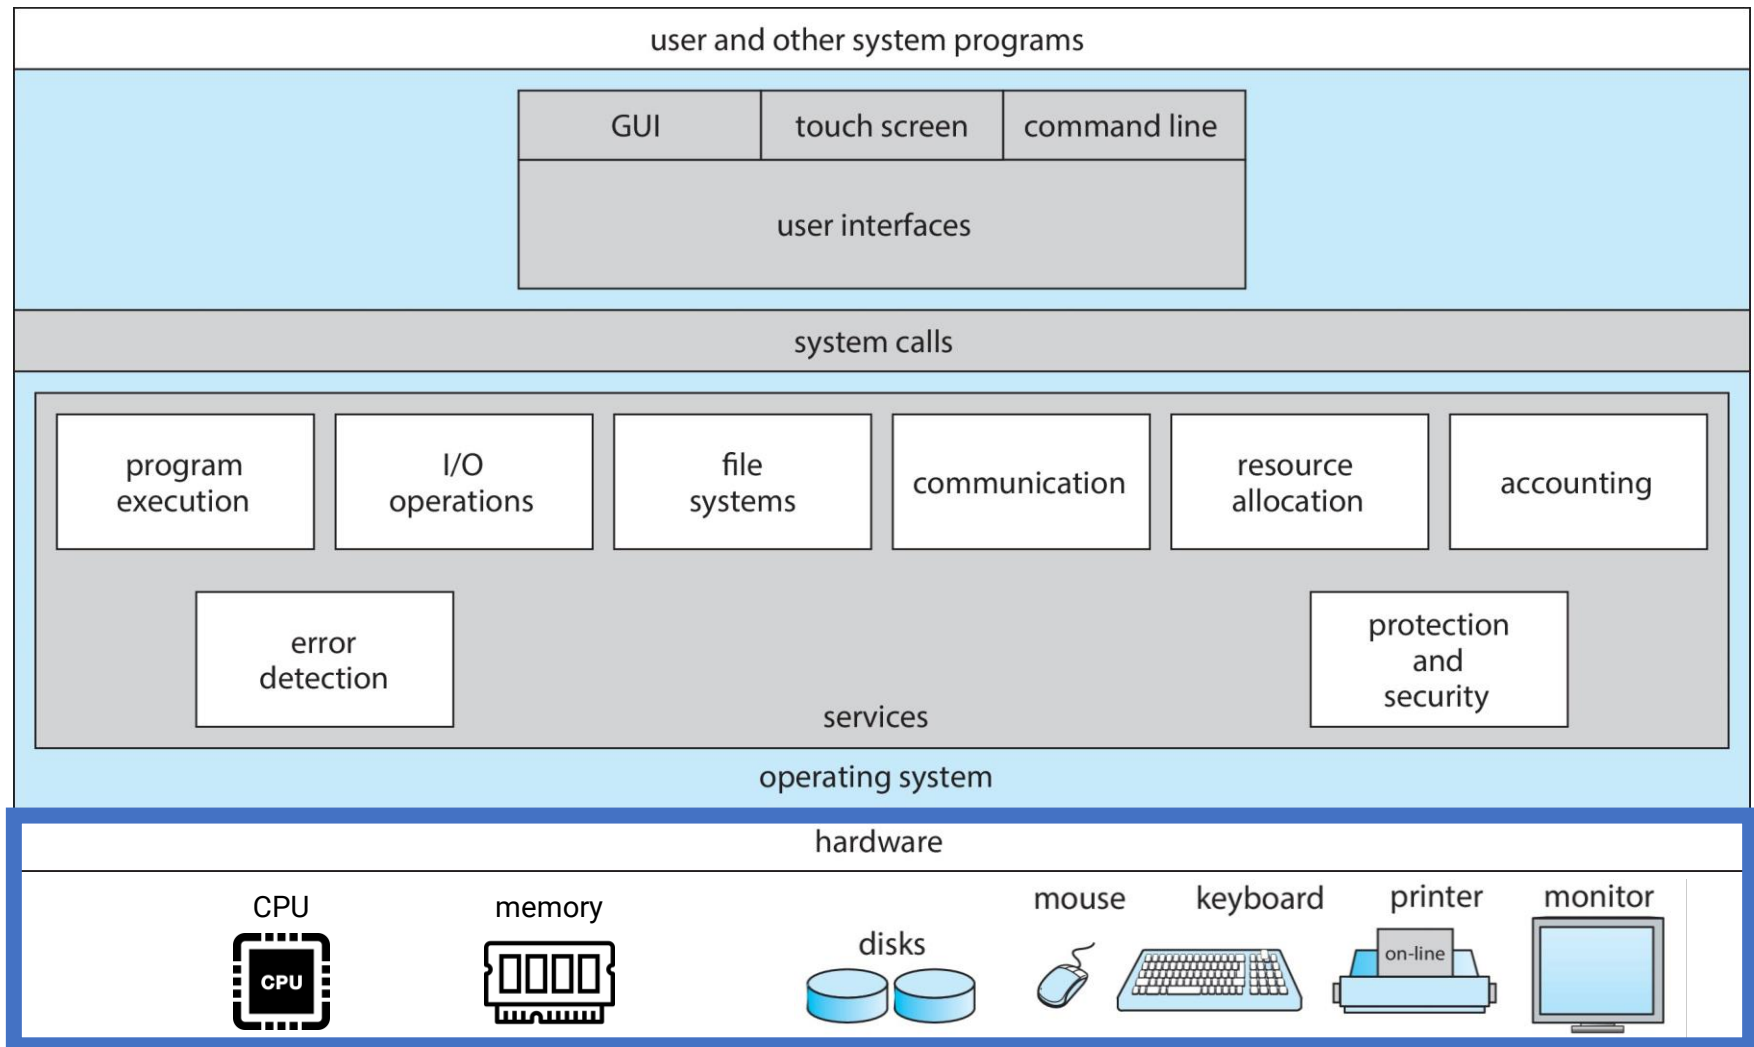


- **Control program**
  - Execute users' program
  - Make the computer system convenient to use
- **Resource allocator**
  - Use the computer hardware in an efficient manner

# Operating System Services

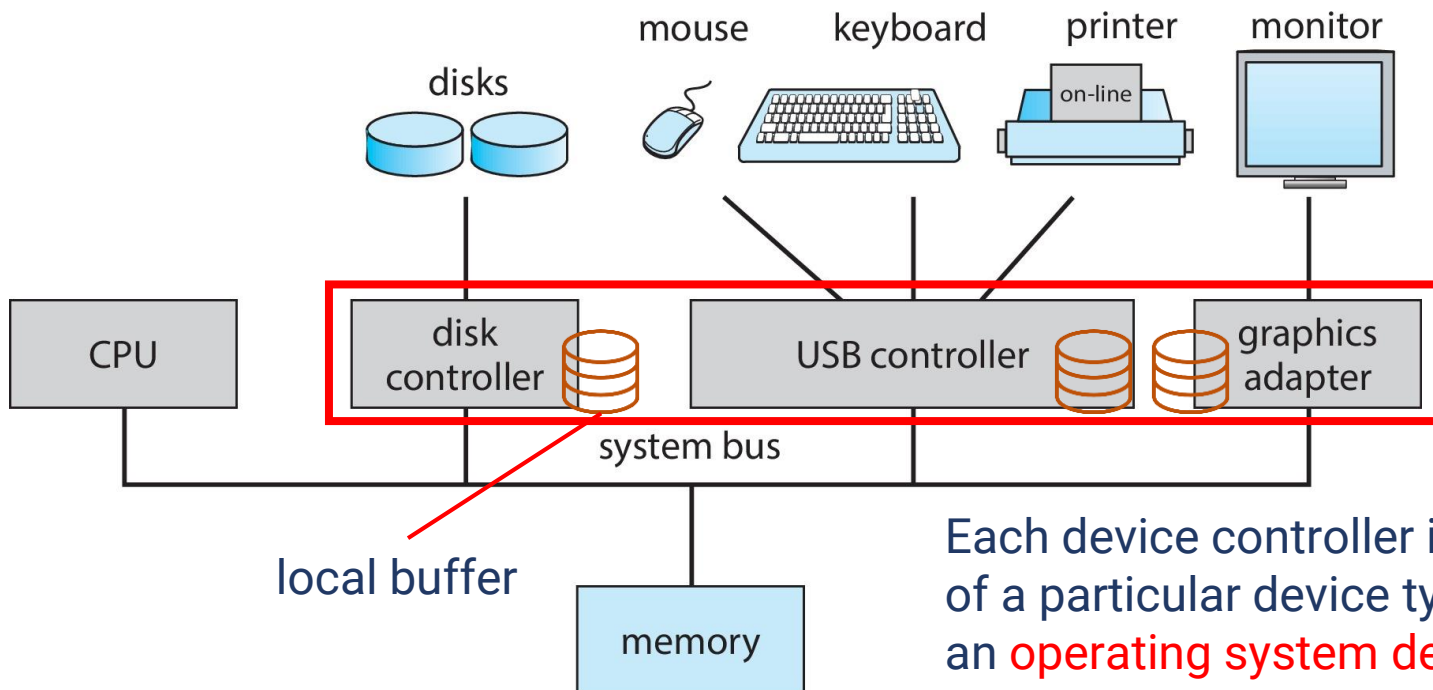


# Computer System Organization



# Computer System Operations

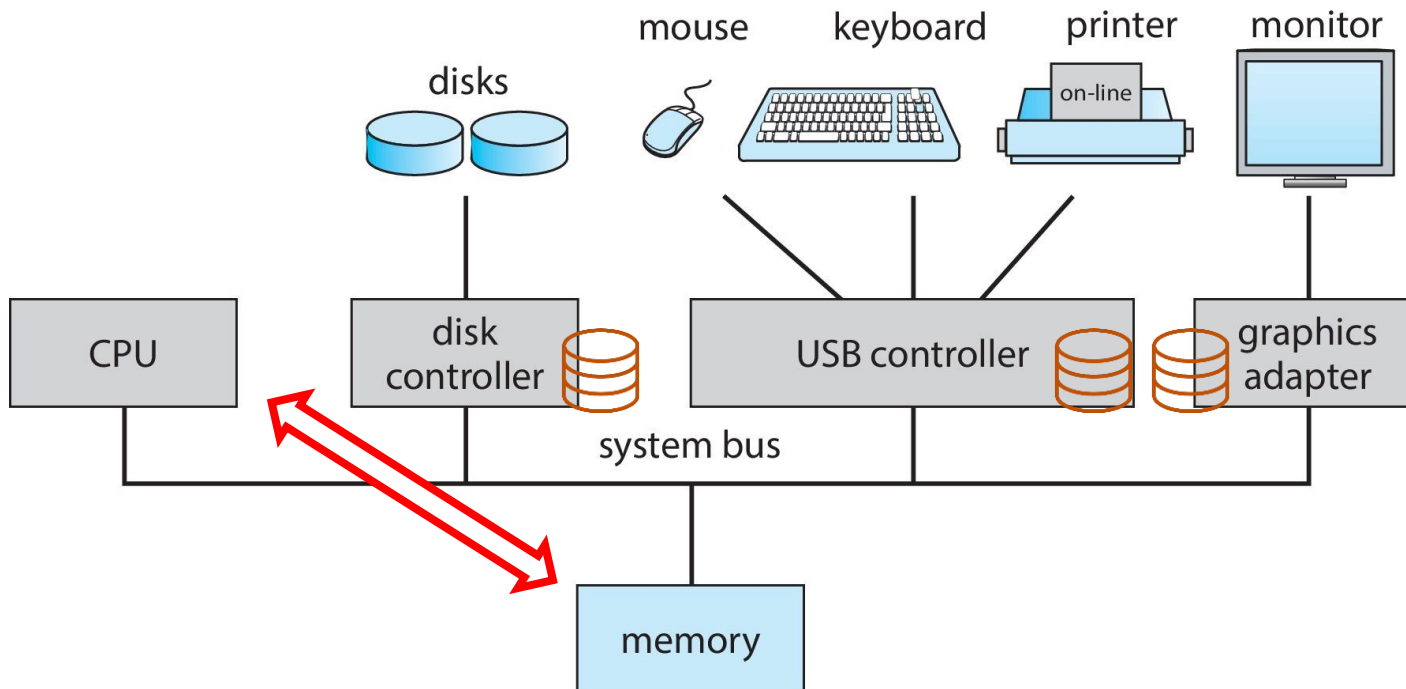
- CPU (or CPUs) and device controllers connect through common **bus**, which provides access to memory



Each device controller is in charge of a particular device type, and has an **operating system device driver** to manage it

# Computer System Operations (cont.)

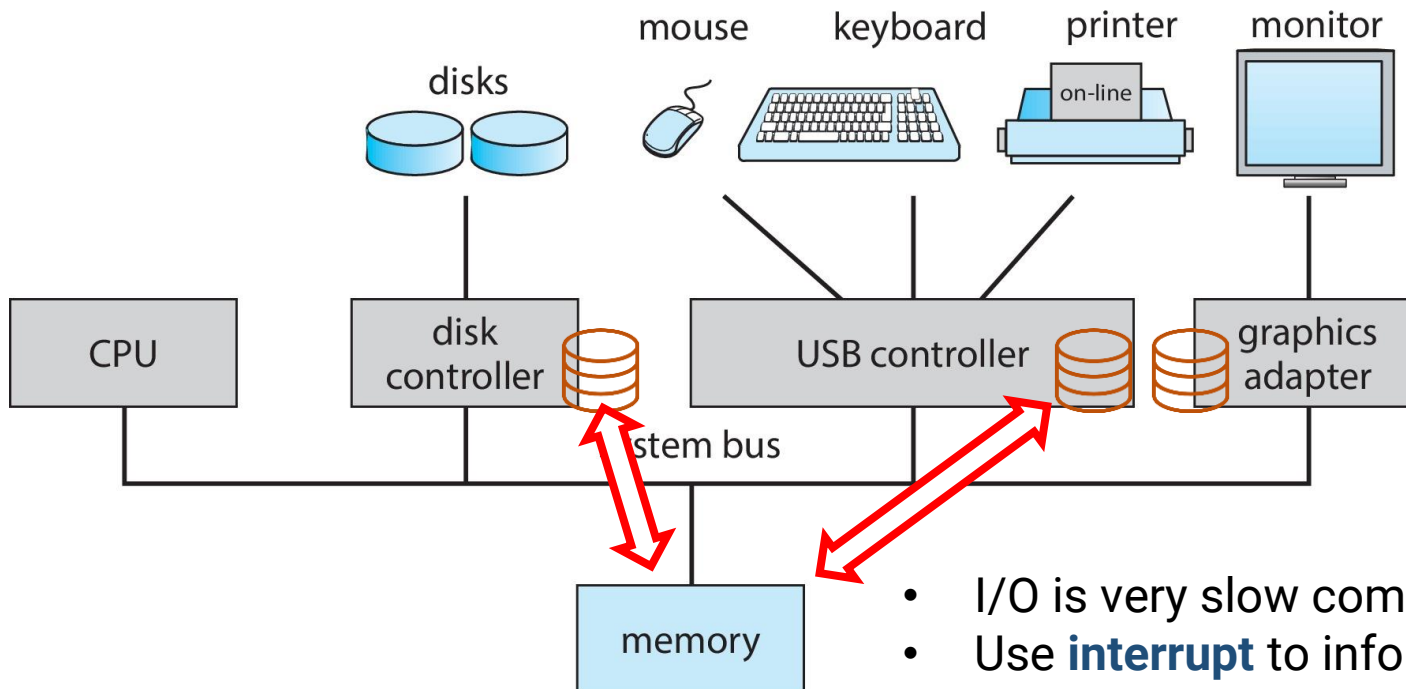
- CPU moves data from/to main memory to/from local buffer for executing programs
  - Main memory is the **only** storage that CPU can access





# Computer System Operations (cont.)

- I/O: from the device to local buffer of controller
  - **Read:** devices → controller buffers → memory
  - **Write:** memory → controller buffers → devices



- I/O is very slow compared to CPU
- Use **interrupt** to inform CPU that it has finished its task

# Interrupt

- Interrupt provides a way to change the **flow of control** in the CPU
- **Hardware interrupt (signal)**
  - Service requests from one of the devices
    - Examples: keyboard, mouse click, etc.
- **Software interrupt (trap)**
  - Invalid memory access
  - Software error
    - Example: division by zero
  - **System calls**
    - Request for system services

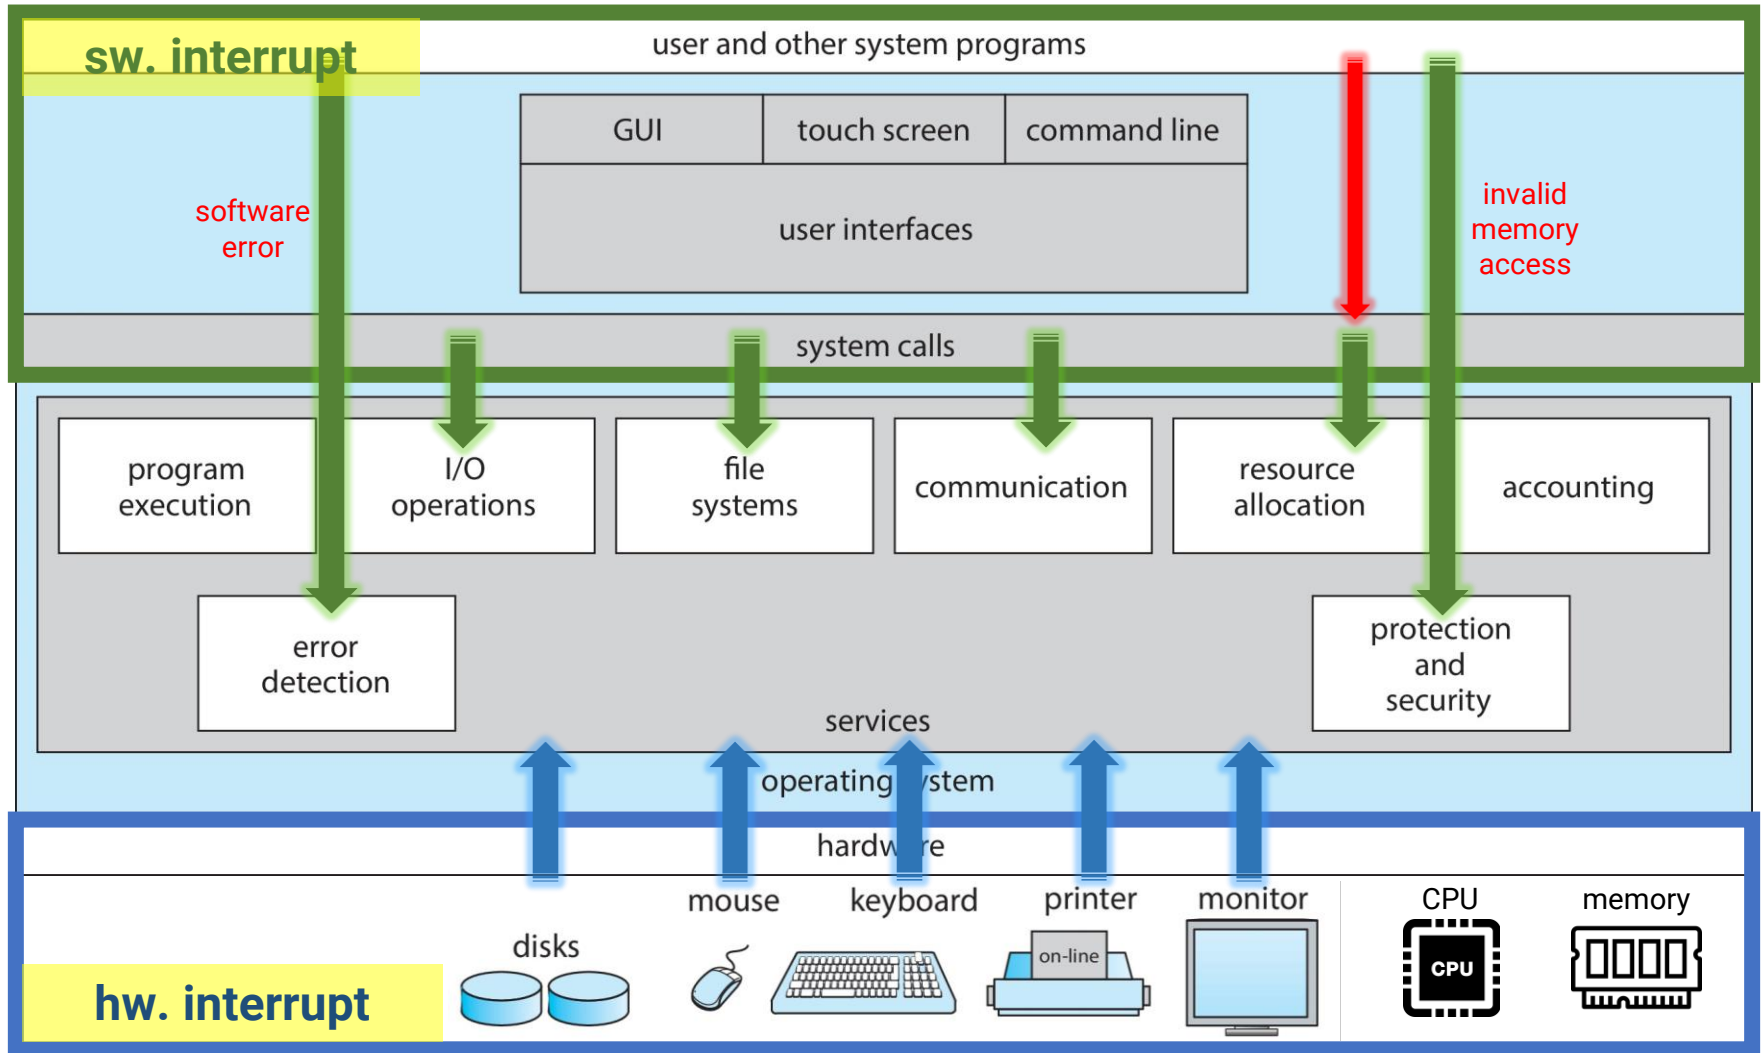


```
#include <stdio.h>

int main(int argc, const char * argv[]) {
    FILE* fp = fopen("test.txt", "r");
    if (fp) {
        printf("\nNot NULL");
    } else {
        printf("NULL");
    }
}
```


➔ system call (open)

# Computer System Organization



# Interrupt Timeline

- Must **save the address** of the interrupted instruction
  - Usually done by using **stacks**
- Transfer control to the **interrupt service routine**
  - Generic handler / interrupt vector / hybrid



# **A Brief Introduction to Hardware Architecture**


# Processor

- **Single-processor system**
  - One main CPU per system
  - Example: earlier desktop or mobile device
- **Multiple-processor system**
  - Multiple CPUs per system

# Multi-Processor System


- **Tightly coupled**

- More than one processor in close communication sharing bus, memory, and peripheral devices



- **Loosely coupled**

- Otherwise (such as distributed systems)
- Each machine has its own memory



# Multi-Processor System (cont.)

- **Symmetric model**

- Each processor in the system runs an identical copy of the OS


- **Asymmetric model**

- Master-slave fashion
- Commonly seen in extremely large systems



# Multi-Processor System (cont.)

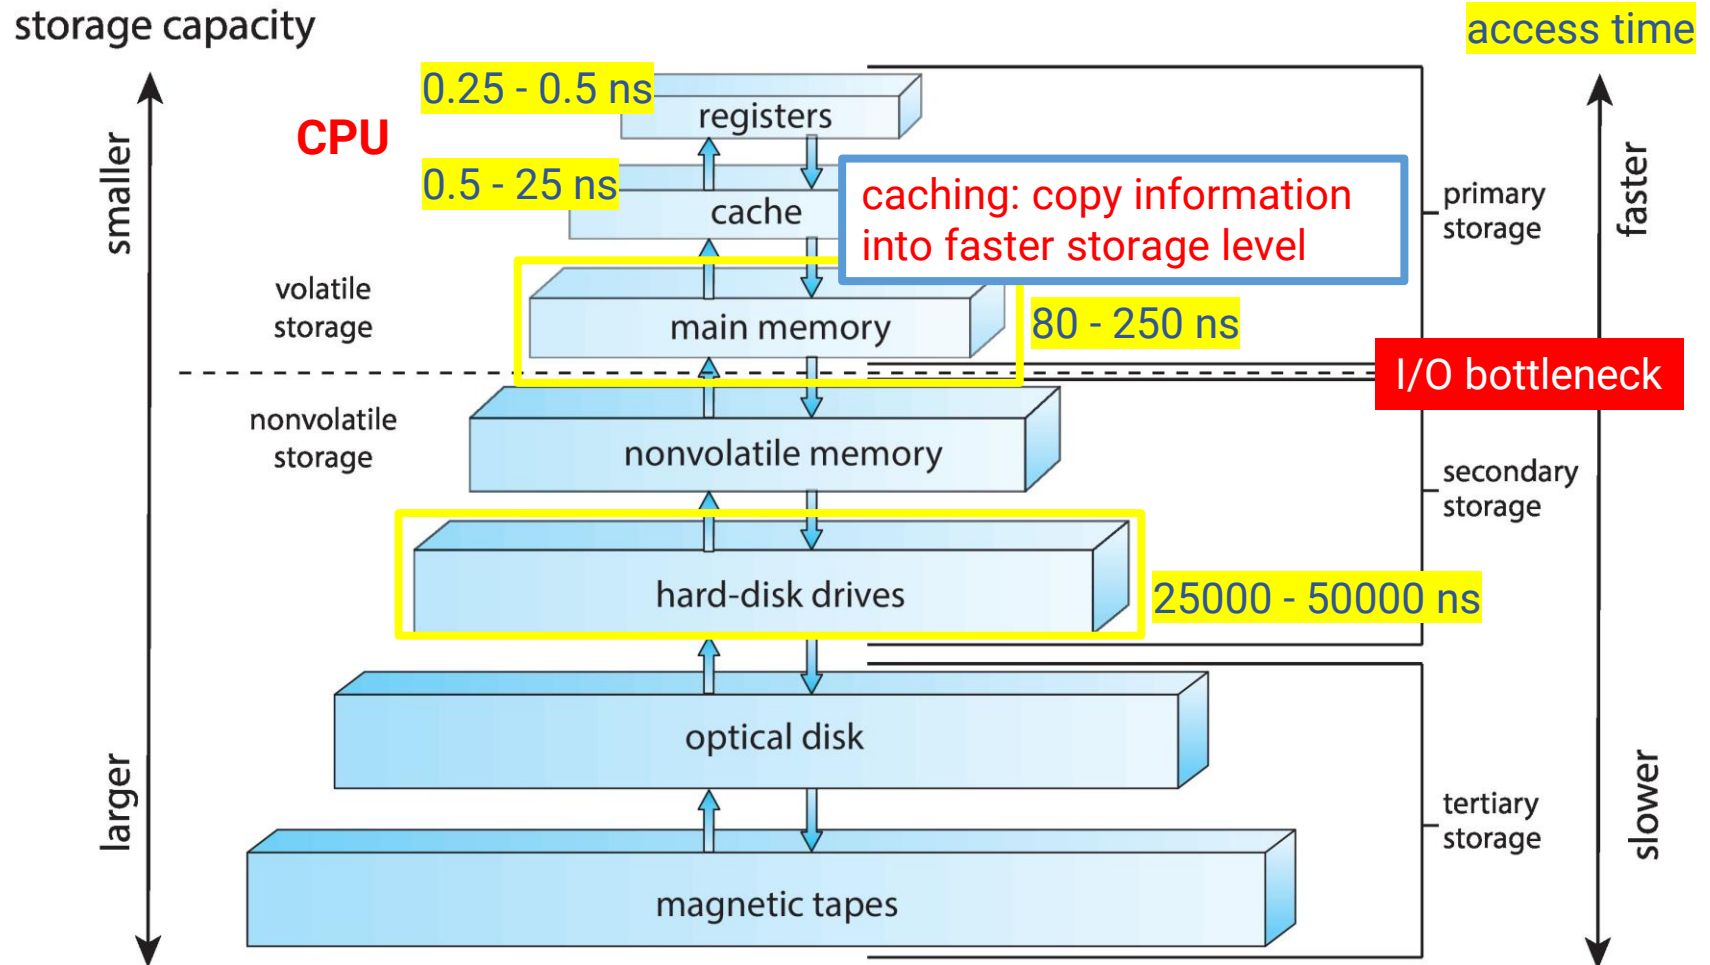
- **Advantages of multi-processor systems**
  - **Speedup**: better throughput
  - **Lower cost**: building one small fast chip is very expensive
  - **More reliable**: Graceful degradation and fail soft
- The recent trend: from a fast single processor to lots of processors
  - **Multiple cores** over a single chip
  - Hyper-threading (logical core)



# Storage Structure


- Organized in hierarchy based on
  - **Speed**
  - **Cost**
  - **Volatility**
- **Main memory**
  - The **only large storage media** that the CPU can access directly
  - Typically volatile
- **Secondary storage** (ex: HDD, USB sticks, CD, DVD, ...)
  - Extension of main memory that provides large storage capacity
  - Typically nonvolatile

# Storage Structure (cont.)



# Caching


- Information is **copied** to a faster storage system on a **temporary** basis
- Assumption: data will be used again soon (locality)
- Cache management
  - Cache size
  - Replacement policy



# Memory Access Architecture


- **Uniform Memory Access (UMA)**

- Most commonly represented today by **symmetric multi-processor (SMP)** machines
- Equal access times to memory
- Example: most commodity computers



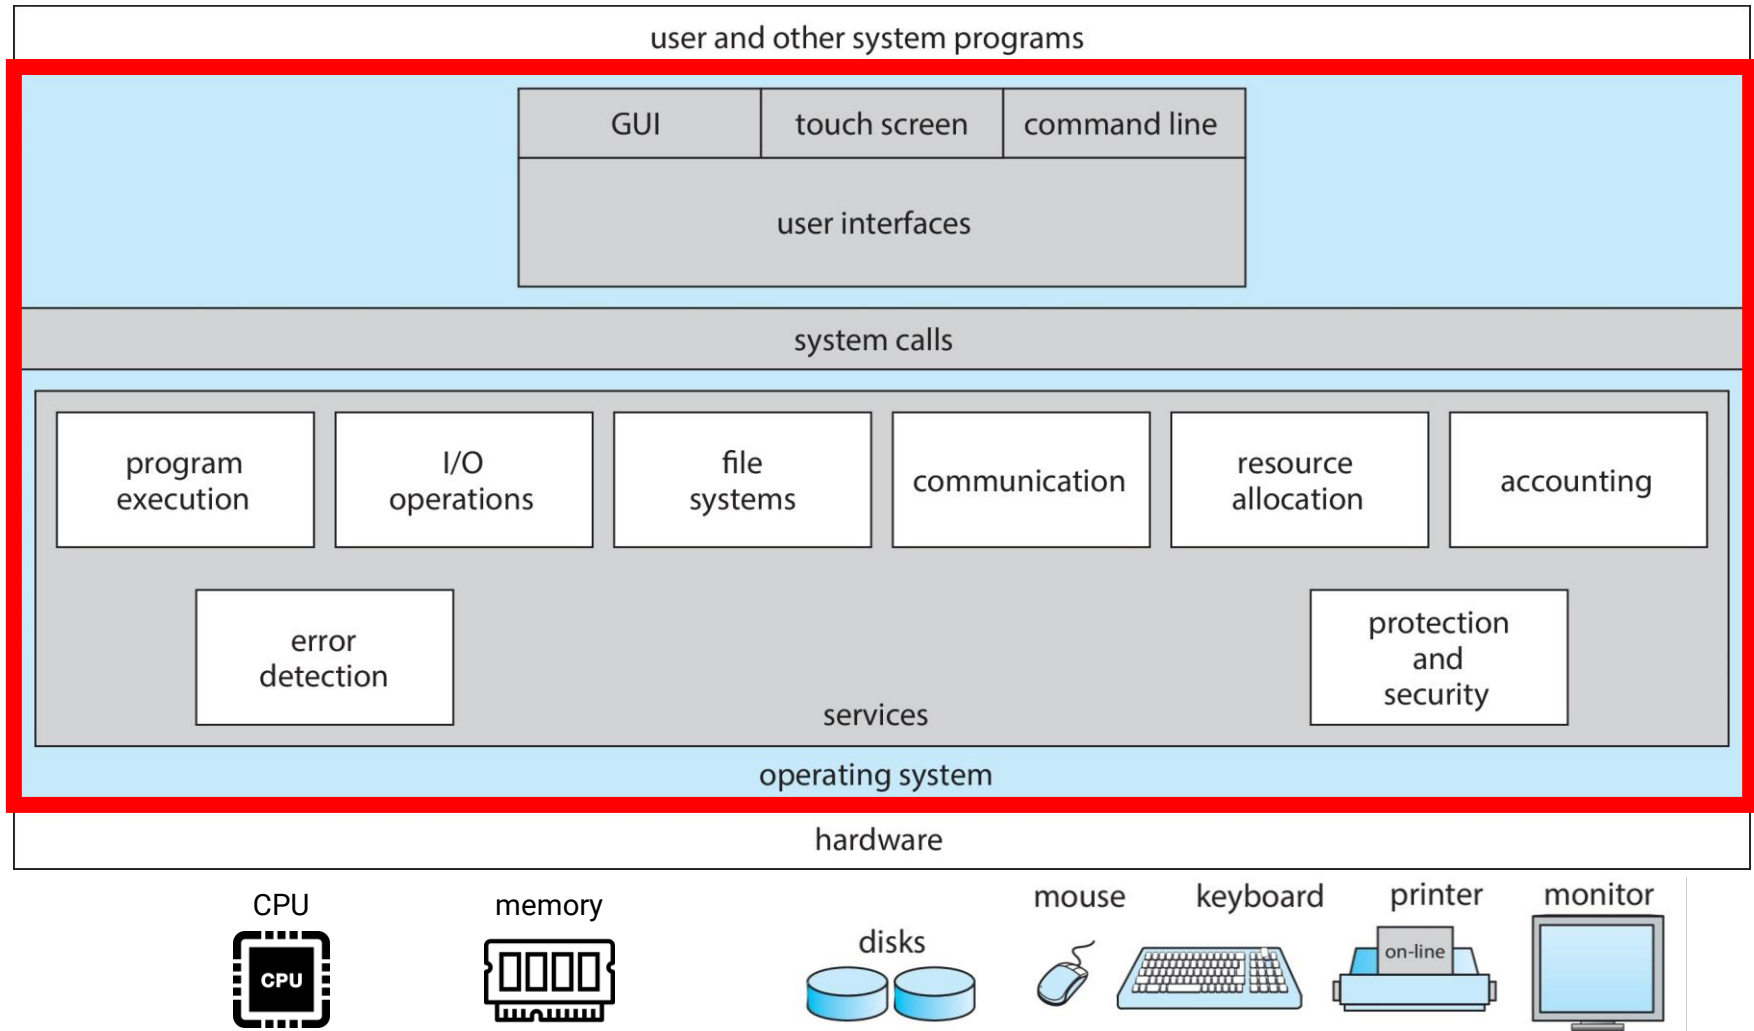
- **Non-Uniform Memory Access (NUMA)**

- Often made by physically **linking two or more SMPs**
- One SMP can directly access memory of another SMP
- Memory access across link is slower
- Example: IBM Blade server



# Operating System Services

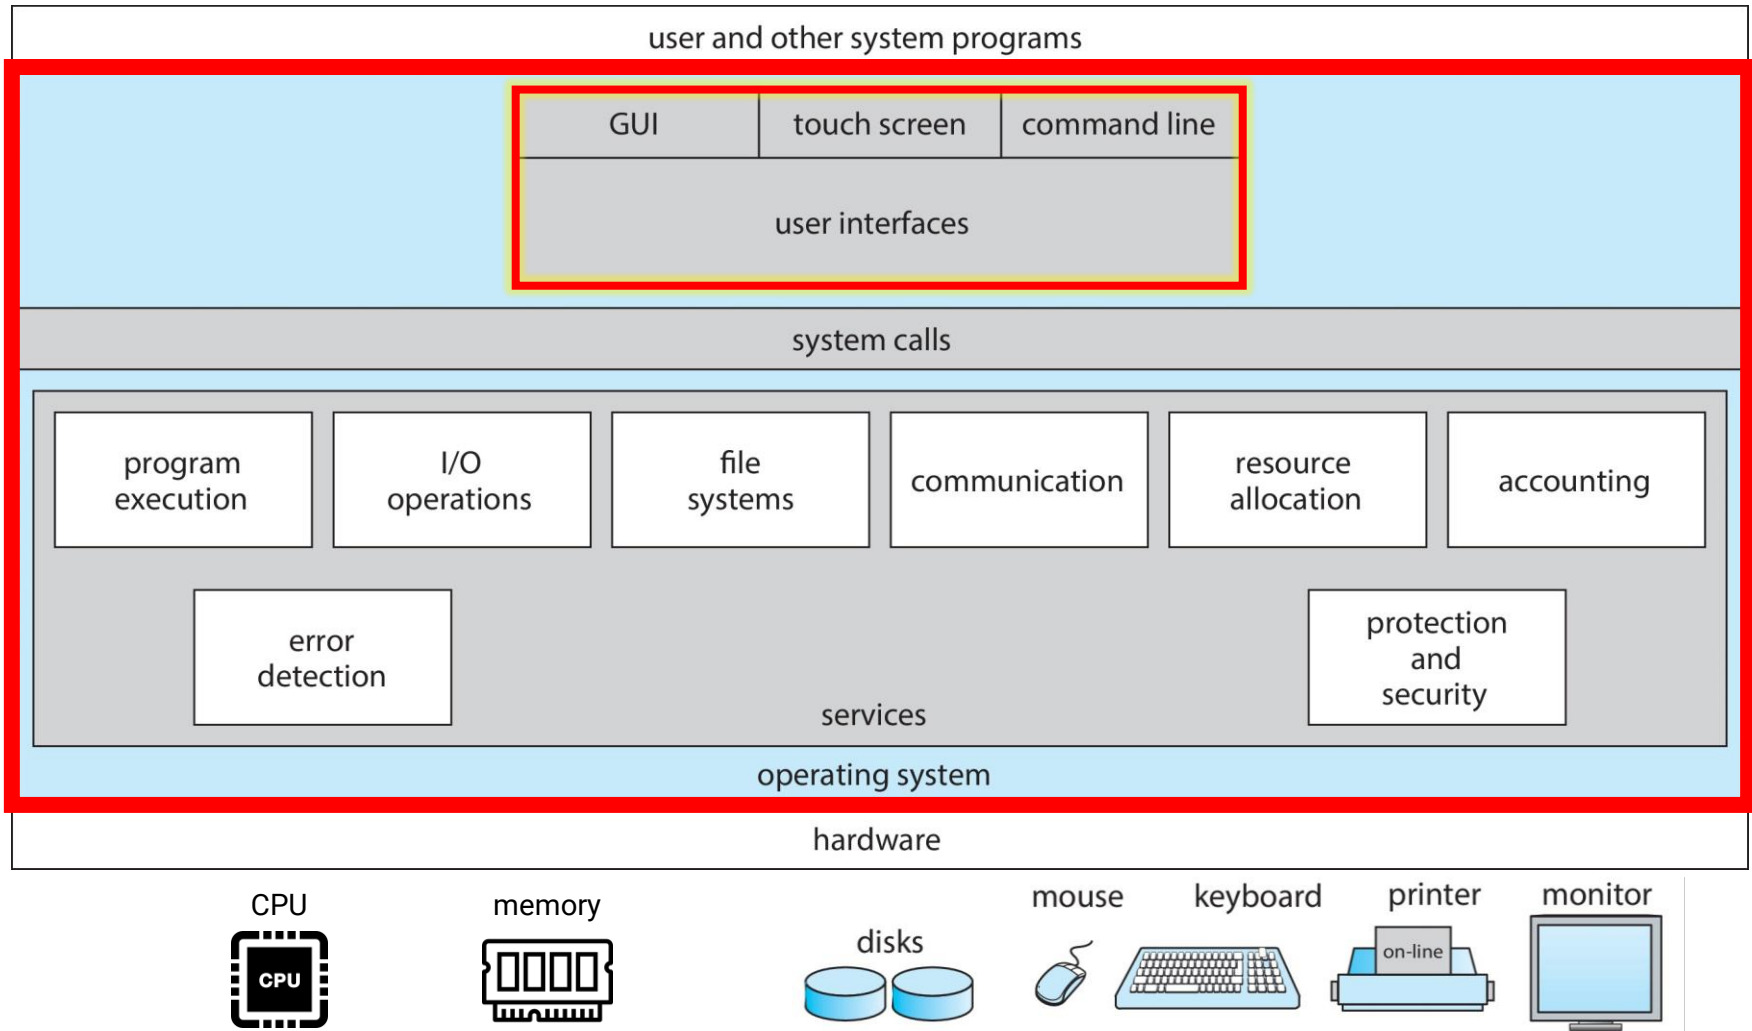
# Operating System Services



# **Operating System Service: User Interfaces**



# Operating System Services



# User Interface

- **Command line interface (CLI)**

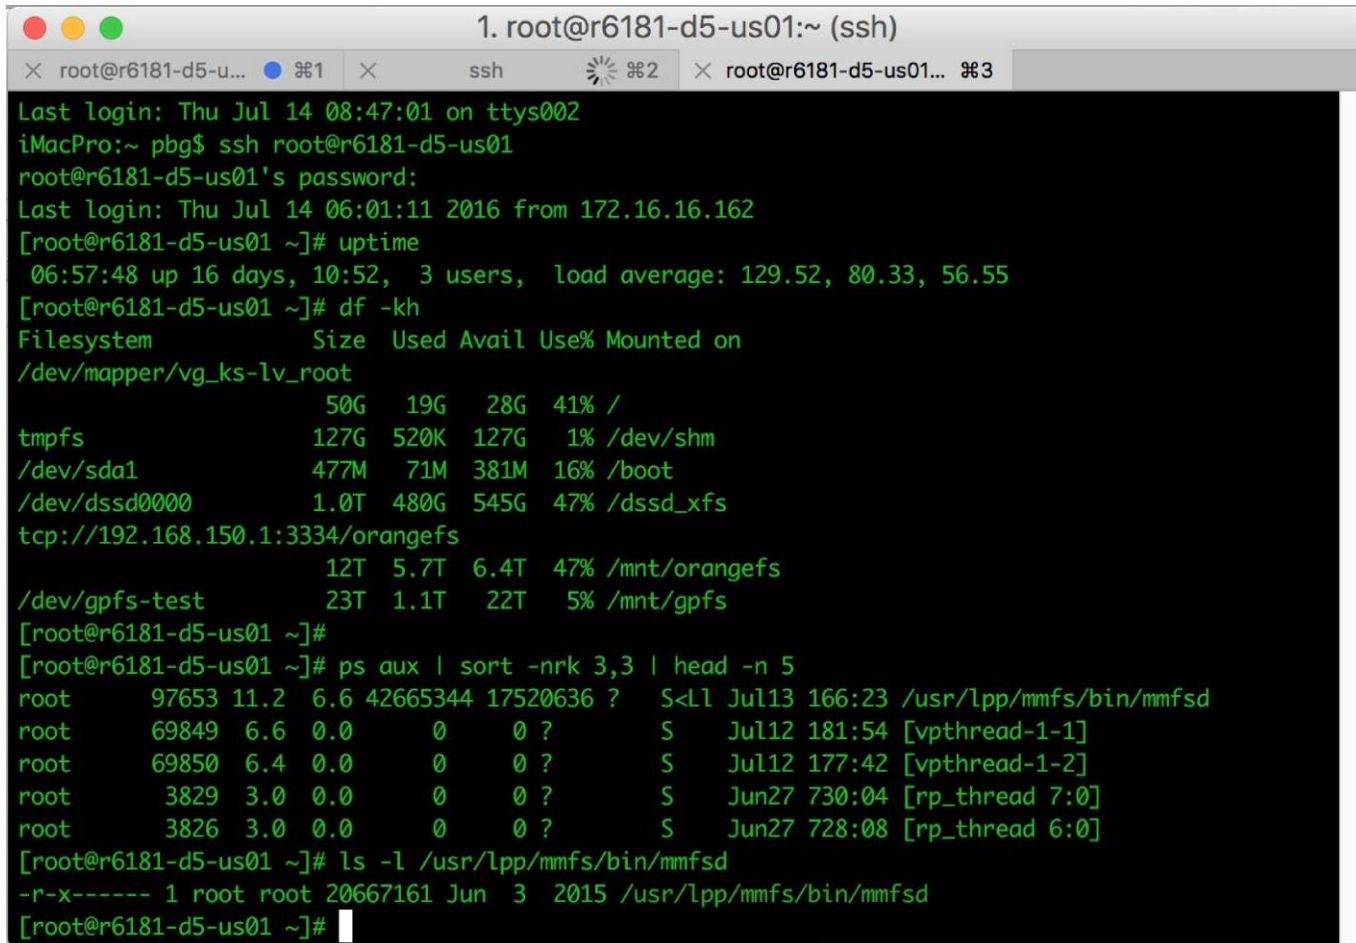
- Fetch a command from user and execute it
- Shell (command-line interpreter)
  - Ex: CSHELL, BASH
  - Allow to some modification based on user behavior and preference

- **Graphic user interface (GUI)**

- Usually with mouse, keyboard, and monitor
- Icons are used to represent files, directories, programs, etc.
- Usually built on CLI

- Most systems have both CLI and GUI

# Command Line Interface




```

1. root@r6181-d5-us01:~ (ssh)
X root@r6181-d5-u...  X ssh  X root@r6181-d5-us01...
Last login: Thu Jul 14 08:47:01 on ttys002
iMacPro:~ pbg$ ssh root@r6181-d5-us01
root@r6181-d5-us01's password:
Last login: Thu Jul 14 06:01:11 2016 from 172.16.162
[root@r6181-d5-us01 ~]# uptime
 06:57:48 up 16 days, 10:52,  3 users,  load average: 129.52, 80.33, 56.55
[root@r6181-d5-us01 ~]# df -kh
Filesystem                Size      Used Avail Use% Mounted on
/dev/mapper/vg_ks-lv_root   50G       19G   28G   41% /
tmpfs                      127G      520K   127G    1% /dev/shm
/dev/sda1                   477M       71M   381M   16% /boot
/dev/dssd0000               1.0T     480G   545G   47% /dssd_xfs
tcp://192.168.150.1:3334/orangefs 12T     5.7T   6.4T   47% /mnt/orangefs
/dev/gpfs-test             23T     1.1T   22T    5% /mnt/gpfs
[root@r6181-d5-us01 ~]#
[root@r6181-d5-us01 ~]# ps aux | sort -nrk 3,3 | head -n 5
root      97653 11.2  6.6 42665344 17520636 ?    S<Ll  Jul13 166:23 /usr/lpp/mmfs/bin/mmfstd
root      69849  6.6  0.0      0      0 ?        S    Jul12 181:54 [vpthread-1-1]
root      69850  6.4  0.0      0      0 ?        S    Jul12 177:42 [vpthread-1-2]
root       3829  3.0  0.0      0      0 ?        S    Jun27 730:04 [rp_thread 7:0]
root       3826  3.0  0.0      0      0 ?        S    Jun27 728:08 [rp_thread 6:0]
[root@r6181-d5-us01 ~]# ls -l /usr/lpp/mmfs/bin/mmfstd
-r-x----- 1 root root 20667161 Jun  3  2015 /usr/lpp/mmfs/bin/mmfstd
[root@r6181-d5-us01 ~]#

```

Bourne Shell (default shell of UNIX ver. 7)



# Graphic User Interface



Mac OS X GUI

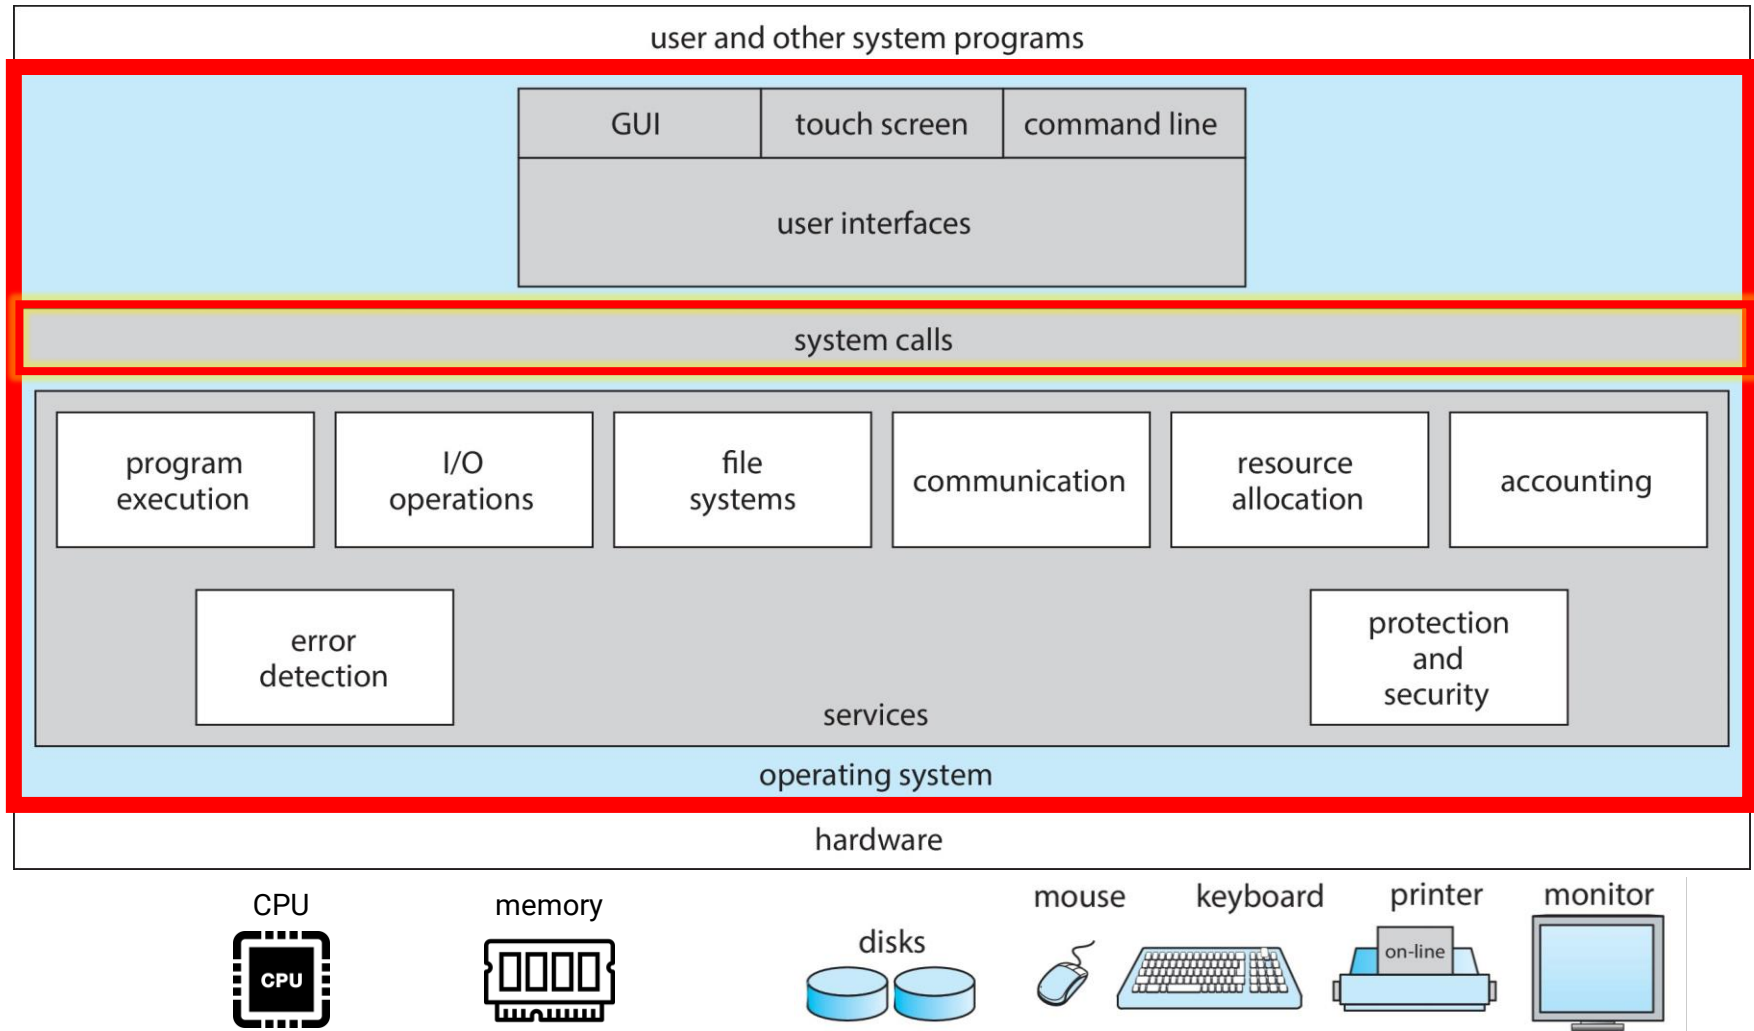
# Other Interfaces

- Batch
- Touch-screen
- Voice control



# System Calls and API

# Operating System Services



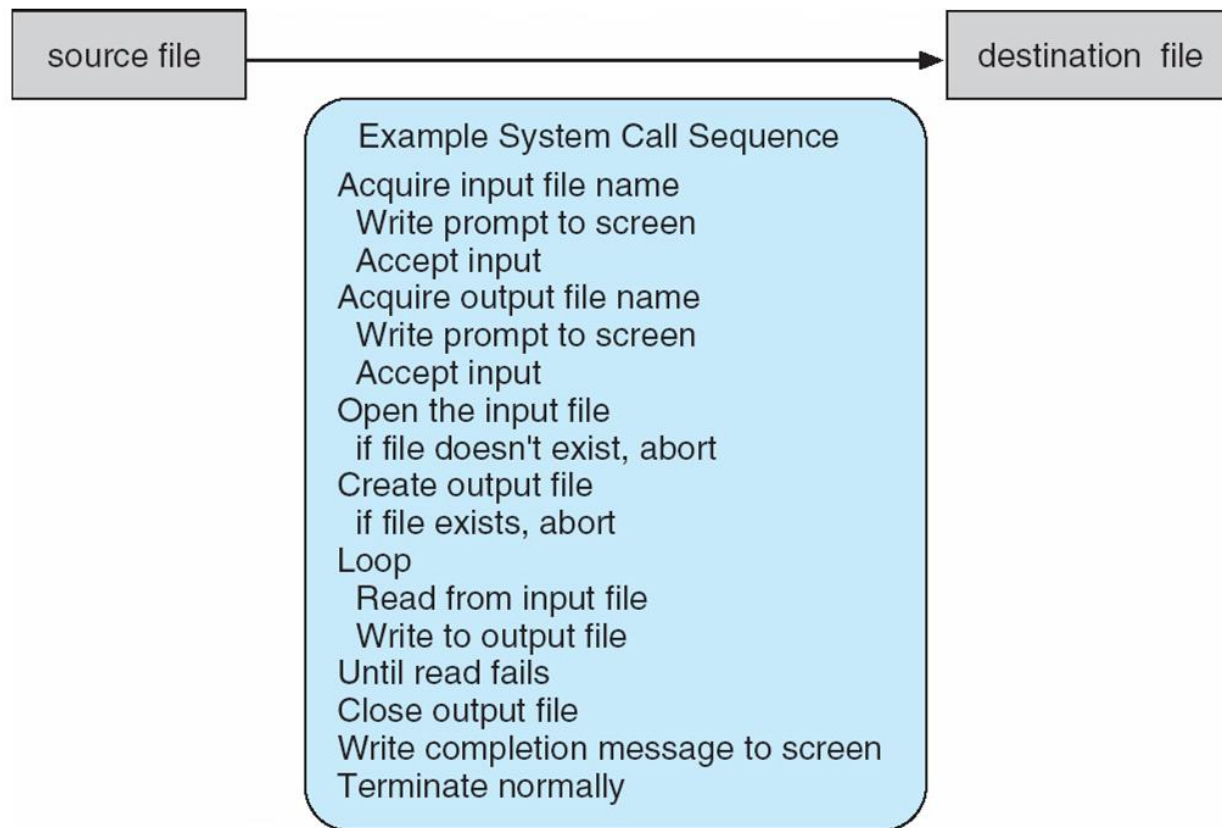
# System Calls

- **Programming interface to the services provided by the OS**
  - An explicit request to the kernel made via **software interrupt**
  - Generally available as assembly-language instructions
- Mostly accessed by programs via a high-level **Application Programming Interface (API)** rather than direct system call use



# System Calls (cont.)

- Example: a sequence of system calls for copying a file



# System Calls (cont.)

- Request OS services
  - **Process control**
    - End (normal exit) or abort (abnormal)
    - Load and execute
    - Create and terminate
    - Get or set attributes of process
    - Wait for a specific amount of time or an event
    - Memory dumping, profiling, tracing, allocate, and free
  - **File management**
    - Create and delete
    - Open and close
    - Read, write, and reposition
    - Get or set attributes
    - Operations for directories


# System Calls (cont.)

- Request OS services (cont.)
  - **Device management**
    - Request or release
    - Logically attach or detach devices
  - **Information maintenance**
    - Get or set time or date
    - Get or set system data (e.g., maximum memory for a process)
  - **Communications**
    - Send and receive messages
    - Message passing or shared memory
  - **Protection**

# Application Programming Interface (API)

- An **encapsulation** of system calls for user programs
- Provide **portability**
- Usually implemented by high-level languages
  - C library, Java
- Could involve zero or multiple system calls
  - `abs()`: zero
  - `fopen()`: multiple
  - `malloc()`, `free()` → `brk()`

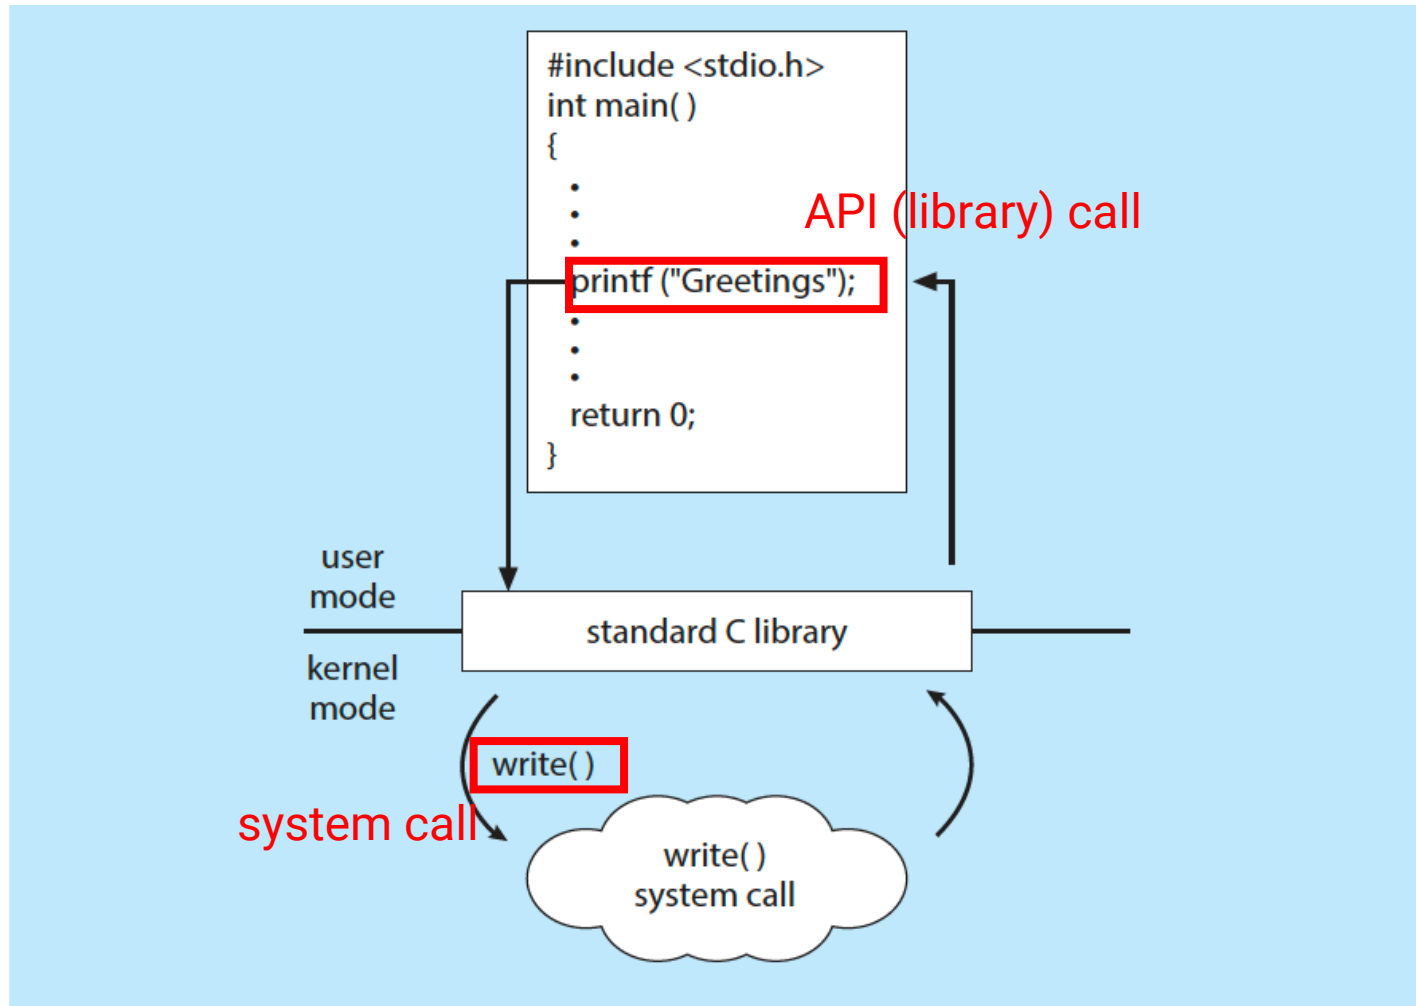
e.g., Win32 API



# API (cont.)

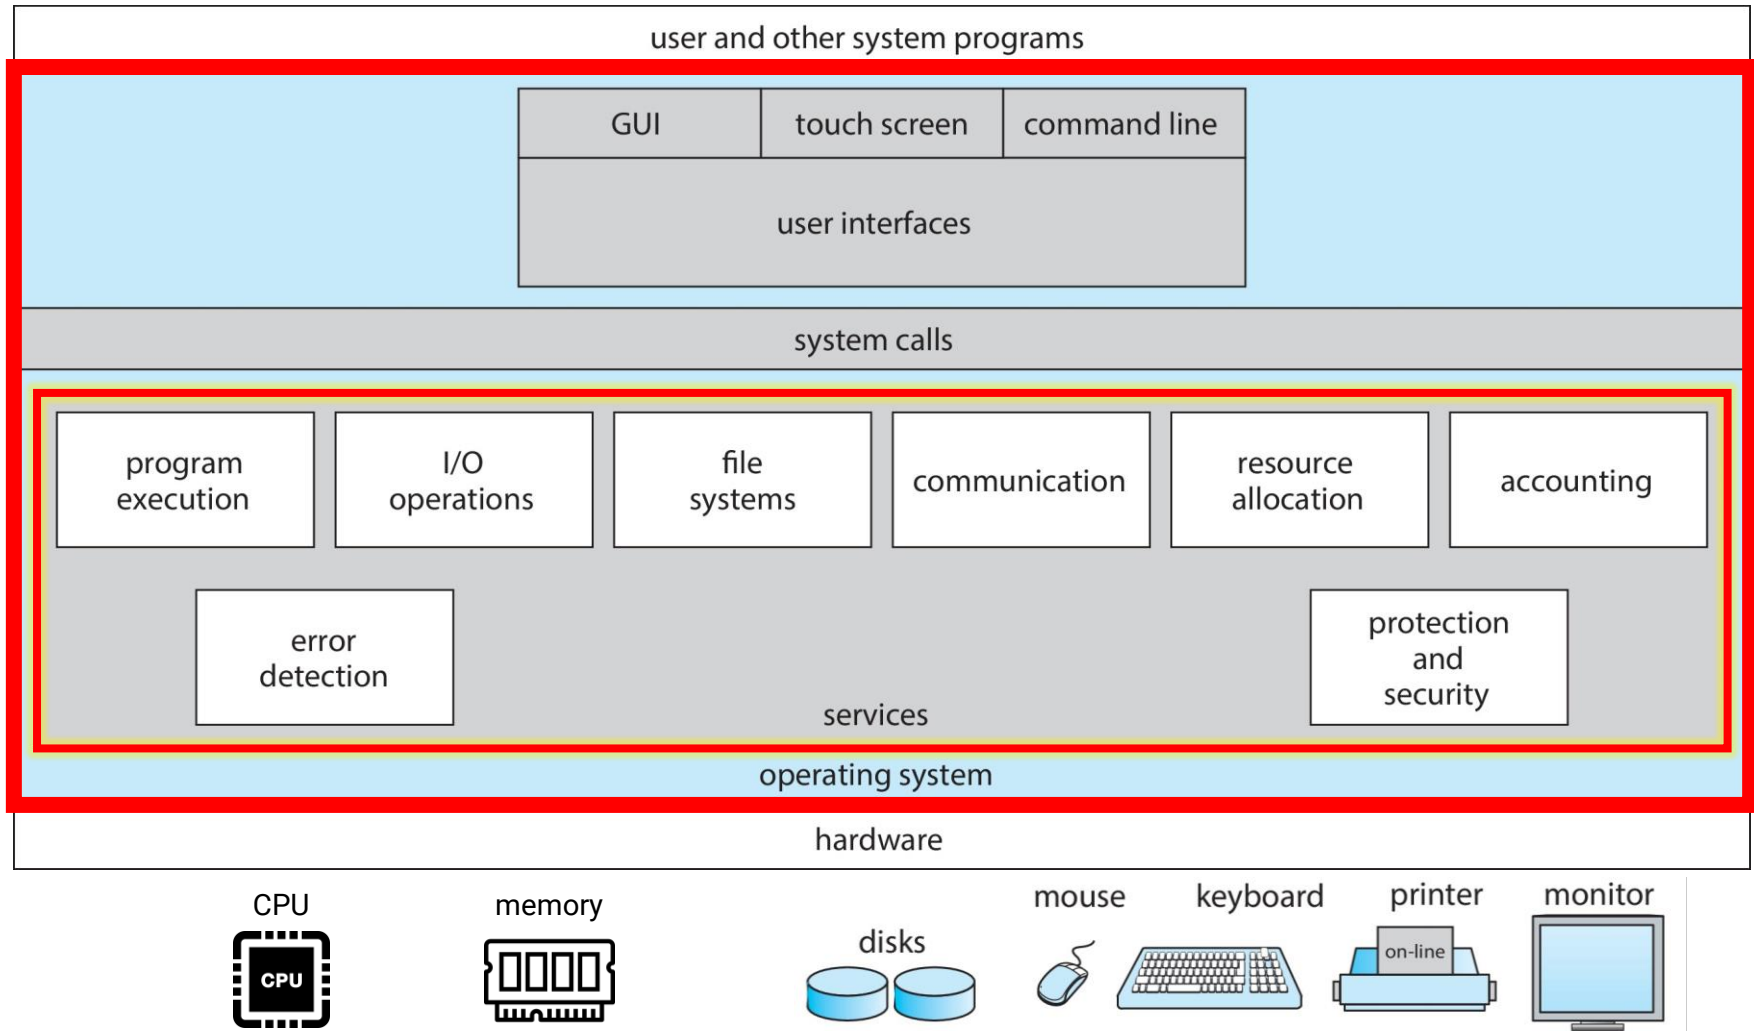
- Three most common APIs
  - **Win32 API**
    - For Microsoft Windows
    - [https://en.wikipedia.org/wiki/Windows\\_API](https://en.wikipedia.org/wiki/Windows_API)
    - <https://docs.microsoft.com/zh-tw/windows/win32/apiindex/windows-api-list?redirectedfrom=MSDN>
  - **POSIX API**
    - POSIX stands for **P**ortable **O**perating **S**ystem Interface for Unix
    - Used by Unix, Linux, and Mac OS X
    - <https://en.wikipedia.org/wiki/POSIX>
  - **Java**
    - For Java virtual machine (JVM)

# System Call and API



# Operating System Design

# Operating System Services



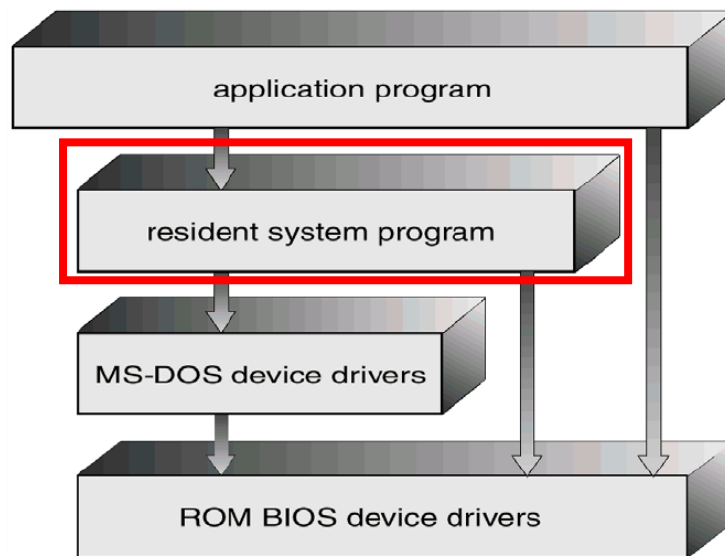


# Overview of OS Structure

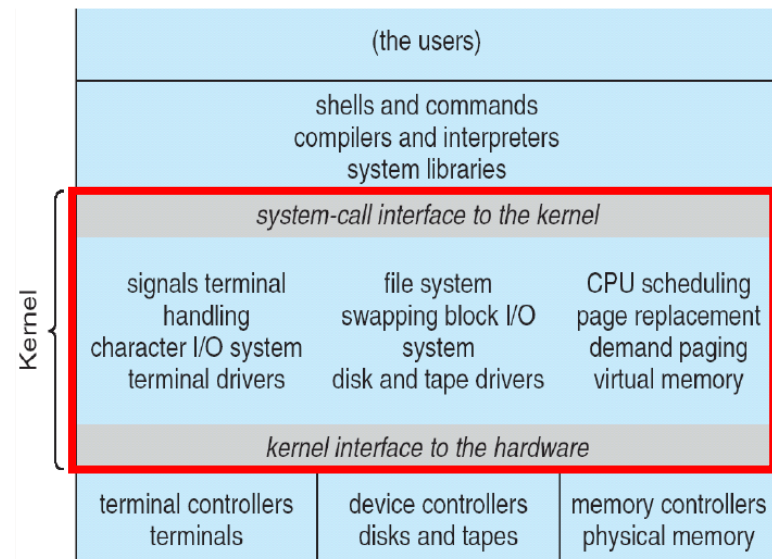
- Simple OS architecture
- Layer OS architecture
- Microkernel OS
- Modular OS architecture
- Hybrid systems
- Virtual machine

# Simple OS Architecture

- Only one or two levels
- Drawbacks
  - Unsafe
  - Difficult to enhance



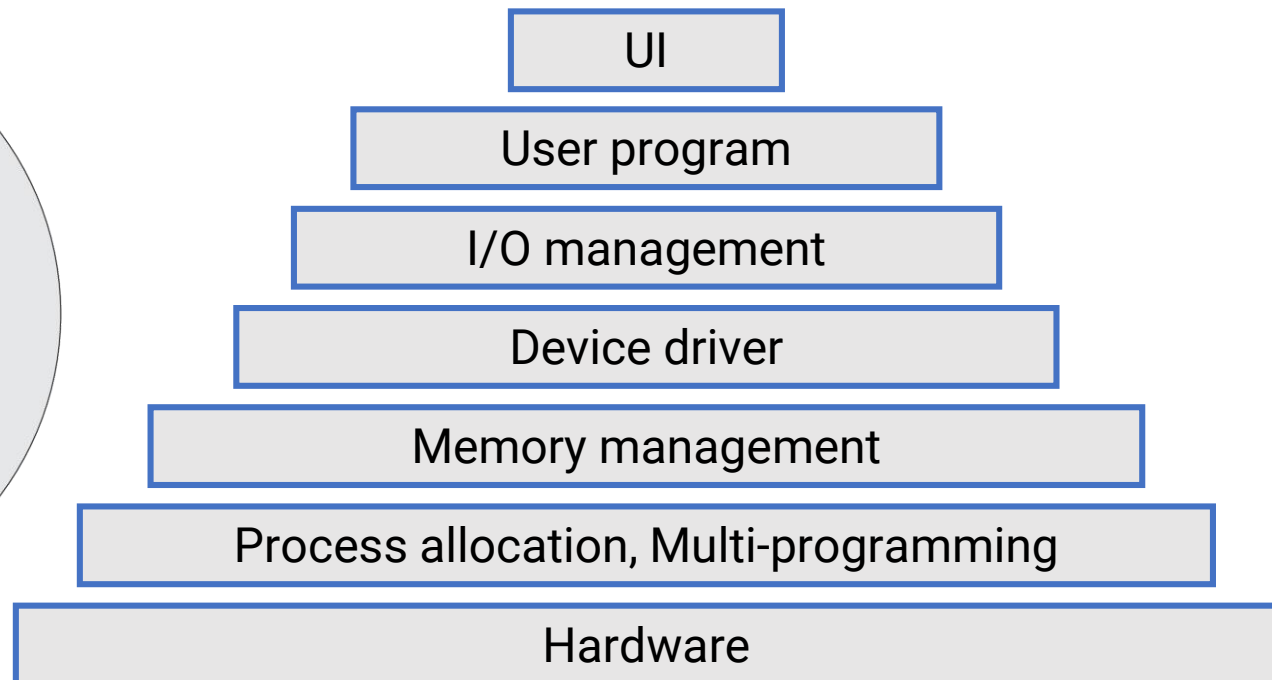

MS-DOS



UNIX


# Layered OS Architecture

- Lower levels are independent of upper levels
- Pros: easier debugging and maintenance
- Cons: less efficient and difficult to define layers



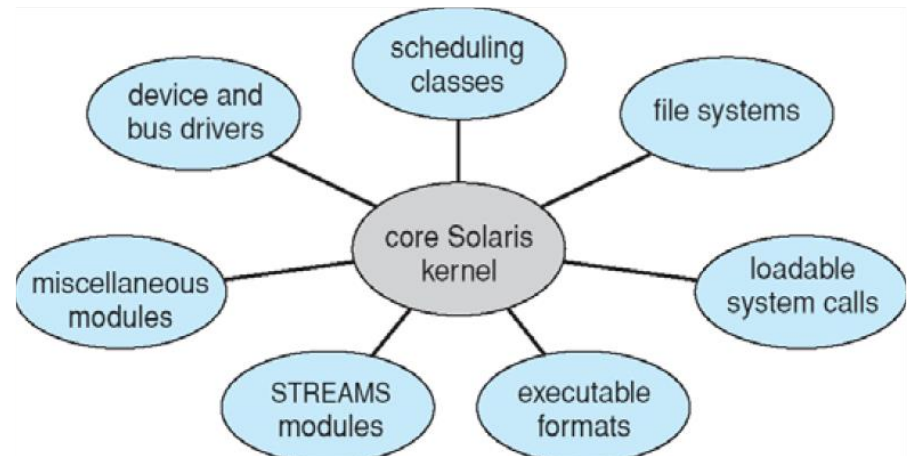
# Microkernel OS

- **Kernel should be as small as possible**
  - Move most parts of the original kernels into user space
- Communication is provided by **message passing**
- Easier for extending and porting
- Slow



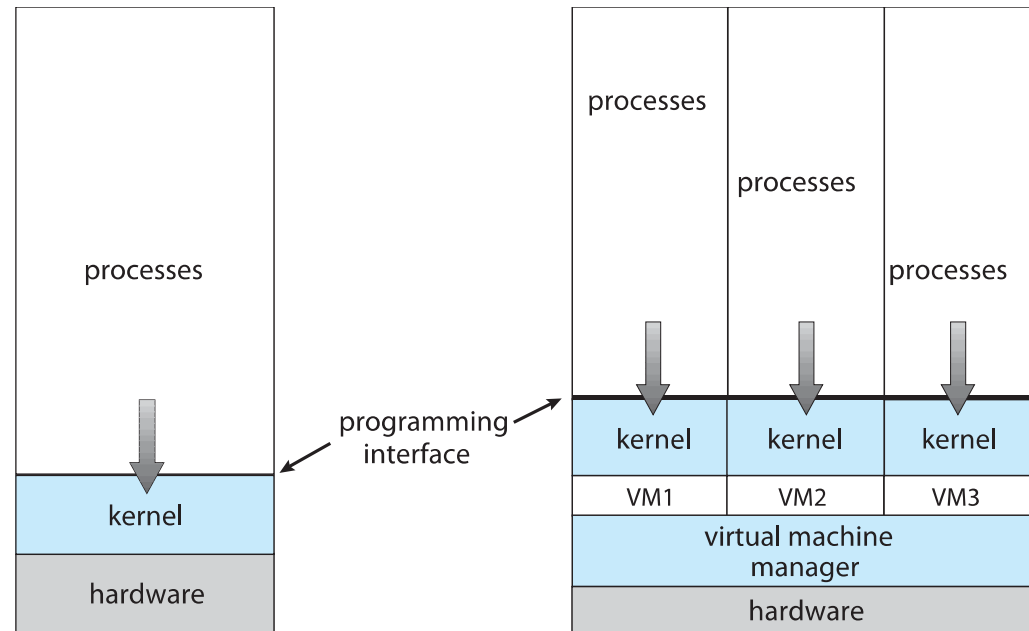
# Modular OS Architecture

- **Employed by most modern OS**
  - Object-oriented approach
  - Each core component is separate
  - Each module talks to the others over known interfaces
  - Each module is loadable as needed **within the kernel**
- Similar to layers but with more flexibility



# Virtual Machine

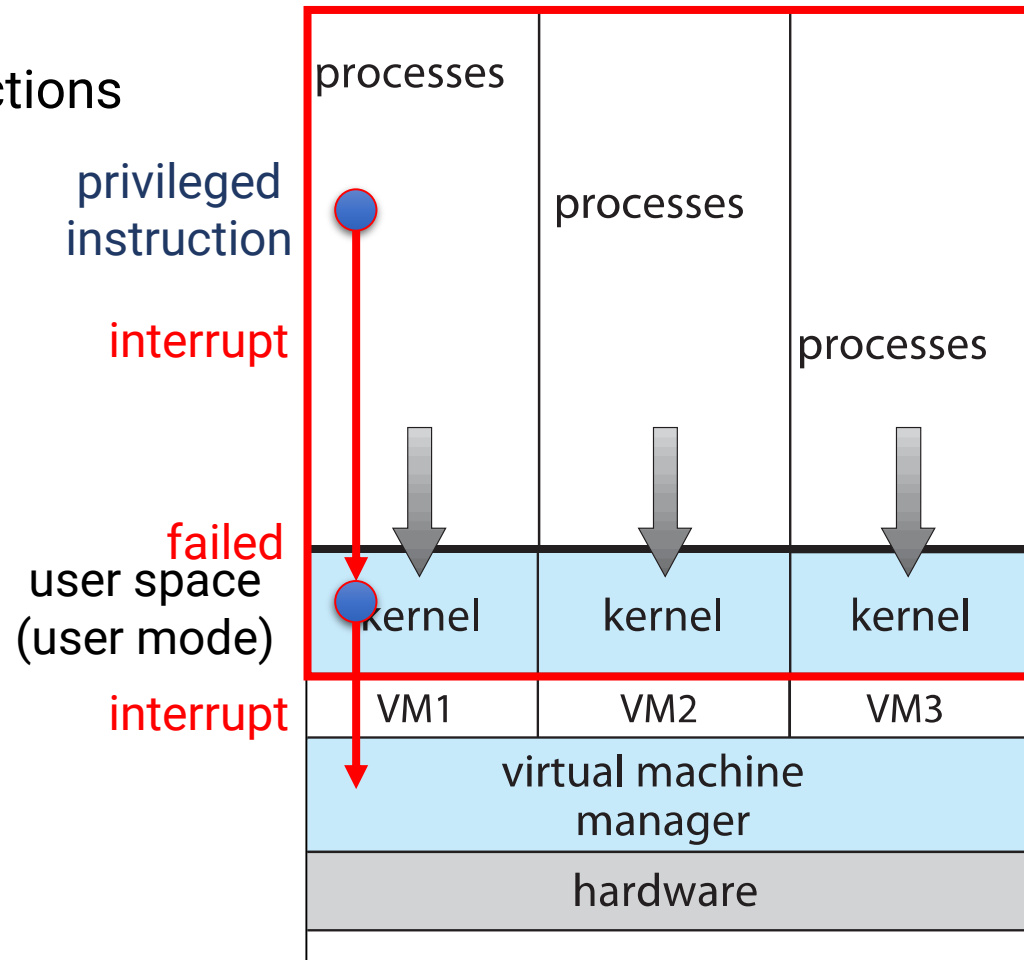
- Provide an interface that is identical to the underlying bare hardware
  - Each process is provided with a (virtual) copy of the underlying computer



# Virtual Machine (cont.)

## • Challenges

- Privileged instructions



# Virtual Machine (cont.)

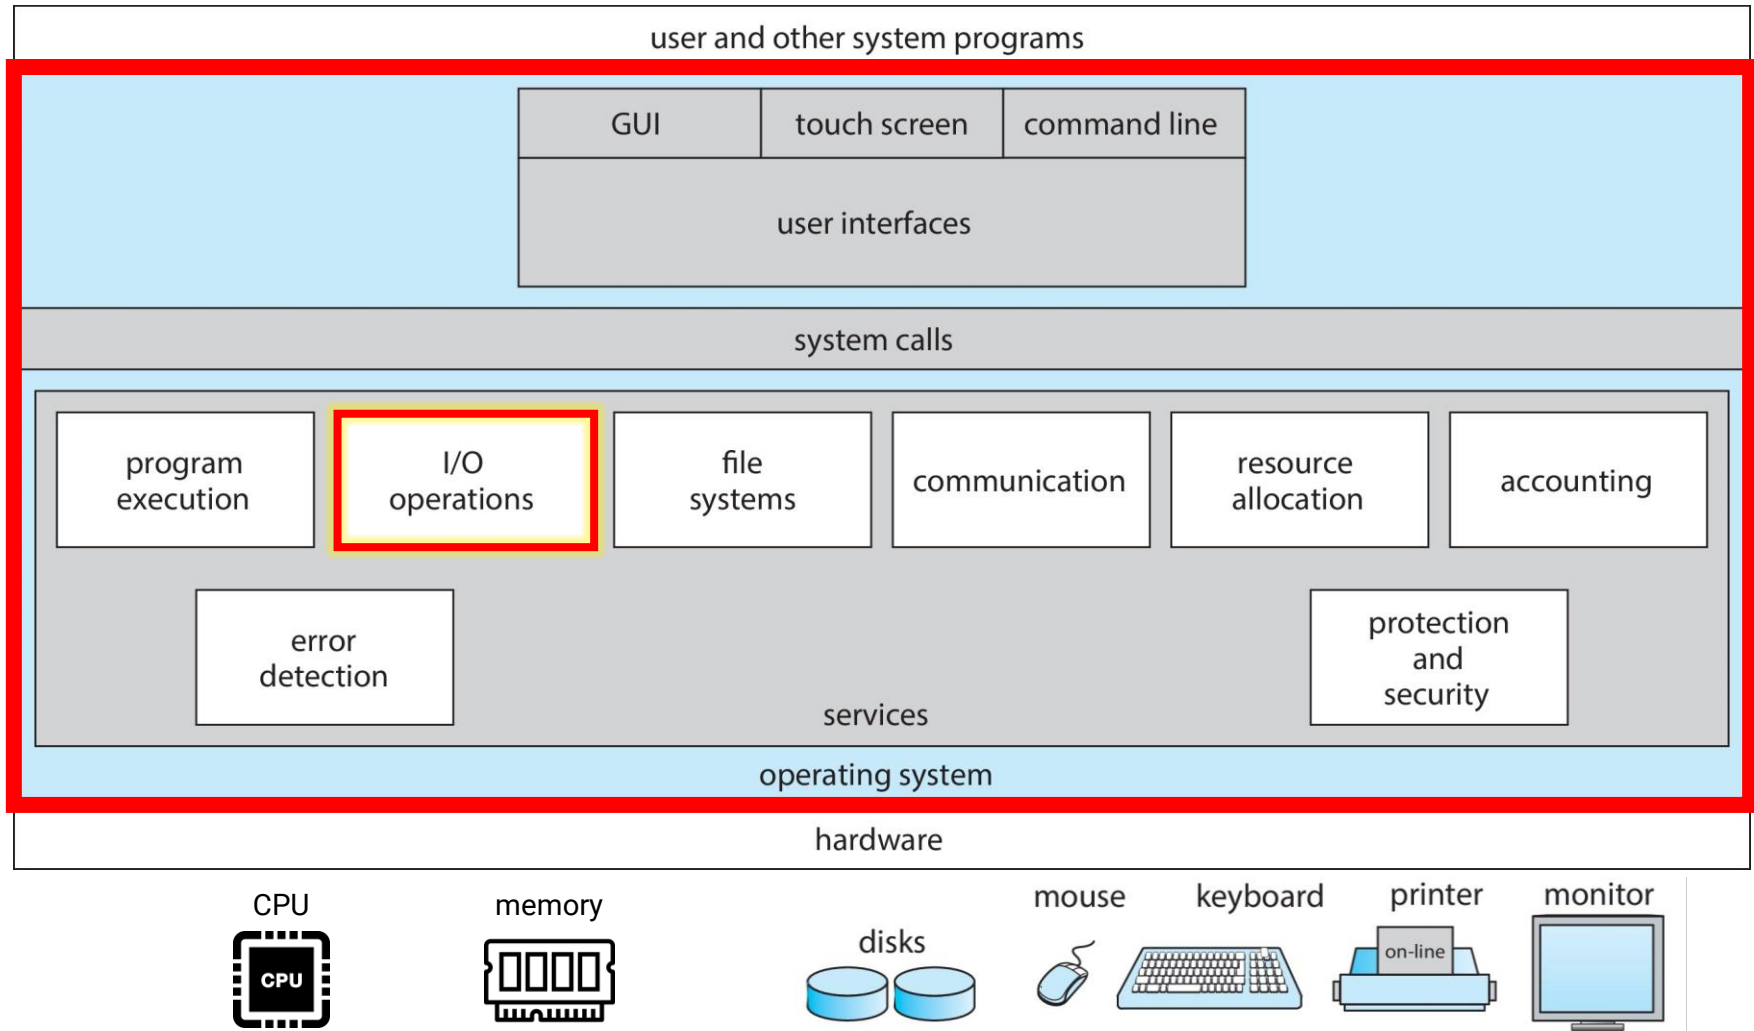
- **Advantages**

- Provide **complete protection** of system resources
- Provide an approach to solve **system compatibility** problems
- Provide a vehicle for **OS research and development**
- Provide a mean for increasing **resource utilization** in cloud computing



# **Operating System Service: I/O Operations**

# Operating System Services



# Strategies for Handling I/O

- **Interrupted-based I/O**


- For slow devices
- OS is informed when jobs have been done
- Example: disk

- **Programmed I/O (pulling)**

- Keep asking if the jobs have been done
- Example: network interface card

- **Memory-mapped I/O**

- Frequently used or very fast devices
- Special I/O instructions are used to move data between memory & device controller registers
- Example: GPU



# DMA: Direct Memory Access

- **Goal**

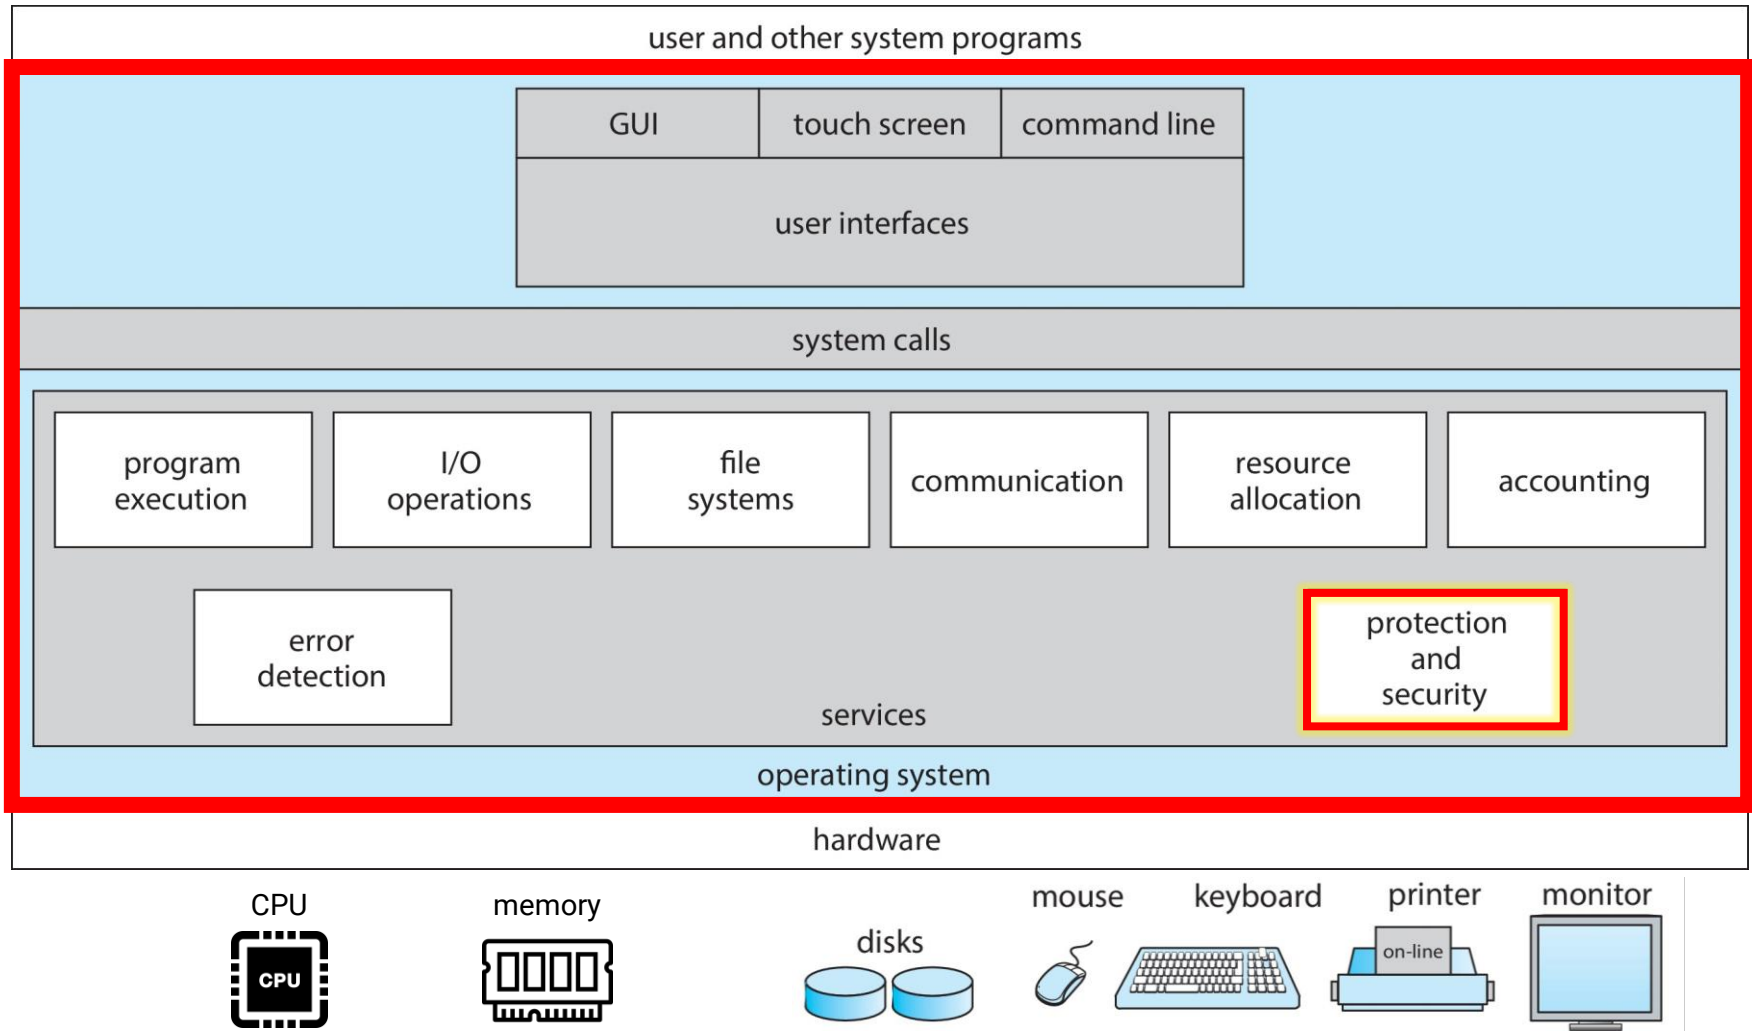
- Device controller can transfer **blocks of data** from buffer storage to main memory **without CPU intervention**
- Only one interrupt is generated per block (rather than per byte), thus avoiding CPU handling excessive interrupts

- **Procedure with DMA**

- Execute the device driver to setup the registers of the DMA controller
- DMA moves blocks of data between the memory and its own buffers
- Transfer from its buffers to its devices
- Interrupt the CPU when the job is done

# **Operating System Service: Protection and Security**

# Operating System Services



# Protection and Security

- **Goal**

- Prevent error and misuse
- Resources are only allowed to be accessed by authorized processes

- **Protection**

- Any mechanism for controlling the access of processes or users to the resources defined by the computer system

- **Security**

- Defense of a system from external and internal attacks
- Examples: viruses, denial of service, identity theft

# Protection

- **Dual-mode operations**

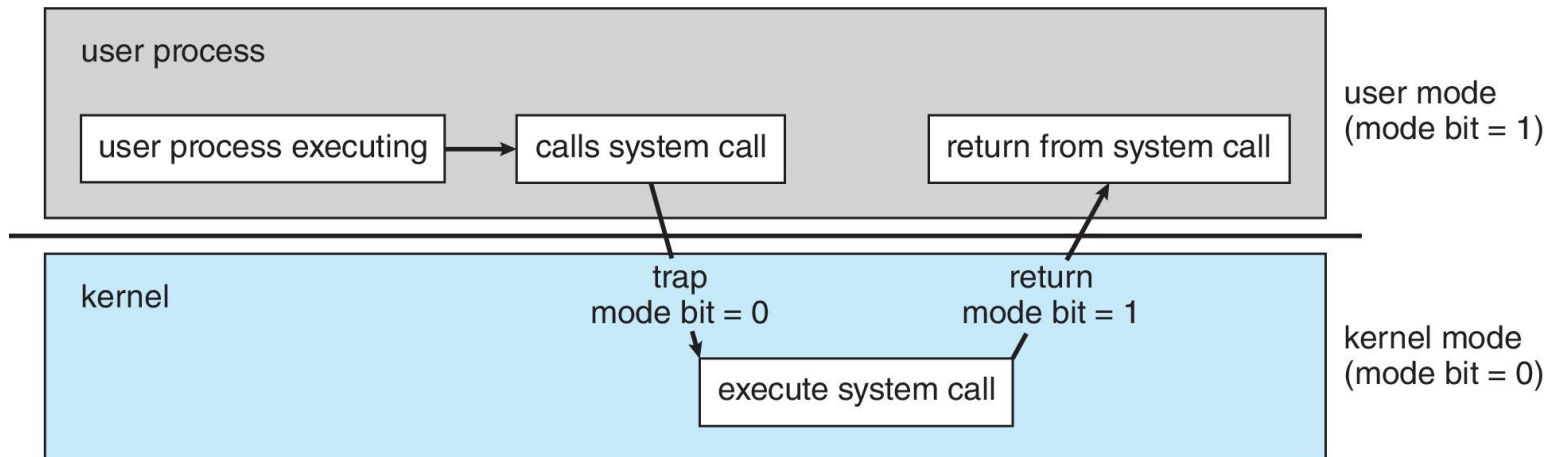
- **User mode**

- Executions except those after a trap or an interrupt occurs

- **Monitor mode (system mode, privileged mode)**

- Can execute all instructions including privileged ones (machine instructions that may cause harm)

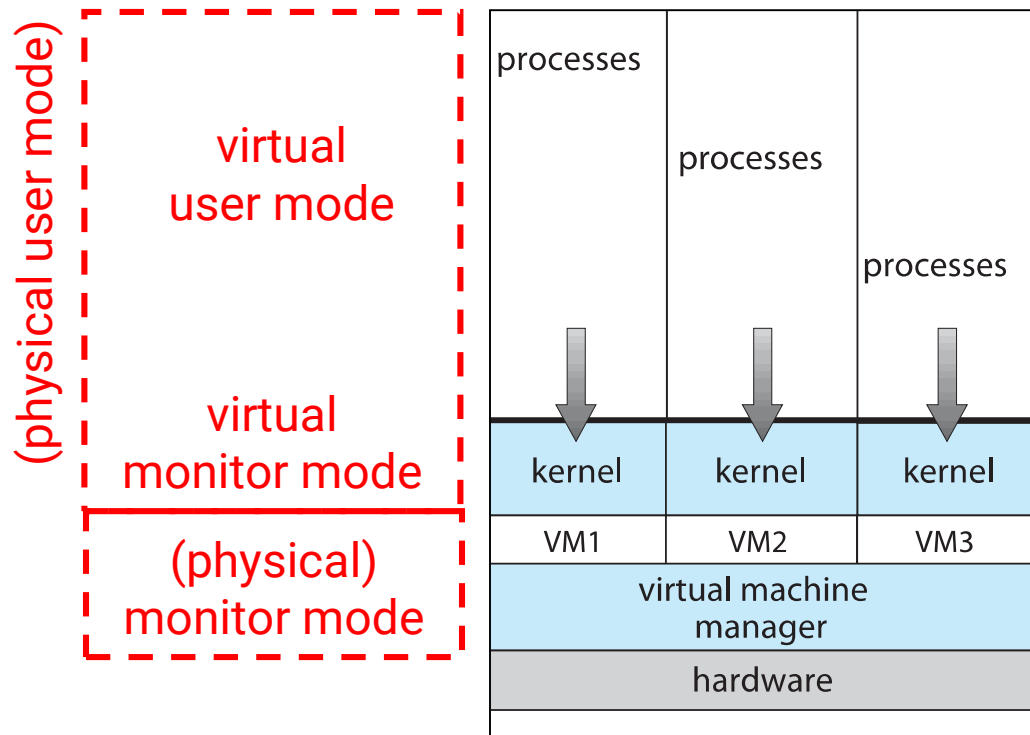
- Implemented by a **mode bit** and **system calls**





# Protection (cont.)

- **Virtual machine** has more than two modes



# Protection (cont.)

- **I/O protection**

- I/O devices are scarce resources, user programs must issue I/O through OS
- Examples: fopen (open), gets (read), puts (write)

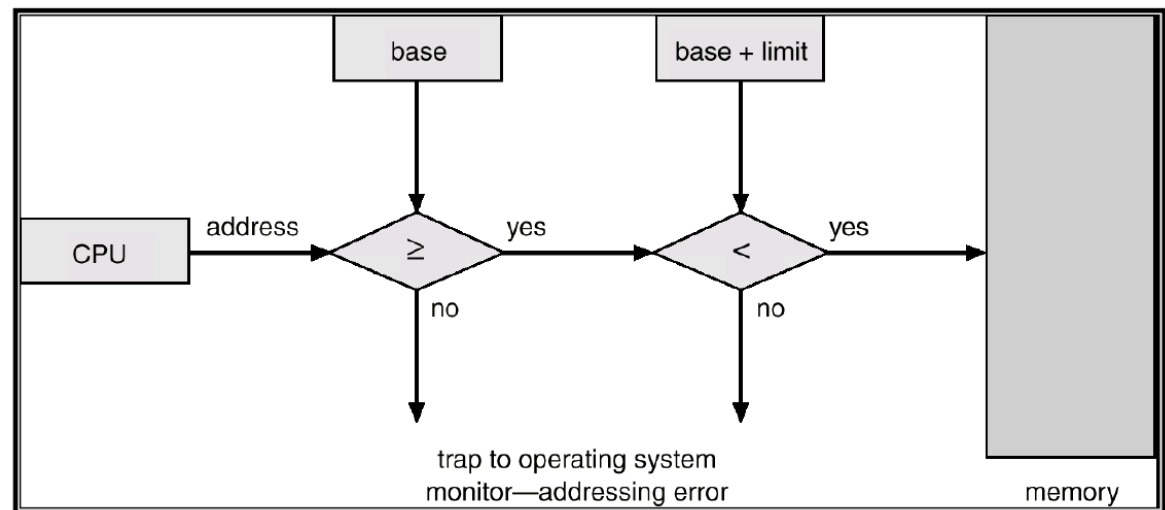
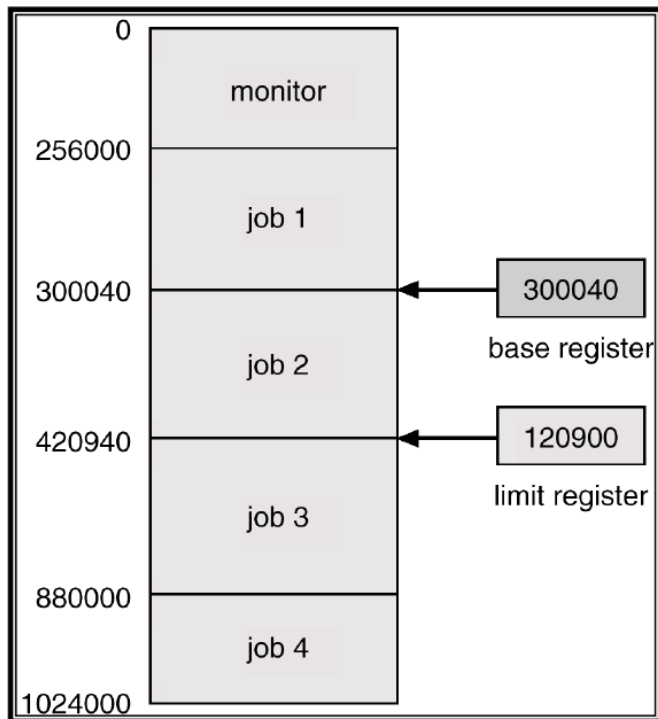
- **Memory protection**

- Prevent a user program from modifying the code or data structures of either the OS or other users
- Example: instructions to modify the memory space

- **CPU protection**

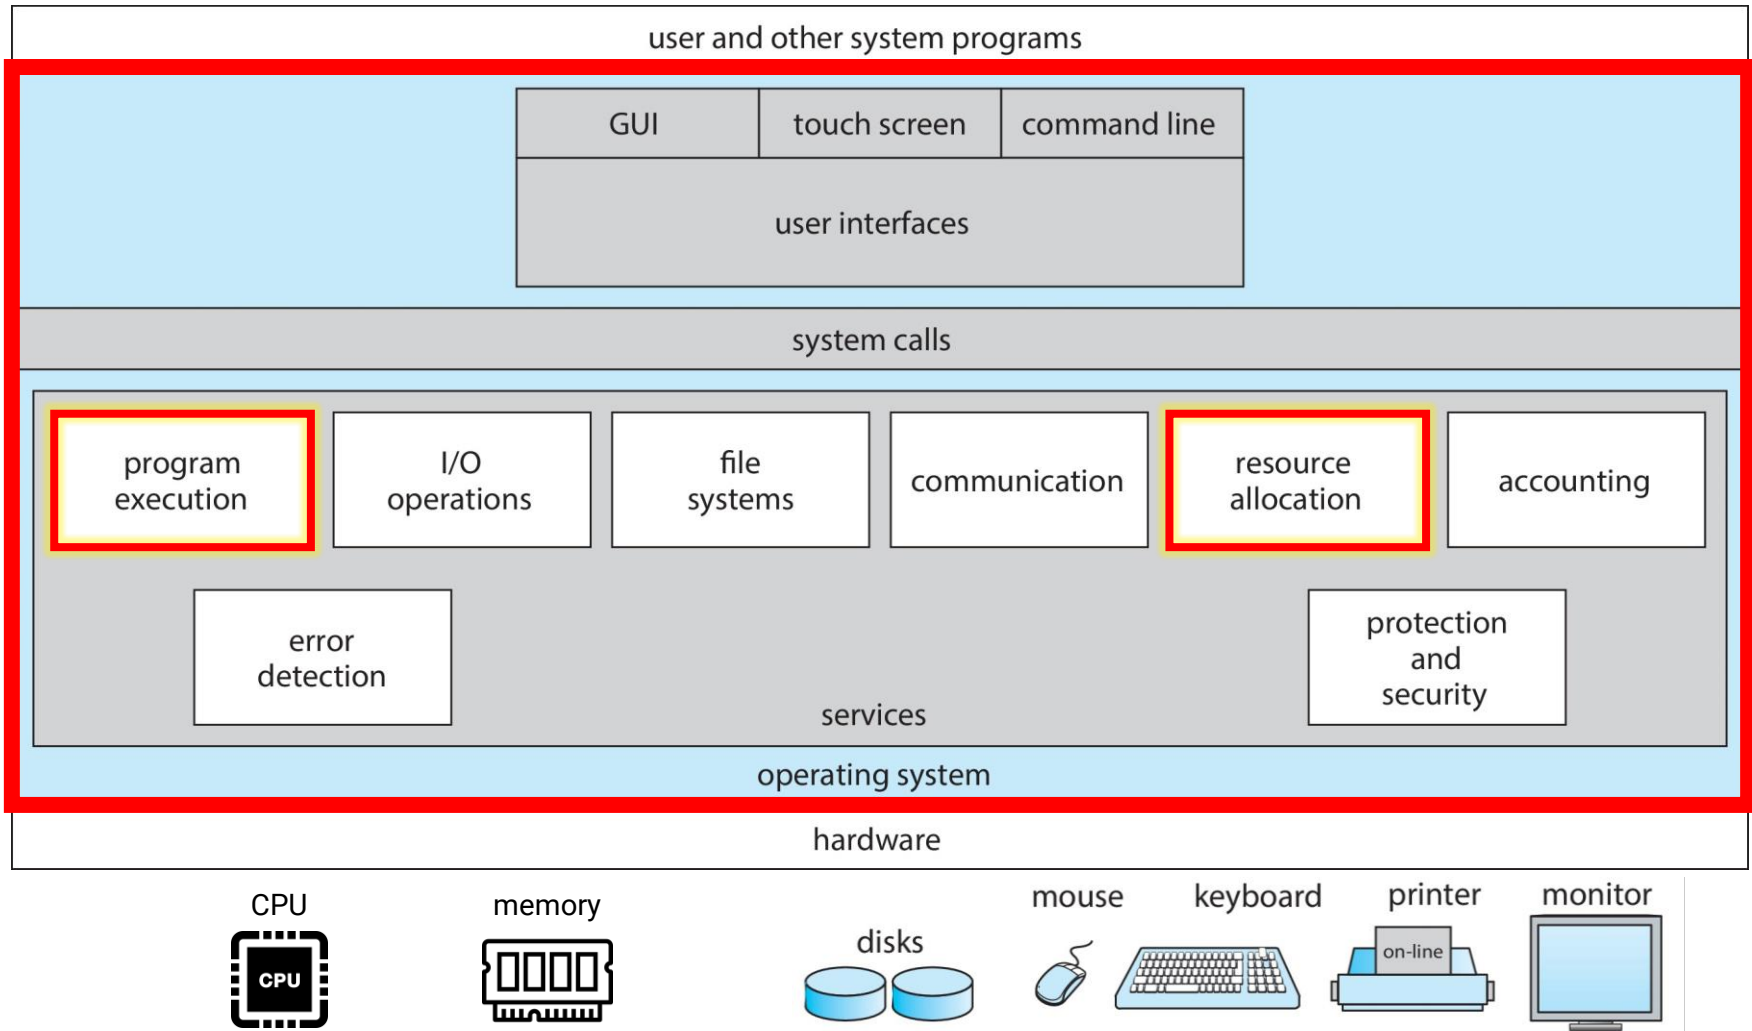
- Prevent user programs from sucking up CPU power
- Implement by **timers** and **time-sharing**
  - Need context switch

# Protection (cont.)



# **Operating System Service: Program Execution & Resource Allocation**

# Operating System Services




# Process Concepts

# Process Concepts

- An operating system concurrently executes a variety of programs
  - **Program: passive entity**, binary file stored **in disk**
  - **Process: active entity**, a running program **in memory**
- A process includes
  - **Code segment** (text section)
  - **Data section**: global variables
  - **Stack**: temporary local variables and functions
  - **Heap**: dynamic allocated variables or classes
  - **Current activity** (e.g., program counter, register contents)
  - **Associated resources** (e.g., handlers of open files)

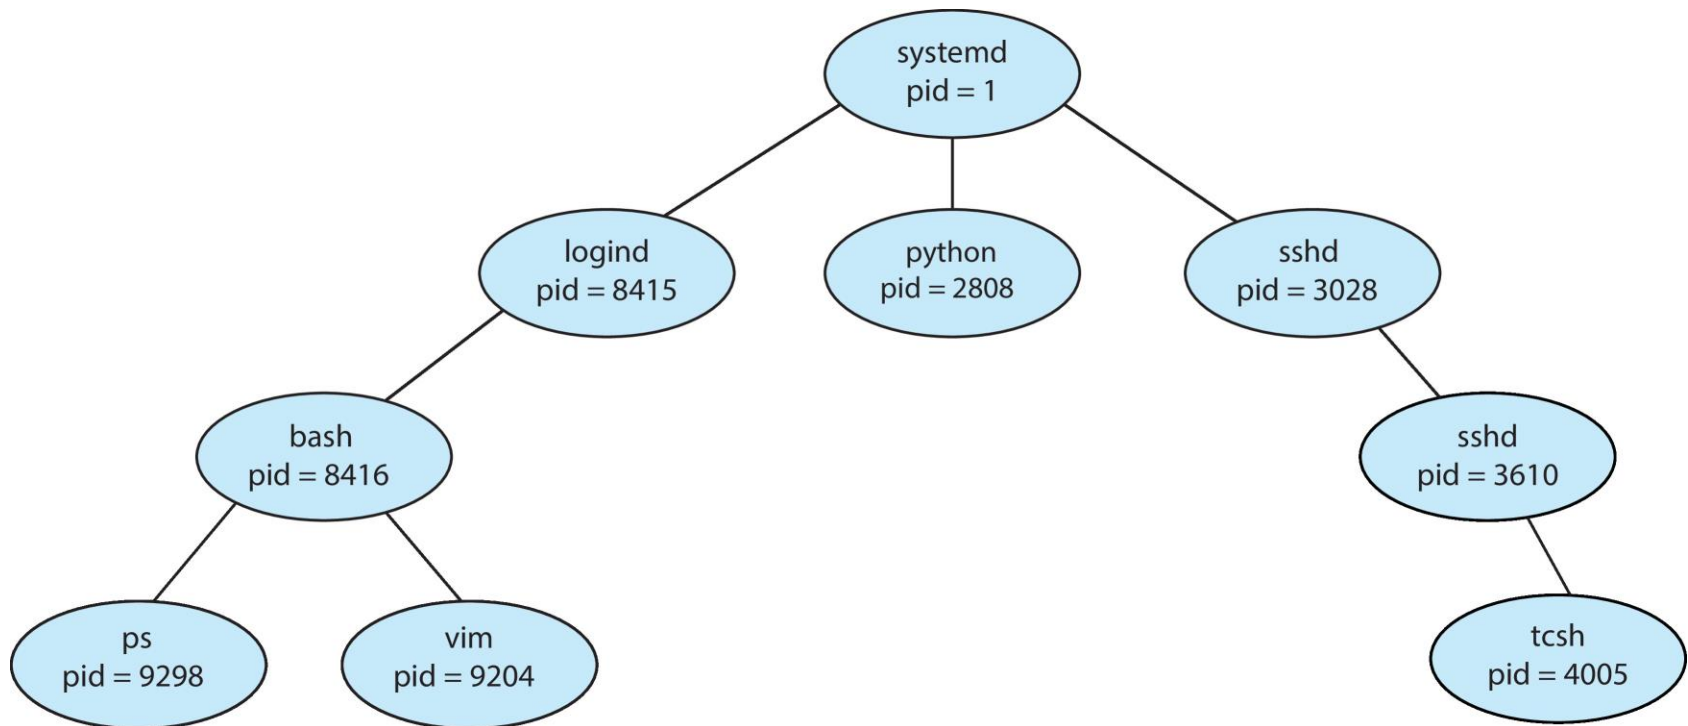
# Process in Memory





# Tree of Processes

- Each process is identified by a unique **processor identifier (pid)**



# Process Creation

- **Resource sharing** (three possibilities)
  - Parent and child processes share **all** resources
  - Child process shares **subset** of parent's resources
  - Parent and child share **no** resources
- **Execution order** (two possibilities)
  - Parent and children execute **concurrently**
  - Parent **waits until children terminate**
- **Address space** (two possibilities)
  - Children **duplicate** of parent, communication via sharing variables
  - Child **has a program loaded into it**, communication via message passing

# UNIX / Linux Process Creation

- **fork system call**

- Create a new (child) process
- The new process **duplicates** the address space of its parent
- Child and parent **execute concurrently** after fork
- Child: return value of fork is 0
- Parent: return value of fork is PID of the child process

- **execvp system call**


- Load a new binary file into memory
- Destroy the old code

- **wait system call**

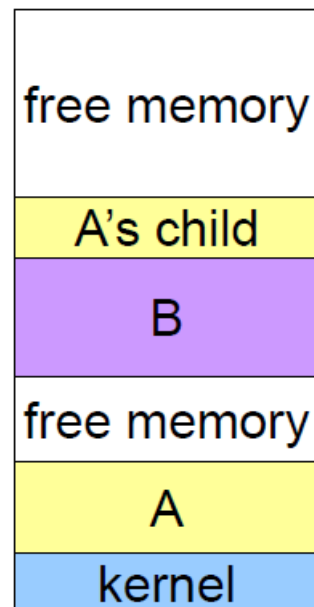

- The parent **waits** for one of its child processes to complete

# UNIX / Linux Process Creation (cont.)

- Memory space of fork()
  - Old implementation: A's child is an exact copy of parent
  - Current implementation: use **copy-on-write** technique to store differences in A's child address space



originally

after A does a **fork**after the child does an **execp**

# UNIX / Linux Example

```
#include <sys/types.h>
#include <stdio.h>
#include <unistd.h>

int main()
{
    pid_t pid;

    /* fork a child process */
    pid = fork();


    if (pid < 0) { /* error occurred */
        fprintf(stderr, "Fork Failed");
        return 1;
    }
    else if (pid == 0) { /* child process */
        execlp("/bin/ls", "ls", NULL);
    }
    else { /* parent process */
        /* parent will wait for the child to complete */
        wait(NULL);
        printf("Child Complete");
    }

    return 0;
}
```

# UNIX / Linux Example Quiz

- How many processors are created?

```
#include <stdio.h>
#include <unistd.h>
int main()
{
    for (int i = 0; i < 3 ; i++)
        fork();
    return 0;
}
```



# Process Termination

- Terminate when the **last statement is executed** or **exit()** is called
  - Return status data from child to parent
  - All resources of the process, including physical and virtual memory, open files, I/O buffers, are deallocated by the OS
- Parent may terminate execution of children processes by specifying its PID (**abort**)
  - Children has exceeded allocated resources
  - Task assigned to child is no longer required
  - The parent is exiting, and the OS does not allow a child to continue if its parent terminates


# Thread Concepts




# Thread Concepts

- A thread is a **lightweight** process
  - Basic unit of CPU utilization
- All threads belonging to the same process **share**
  - Code section
  - Data section (including dynamic allocated variables)
  - OS resources (e.g., open files)
- **But each thread has its own** (thread control block)
  - Thread ID
  - Program counter
  - Register set
  - Stack

# Thread Concept (cont.)



single-threaded process



multithreaded process

# Benefits of Threads

- **Responsiveness**

- Allow continued running of a process even if part of it is blocked or is performing a lengthy operation

- **Resource sharing**

- Thread share resources of process

- **Scalability**

- Take advantage of multicore architectures


- **Economy**

- Thread creation is cheaper than process creation
  - Ex: creating a thread is 30x cheaper than a process in Solaris
- Thread switching has lower overhead than context switching

# Parallelism


- **Data parallelism**

- Distribute subsets of the same data across multiple cores
- Same operation on each



- **Task parallelism**

- Distribute threads across cores, each thread performing unique operation



# Amdahl's Law

- Identify performance gains from adding additional cores to an application that has both serial and parallel components
- Assume **S** is the serial portion and the system has **N** processing cores

$$speedup \leq \frac{1}{S + \frac{(1-S)}{N}}$$

- Example: 75% parallel / 25% serial, moving from 1 to 2 cores results in a speedup of 1.6 times
- As  $N$  approaches infinity, speedup approaches  $1/S$

# User Threads and Kernel Threads

- **User threads**

- Thread management done by user-level threads library
- Ex: POSIX Pthreads, Win32 threads, Java threads


- **Kernel threads**

- Supported by kernel (OS) directly
- Ex: Windows 2000 (NT), Solaris, Linux, Tru64 UNIX

- Programmers create **user threads** using thread APIs and then the threads are bounded to **kernel threads** by the OS

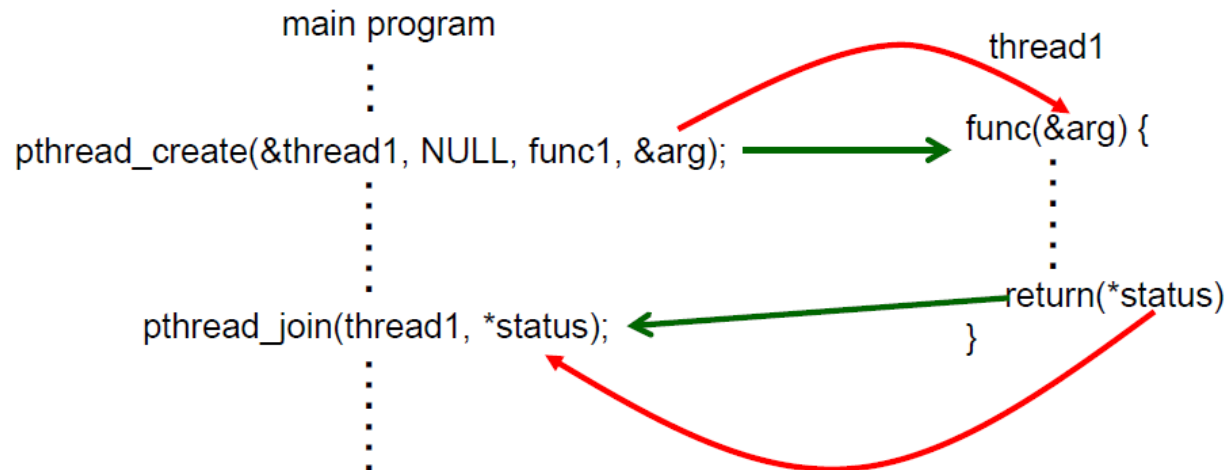
# Multi-Threading Models

- **Many-to-One**
- **One-to-One**
- **Many-to-Many**



# Pthread Creation

- Library call: ***pthread\_create(thread, attr, routine, arg)***
  - *thread*: an unique identifier (token) for the new thread
  - *attr*: used to set thread attributes
  - *routine*: the routine that the thread will execute once it is created
  - *arg*: a single argument that may be passed to routine






# Pthread Joining and Detaching

- Library call: ***pthread\_join(threadId, status)***
  - Blocks until the specified thread (threadId) terminates
  - One way to accomplish synchronization between threads
  - Example:  

```
for (int i = 0; i < n ; ++i) pthread_join(thread[i], NULL);
```
- Library call: ***pthread\_detach(threadId)***
  - Once a thread is detached, it can never be joined
  - Detach a thread could free some system resources



# Linux Kernel Threads

- Linux implements threads by **extending** processes
- **The fork system call**
  - Create a new process and a copy of the associated data of the parent process
- **The clone system call**
  - Create a new process and a link that points to the associated data of the parent process

# Linux Threads (cont.)


- A set of **flags** is used in the clone system call for indication of the level of the sharing
  - None of the flags is set → clone = fork
  - All flags are set → parent and child share everything

flag	meaning
CLONE_FS	File-system information is shared.
CLONE_VM	The same memory space is shared.
CLONE_SIGHAND	Signal handlers are shared.
CLONE_FILES	The set of open files is shared.

# Process Execution Strategy


# Simple Batch System

- Workflow
  - Users submit data (program, data, control card)
  - Operator sort jobs with similar requirement
  - OS simply transfer control from one job to the next one
    - Resident monitor: automatically transfer control from one job the next



# Simple Batch System (cont.)

- **Problem of batch systems**
  - One job at a time → multi-programming
  - No interaction between users and jobs → time sharing
  - CPU is often idle
- **Spooling** (Simultaneous Peripheral Operation On-Line)
  - Replace sequential-access devices with random-access devices (disks)




- Ex: printer
  - Data for printing is first written into a directory as files, then printed by the printer

# Multi-Programming

- Single user cannot always keep CPU and I/O devices busy
  - Even with spooling, disk I/O is still too slow compared to CPU and memory
- **Put multiple programs in memory**
- **OS organizes jobs** so that the CPU always has one to execute
  - When job has to wait (e.g., for I/O), OS switches to another job according to a **scheduling algorithm**
  - Increase CPU utilization

# Time-Sharing (Multi-Tasking)

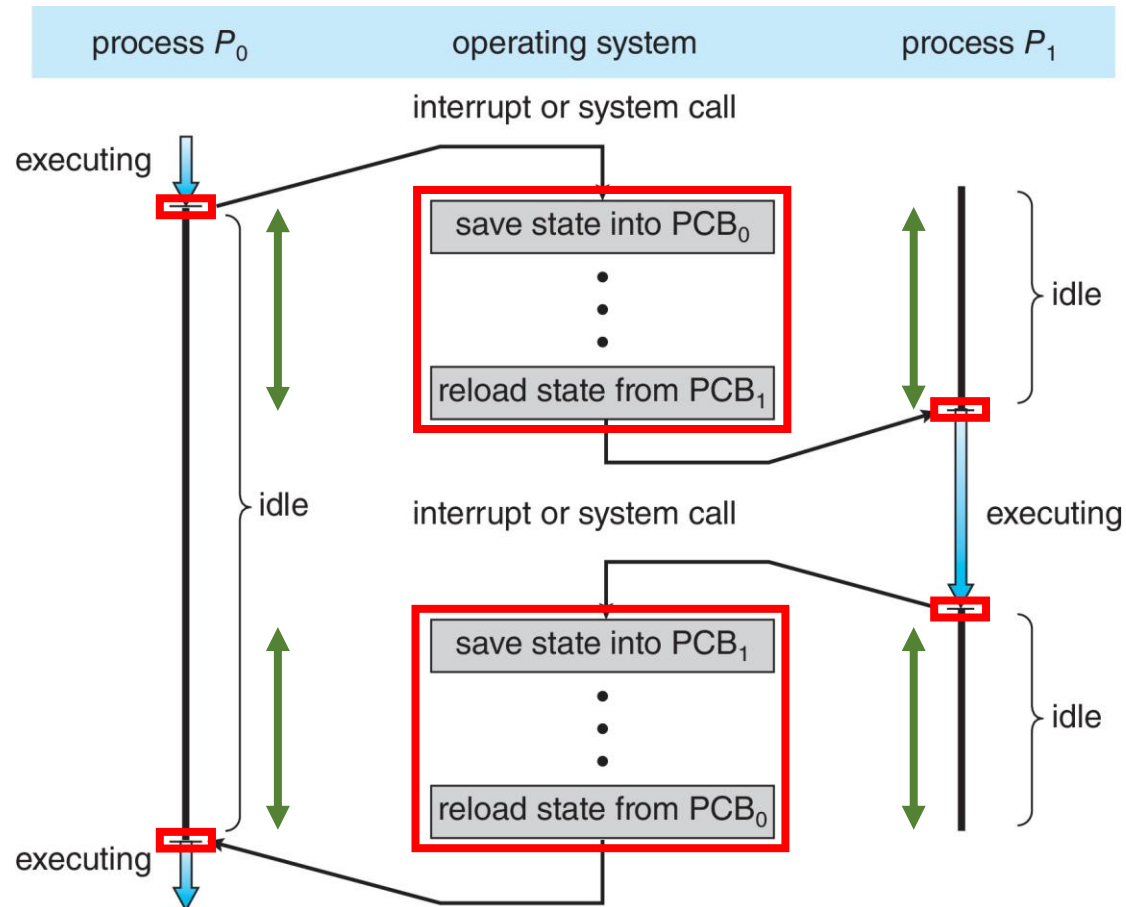
- CPU switches jobs frequently so that users can interact with each job while it is running
  - A logical extension of multi-programming
  - **Interactivity** !
    - Response time should be less than 1 sec.





# Context Switch

- Occurs when the CPU switches from one process to another

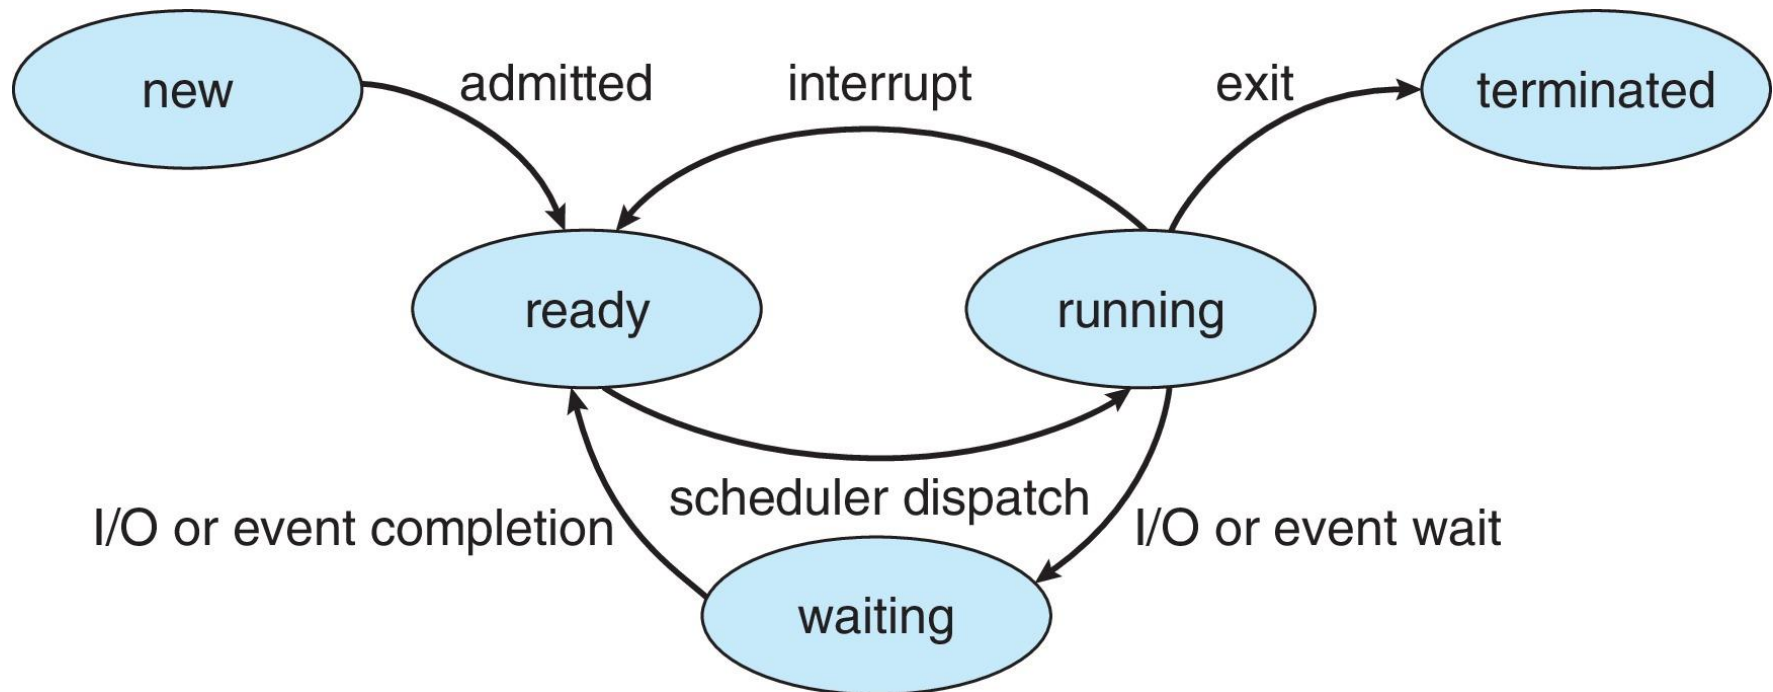


# Process Scheduling Overview

# Process State

- Types of states
  - **New**
    - The process is being created
  - **Ready**
    - The process is in the memory waiting to be assigned to a processor
  - **Running**
    - The process whose instructions are being executed by CPU
  - **Waiting**
    - The process is waiting for events to occur
  - **Terminated**
    - The process has finished execution

# Process State (cont.)



**Only one** process is running on any processor at any instant  
However, many processes may be ready or waiting (put into a queue)

# Process Schedulers

- **Short-term scheduler (CPU scheduler)**

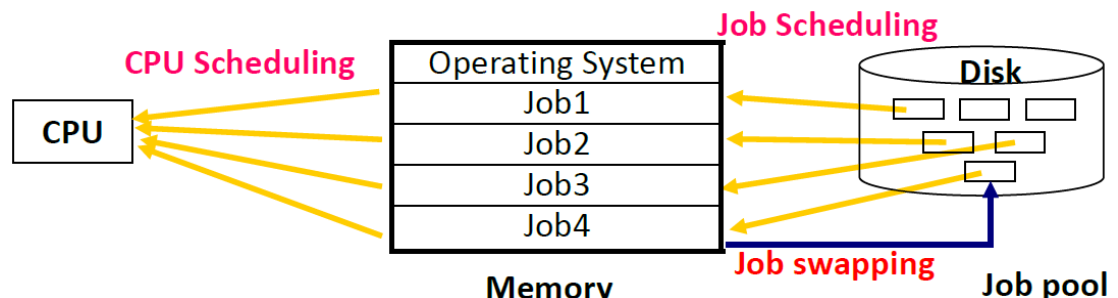
- Select which process should be executed and allocated CPU  
(Ready state → Running state)

- **Long-term scheduler (job scheduler)**

- Select which processes should be loaded into memory and brought into the ready queue (New state → Ready state)

- **Medium-term scheduler**

- Select which processes should be swapped in/out memory  
(Ready state → Waiting state)

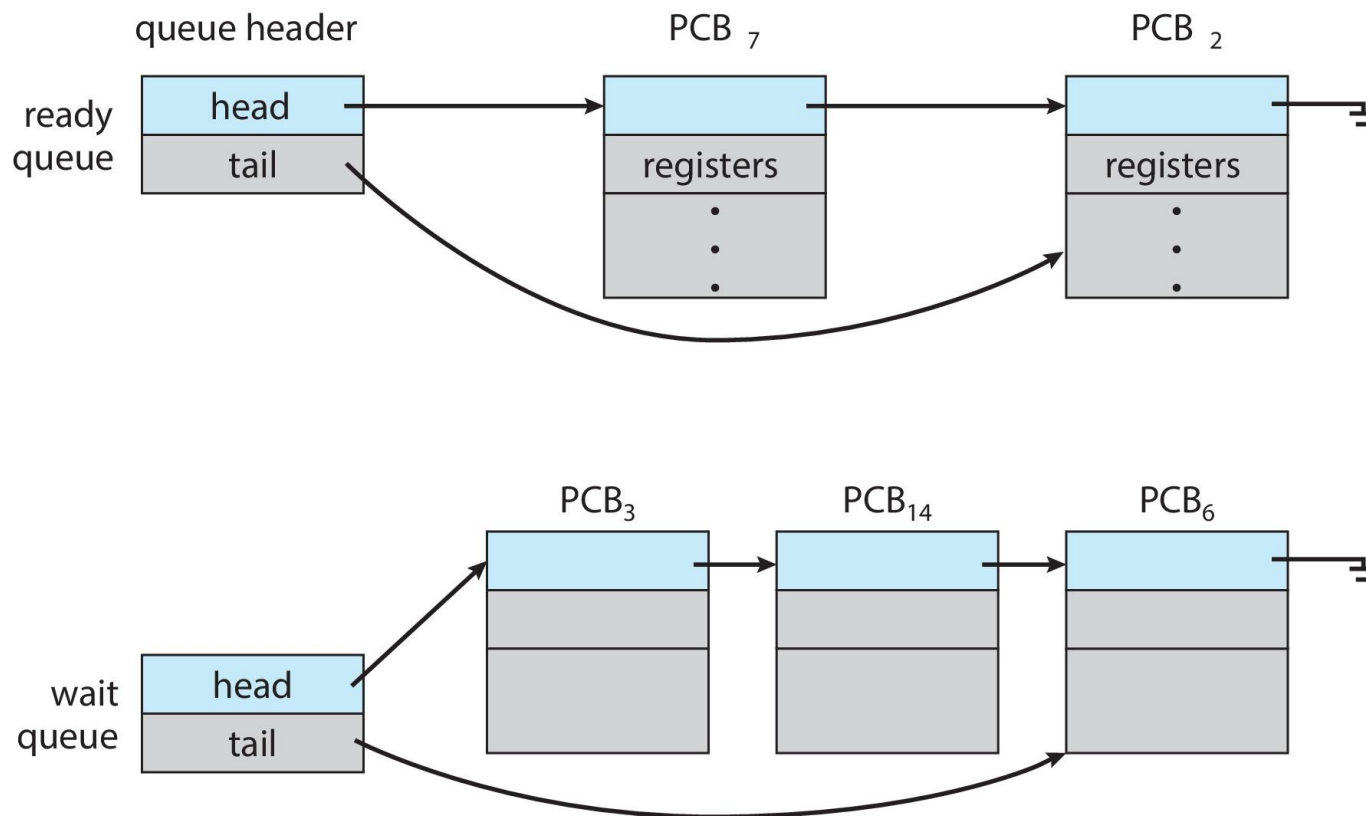


# Process Scheduling Queues

- Maintain **scheduling queues** of processes
  - **Job queue (New state)**
    - Set of all processes in the system
  - **Ready queue (Ready state)**
    - Set of all processes residing in main memory
    - Ready and waiting to execute
  - **Waiting queue (Wait State)**
    - Set of processes waiting for an event (e.g., I/O)
- Processes migrate among the various queues

# Process Scheduling Queues (cont.)

- Ready queue and waiting queue



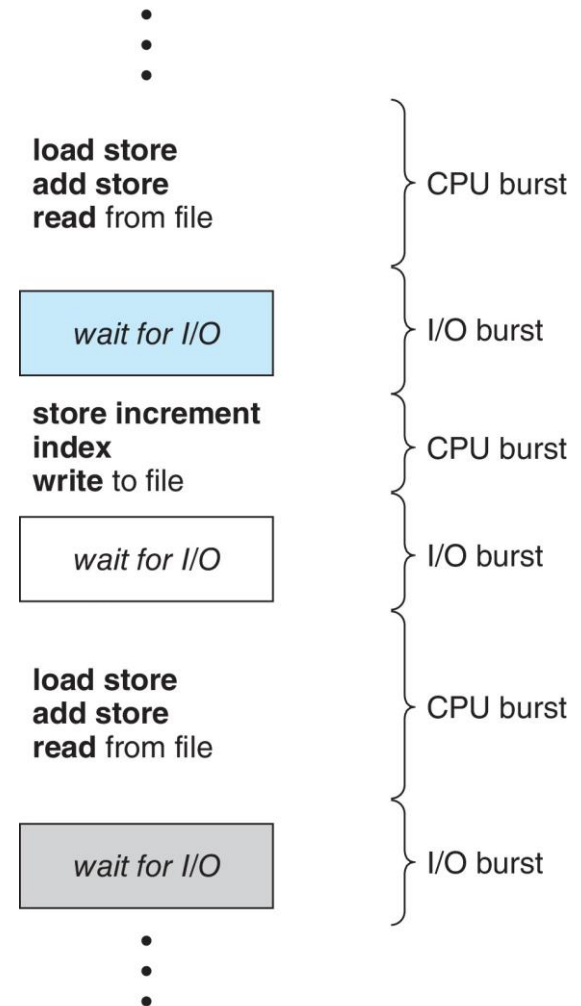
# CPU Scheduling Algorithms



# Basic Concepts

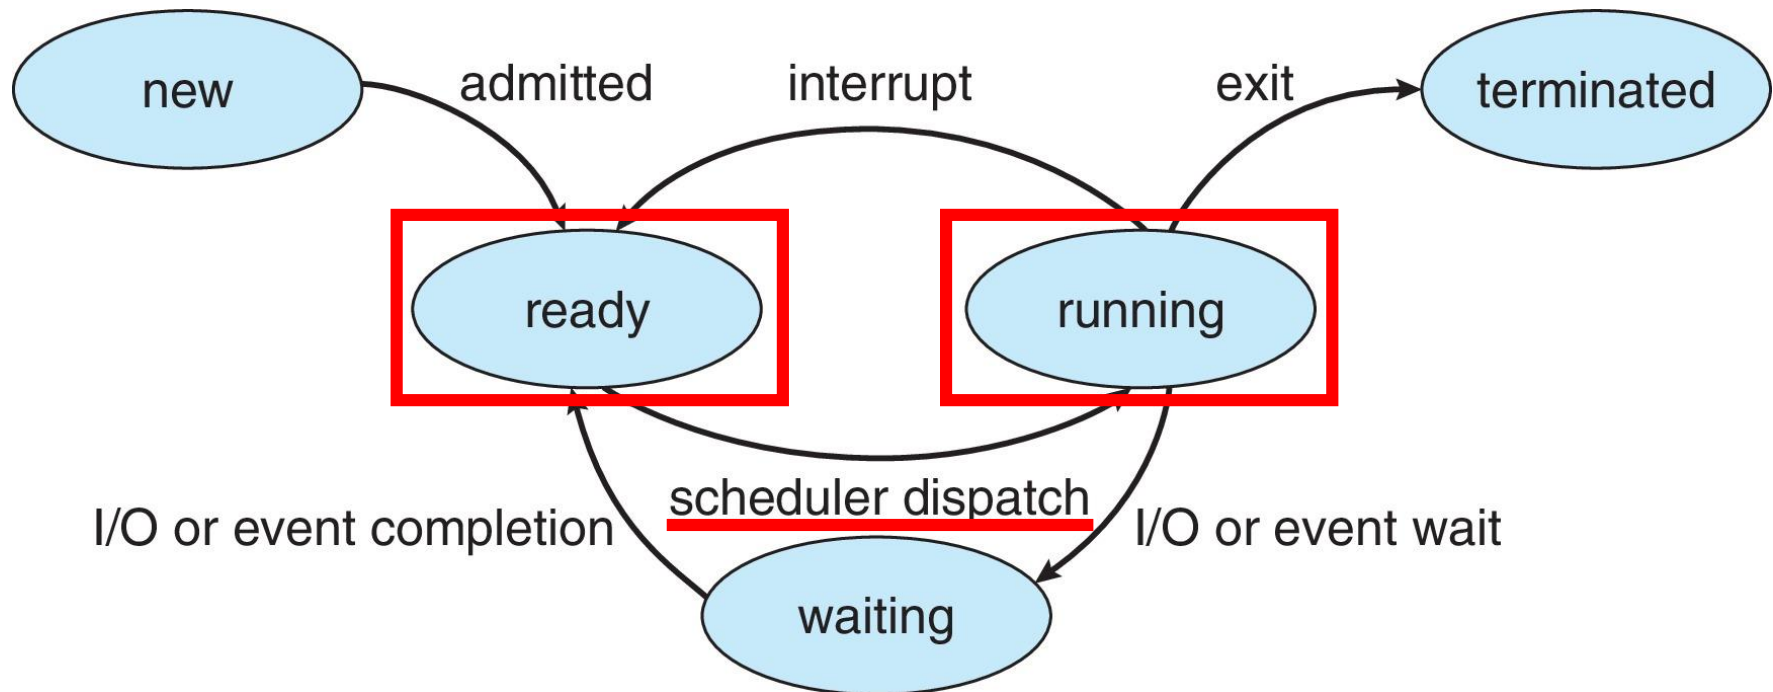
## • CPU-I/O burst cycle

- Process execution consists of a cycle of **CPU execution (CPU burst)** and **I/O wait (I/O burst)**
  - Generally, there is a large number of short CPU bursts, and a small number of long CPU bursts
  - An I/O-bound program would typically have many very short CPU bursts
  - A CPU-bound program might have a few long CPU bursts



# CPU Scheduler

- Selects process from ready queue to execute
  - Allocate a CPU for the selected process



# Preemptive vs Non-preemptive

- CPU scheduling decisions may take place when a process
  - Switches from running to waiting state
  - Switches from running to ready state
  - Switches from waiting to ready
  - Terminates
- **Non-preemptive scheduling**
  - Scheduling under 1 and 4 (no choice in terms of scheduling)
  - The process keeps the CPU until it is **terminated** or **switched to the waiting state**
- **Preemptive scheduling**
  - Scheduling under all cases
  - E.g., Windows 95 and subsequent versions, Mac OS X

# Scheduling Criteria

- **CPU utilization**

- Theoretically 0% ~ 100%
- Real systems: 40% (light) ~ 90% (heavy)

- **Throughput**

- Number of completed processes per time unit

system view

- **Turnaround time**

- Submission ~ completion

- **Waiting time**

- Total waiting time in the ready queue

- **Response time**

- Submission ~ the first response is produced

single job view

# Algorithms

- First-Come, First-Served (FCFS) scheduling
- Shortest-Job-First (SJF) scheduling
- Priority scheduling
- Round-Robin scheduling
- Multi-level queue scheduling
- Multi-level feedback queue scheduling

# FCFS Scheduling

- Process (burst time) in arriving order
  - P1 (24), P2 (3), P3 (3)
- The Gantt Chart of the schedule



- Waiting time: P1 = 0, P2 = 24, P3 = 27
- Average Waiting Time (AWT):  $(0 + 24 + 27) / 3 = 17$
- **Convoy effect**
  - Short processes behind a long process

# Shortest-Job-First (SJF) Scheduling

- Associate with each process the length of its next CPU burst
- A process with shortest burst length gets the CPU first
- SJF provides the **minimum (optimal) average waiting time**
- Two schemes
  - **Non-preemptive**
    - Once the CPU is given to a process, it cannot be preempted until its completion
  - **Preemptive**
    - If a new process arrives with shorter burst length, preemption happens

# Preemptive SJF Example

Process	Arrival Time	Burst Time
P1	0	7
P2	2	4
P3	4	1
P4	5	4

Ready queue at t = 11

P1 (5)

Scheduling



- **Wait time = completion time - arrival time - run time**
- **AWT** =  $[(16-0-7) + (7-2-4) + (5-4-1) + (11-5-4)] / 4 = 3$
- **Response time:** P1 = 0, P2 = 0, P3 = 0, P4 = 2



# Approximate Shortest-Job-First (SJF)

- **SJF difficulty**

- No way to know length of the next CPU burst

- **Approximate SJF**

- The next burst can be predicted as an exponential average of the measured length of previous CPU bursts


$$\tau_{n+1} = \alpha \underbrace{t_n}_{\text{new one}} + (1 - \alpha) \underbrace{\tau_n}_{\text{history}}$$

$$= \alpha t_n + (1 - \alpha) \alpha t_{n-1} + (1 - \alpha)^2 \alpha t_{n-2} + \dots$$

Example:

$$\alpha = 1/2 \quad = \left(\frac{1}{2}\right)t_n + \left(\frac{1}{2}\right)^2 t_{n-1} + \left(\frac{1}{2}\right)^3 t_{n-2}$$

# Exponential Prediction of Next CPU Burst



CPU burst ( $t_i$ )		6	4	6	4	13	13	13	...
"guess" ( $\tau_i$ )	10	8	6	6	5	9	11	12	...

# Priority Scheduling

- A **priority number** is associated with each process
- The CPU is allocated to the highest priority process
  - Preemptive
  - Non-preemptive
- **SJF is a priority scheduling** where priority is the predicted next CPU burst time
- Problem: **starvation**
  - Low priority processes never execute
  - Example: IBM 7094 shutdown at 1973, a 1967-process never run
  - Solution: **aging**
    - As time progresses, increase the priority of processes

# Round-Robin (RR) Scheduling


- Each process gets a small unit of CPU time (**time quantum**, TQ), usually 10 ~ 100 ms
  - Context switch time usually < 10 microseconds
- After TQ elapsed, process is **preempted** and **added to the end of the ready queue**
- If there are  $n$  processes in the ready queue and the time quantum is  $q$ , each process gets  $1/n$  of the CPU time ( $q$  time units)
  - No process waits more than  $(n-1)q$  time units
- Performance
  - TQ large → FCFS
  - TQ small → (context switch) overhead increases

# Priority Scheduling with Round-Robin

Process	Burst Time	Priority
P1	4	3
P2	5	2
P3	4	2
P4	7	1
P5	3	3

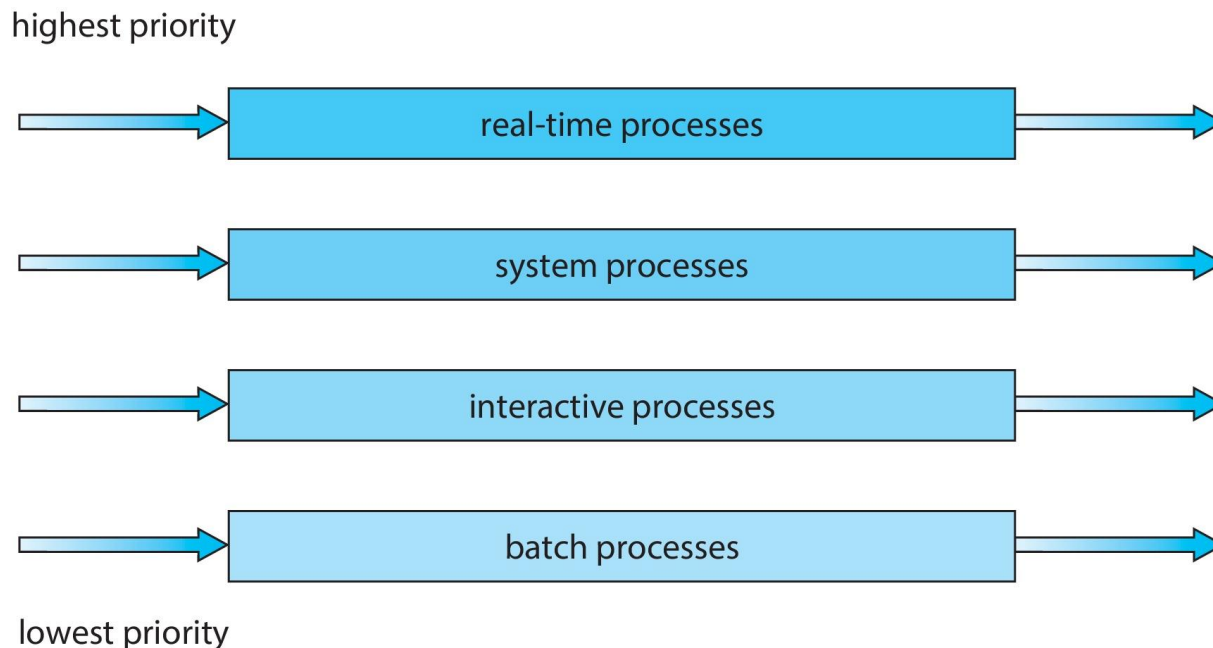
- Run the process with the highest priority
- Processes with the same priority run round-robin

Gantt Chart (TQ = 2)




# Multi-level Queue Scheduling

- Ready queue is partitioned into separate queues
- Each queue has its own scheduling algorithm
- Scheduling must be done **between** queues
  - Time slice: each queue gets a certain amount of CPU



# Multi-level Feedback Queue Scheduling


- A process can move between the various queues
  - Aging must be implemented
- Idea: separate processes according to the characteristic of their CPU burst
  - **I/O-bound** and **interactive processes** in **higher priority** queue → short CPU burst
  - **CPU-bound** processes in **lower priority** queue → long CPU burst



# Multi-level Feedback Queue (cont.)

- Three queues

- Q0: RR with TQ 8 ms.
- Q1: RR with TQ 16 ms.
- Q2: FCFS




- Scheduling

- A new process enters queue Q0 which is served in RR
  - When it gains CPU, the process receives 8 ms
  - If it does not finish in 8 ms., the process is moved to queue Q1
- At Q1, job is again served in RR and receives 16 additional ms.
  - If it still does not complete, it is preempted and moved to queue Q2




# Thread Scheduling

## Many-to-One




## Many-to-Many



**Process Contention Scope (PCS) by thread library**

**System Contention Scope (SCS) by OS**

## One-to-One



**System Contention Scope (SCS) by OS**

# Special Scheduling Issues

# Processor Affinity

- **Processor affinity**

- A process has an affinity for the processor on which it is currently running
  - A process populates its recent used data in cache
  - Cache invalidation and repopulation has high cost


- **Solution**

- **Soft affinity**

- Possible to migrate to other processors


- **Hard affinity**

- Not to migrate to other processor



# NUMA and CPU Scheduling

- Occurs in systems containing combined CPU and memory boards
- CPU scheduler and memory-placement works together



# Load Balancing


- Keep the workload evenly distributed across all processors
  - Only necessary on systems where each processor has its own private queue of eligible processes to execute
- **Two strategies**
  - **Push migration**
    - Move (push) processes from overloaded to idle or less-busy processor
  - **Pull migration**
    - Idle processor pulls a waiting task from a busy processor
  - Often implemented in parallel
- Load balancing often counteracts the benefits of processor affinity

# Real-time Scheduling Algorithms

- Must support **preemptive, priority-based** scheduling
  - But only guarantees **soft real-time**
- **Description**
  - $T1 = (0, 4, 10) == (\text{Ready}, \text{Execution}, \text{Period})$
  - $T2 = (1, 2, 4)$
- **Rate-Monotonic (RM) algorithm**
  - **Shorter period  $\rightarrow$  high priority**
  - Fixed-priority real-time system scheduling algorithm
- **Earliest-deadline-first (EDF) algorithm**
  - **Earlier deadline  $\rightarrow$  higher priority**
  - Dynamic priority algorithm


# Rate-Monotonic (RM) Scheduler

- **Fixed-priority scheduling**
  - All jobs of the same task have same priority
  - The task's priority is fixed
- **The shorter period, the higher priority**
- Example (Preempted)
  - $T1 = (4, 1)$ ,  $T2 = (5, 2)$ ,  $T3 = (20, 5)$  for (period, execution)
    - Period:  $4 < 5 < 20$
    - Priority:  $T1 > T2 > T3$



# Early Deadline First (EDF) Scheduler

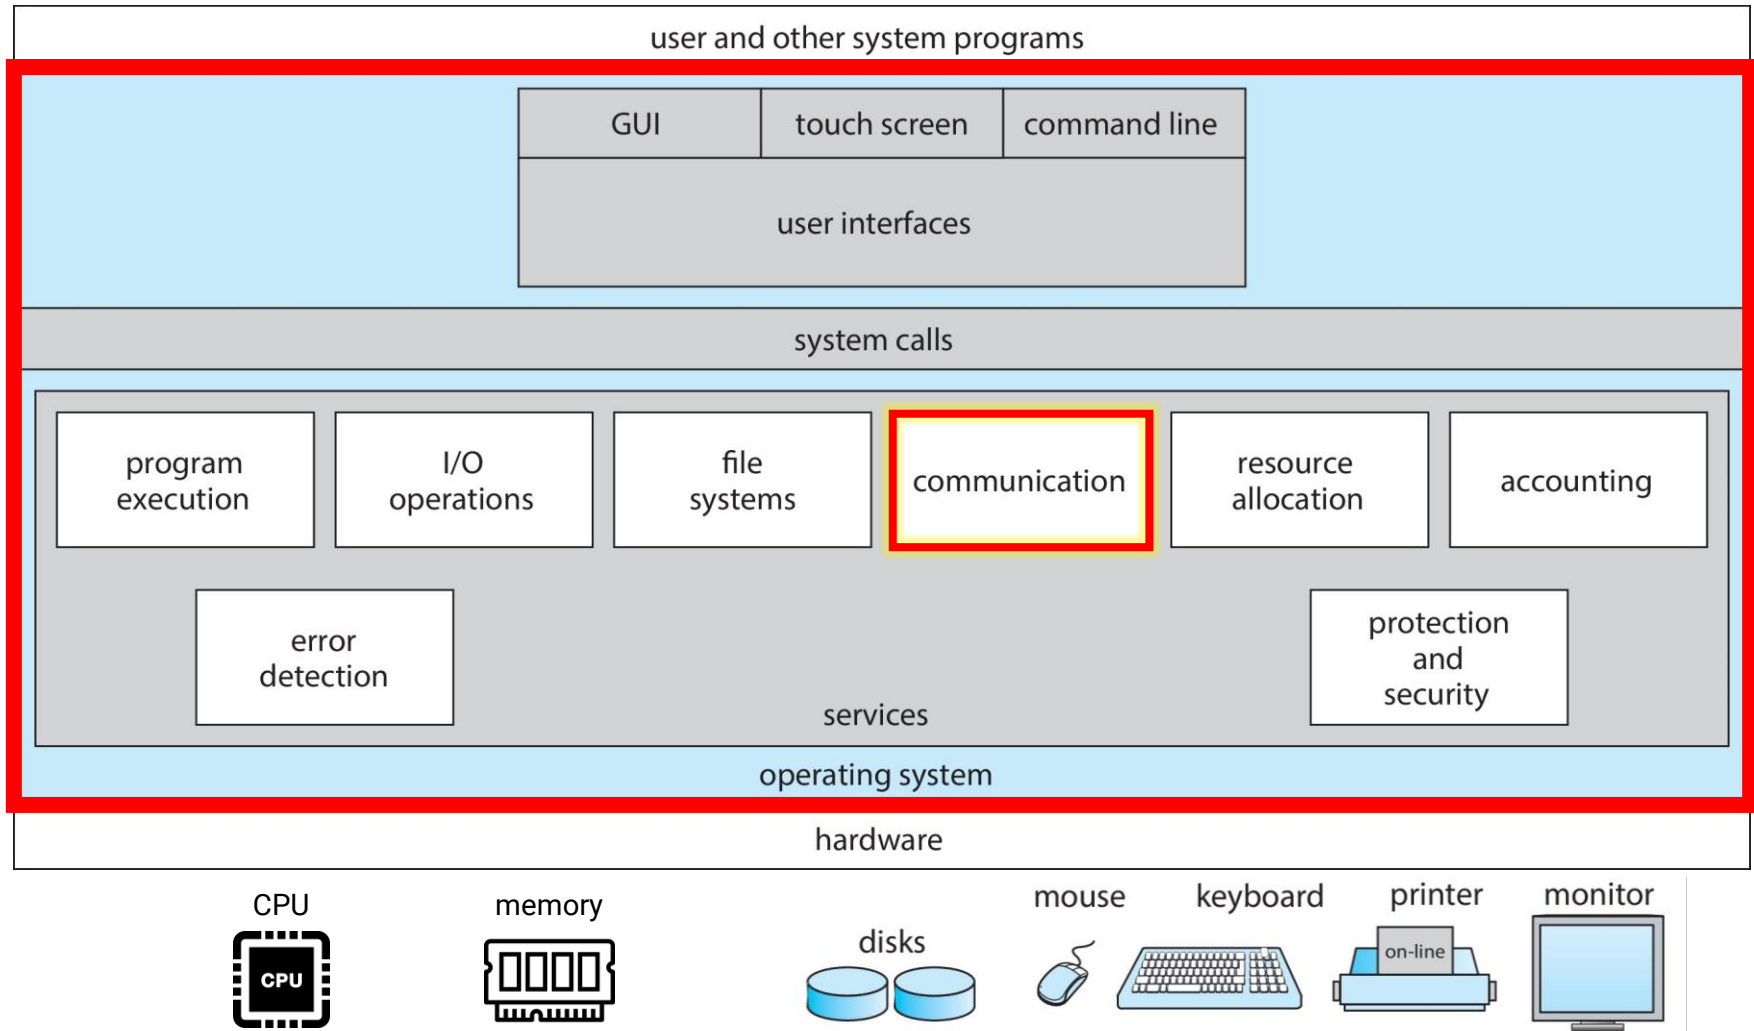
- **Dynamic-priority scheduler**
  - Task's priority is not fixed
  - Task's priority is determined by deadline
- Example (preempted)
  - $T1 = (2, 0.9)$ ,  $T2 = (5, 2.3)$  for (period, execution)





# **Operating System Service: Inter-Process Communication**

# Operating System Services



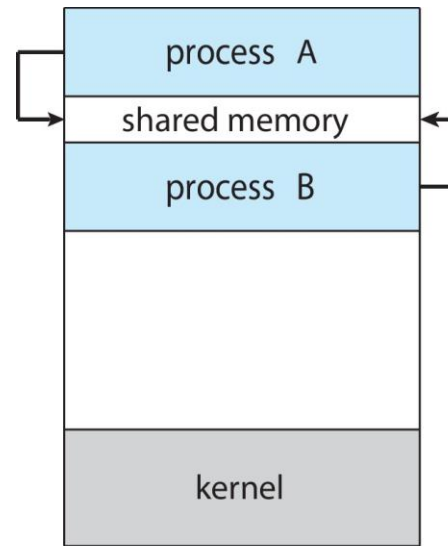
# Communication Methods

- **Shared memory**

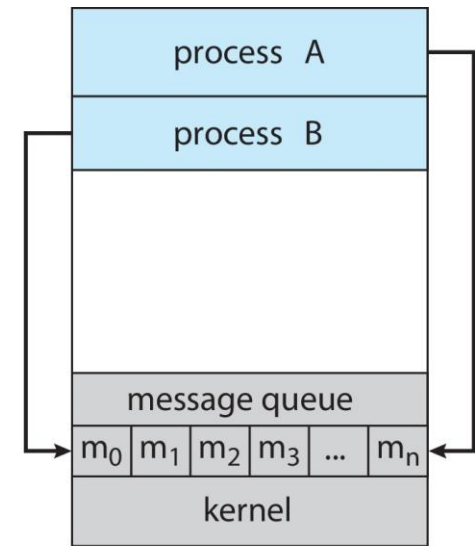
- Require more careful **user synchronization**
- Implemented by memory access (faster)
- Use memory address to access data

- **Message passing**

- No conflict: more efficient for small data
- Use send/recv message
- Implement by system call (slower)



shared memory



message passing

# **Message Passing**

# Message Passing System

- Mechanism for processes to **communicate** and **synchronize** their actions
- IPC provides two operations
  - **Send** (message)
  - **Receive** (message)
- To communicate, processes need to
  - Establish a **communication link**
  - Exchange a message via send/receive


# Message Passing System (cont.)

- Implementation of communication link
  - **Physical**
    - HW bus
    - Network
  - **Logical (properties of the link)**
    - **Direct or indirect communication**
    - Symmetric or asymmetric communication
    - **Blocking or non-blocking**
    - Automatic or explicit buffering
    - Send by copy or send by reference
    - Fixed-sized or variable-sized messages

# Message Passing Methods

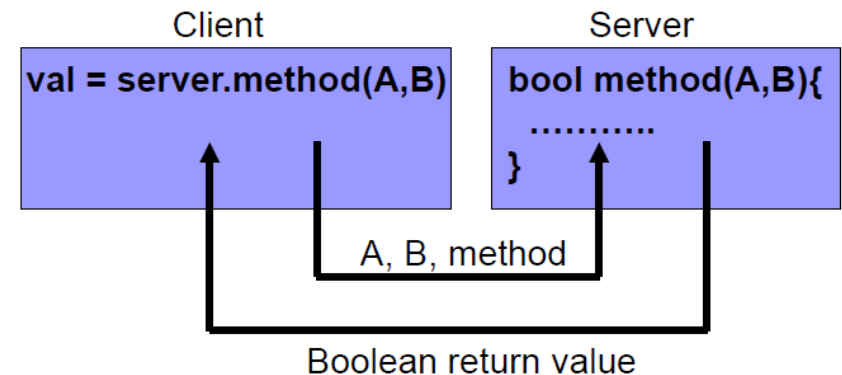
## • Sockets

- A network connection identified by **IP** and **port**
- Exchange **unstructured stream of bytes**



## • Remote Procedure Calls

- Cause a procedure to execute in another address space
- **Parameters** and **return values** are passed by messages



# Blocking and Non-Blocking

- Messages passing may be either **blocking** or **non-blocking**
- **Blocking (synchronous)**
  - **Blocking send**: sender is blocked until the message is received by receiver or by the mailbox
  - **Blocking receive**: receiver is blocking until the message is available
- **Non-blocking (asynchronous)**
  - **Non-blocking send**: sender sends the message and resumes operation
  - **Non-blocking receive**: receiver receives a valid message or a null




# Shared Memory

# Shared Memory

- Processes are responsible for
  - **Establishing a region of shared memory** (ask OS for help)
    - Typically, the created shared-memory regions resides in the address of the process creating the shared memory segment
    - Participating processes must agree to remove memory access constraint from OS
  - **Determining the form of the data and the location**
  - **Synchronization**: ensuring data are not written simultaneously by processes

# Consumer and Producer

- **Producer** process produces information that is consumed by a **Consumer** process
- Buffer as a circular array with size  $B$ 
  - Next free: *in*
  - First available: *out*
  - Empty:  $in = out$
  - Full:  $(in + 1) \% B = out$



- The solution allows at most  $(B - 1)$  item in the buffer
  - Otherwise, cannot tell the buffer is empty or full

# That's All !

